

The Cyclops Forecast Suite Metascheduler User Guide

4.2.0

Released Under the GNU GPL v3.0 Software License

Copyright Hilary Oliver, NIWA, 2008-2011

Hilary Oliver

January 31, 2012



Abstract

Cylc (“silk”) is a metascheduler¹ for cycling environmental forecasting suites containing many forecast models and associated processing tasks. Cylc has a novel self-organising scheduling algorithm: a pool of task proxy objects, that each know just their own inputs and outputs, negotiate dependencies so that correct scheduling emerges naturally at run time. Cylc does not group tasks artificially by forecast cycle² (each task has a private cycle time and is self-spawning - there is no suite-wide cycle time) and handles dependencies within and between cycles equally so that tasks from multiple cycles can run at once to the maximum possible extent. This matters in particular whenever the external driving data³ for upcoming cycles are available in advance: cylc suites can catch up from delays very quickly, parallel test suites can be started behind the main operation to catch up quickly, and one can likewise achieve greater throughput in historical case studies; the usual sequence of distinct forecast cycles emerges naturally if a suite catches up to real time operation. Cylc can easily use existing tasks and can run suites distributed across a heterogenous network. Suites can be stopped and restarted in any state of operation, and they dynamically adapt to insertion and removal of tasks, and to delays or failures in particular tasks or in the external environment: tasks not directly affected will carry on cycling as normal while the problem is addressed, and then the affected tasks will catch up as quickly as possible. Cylc has comprehensive command line and graphical interfaces, including a dependency graph based suite control GUI. Other notable features include suite databases; a fast simulation mode; a structured, validated suite definition file format; dependency graph plotting; task event hooks for centralized alerting; and cryptographic suite security.

¹A metascheduler determines when dependent jobs are *ready to run* and then submits them to run by other means, usually a batch queue scheduler. The term can also refer to an aggregate view of multiple distributed resource managers, but that is not the topic of this document. We drop the “meta” prefix from here on because a metascheduler is also a type of scheduler.

²A *forecast cycle* comprises all tasks with a common *cycle time*, i.e. the analysis time or nominal start time of a forecast model, or that of the associated forecast model(s) for other tasks.

³Forecast suites are typically driven by real time observational data or timely model fields from an external forecasting system.

Contents

1	Introduction: How Cylc Works	10
1.1	Scheduling Forecast Suites	10
1.2	EcoConnect	10
1.3	Dependence Between Tasks	10
1.3.1	Intracycle Dependence	10
1.3.2	Intercycle Dependence	13
1.4	The Cylc Scheduling Algorithm	16
2	Cylc Screenshots	16
3	Required Software	23
3.1	Known Compatibility Issues	23
3.2	Other Software Used Internally By Cylc	23
4	Installation	24
4.1	Install The External Packages	24
4.2	Install The Cylc Release	24
4.3	Configure Your Environment	24
4.4	Create The Central Suite Database	24
4.5	Automated Database Test	25
4.6	Automated Scheduler Test	25
4.7	Complete Non-System-Level Installation	25
4.7.1	Pyro	26
4.7.2	Create The Central Suite Database	26
4.7.3	Graphviz	26
4.7.4	Pygraphviz	27
4.7.5	Jinja2	27
4.8	What Next?	28
4.9	Upgrading To New Cylc Versions	28
4.10	Cross-Version Suite Compatibility	28
4.10.1	Version Compatibility Limitations	29
5	On The Meaning Of <i>Cycle Time</i> In Cylc	29
6	Quick Start Guide	29
6.1	Configure Your Environment	29
6.2	Starting The gcylc GUI	29
6.3	Import The QuickStart Suites	30
6.4	View The QuickStart.a Suite Definition	30
6.5	Plotting The QuickStart.a Dependency Graph	31
6.6	Run The QuickStart.a Suite	32
6.6.1	Viewing The State Of Tasks	37
6.6.2	Triggering Tasks Manually	37
6.6.3	Suite Shut-Down And Restart	37
6.7	QuickStart.b - Handling Cold-Starts Properly	38
6.8	QuickStart.c - Real Task Implementations	39
6.9	Monitoring Running Suites	41
6.9.1	Suite std{out,err}	41

6.9.2	Suite Logs	41
6.9.3	Task stdout And stderr Logs	41
6.10	Searching A Suite	42
6.11	Comparing Suites	43
6.12	Validating A Suite	43
6.13	Other Example Suites	44
6.13.1	Choosing The Initial Cycle Time	45
7	Suite Registration	45
7.1	Private Suite Database	46
7.1.1	Private Database Location	46
7.2	Central Suite Database	46
7.2.1	Central Database Location	47
7.3	Database Operations	47
8	Suite Definition	48
8.1	Suite Definition Directories	48
8.2	Suite.rc File Overview	49
8.2.1	Syntax	49
8.2.2	Include-Files	50
8.2.2.1	Editing Temporarily Inlined Suites	50
8.2.2.2	Include-Files via Jinja2	50
8.2.3	Syntax Highlighting In Vim	50
8.2.4	Gross File Structure	50
8.2.5	Validation	51
8.3	Scheduling - Dependency Graphs	51
8.3.1	Graph String Syntax	52
8.3.2	Interpreting Graph Strings	52
8.3.2.1	Handling Long Graph Lines	54
8.3.3	Graph Types (VALIDITY)	54
8.3.3.1	One-off Asynchronous Tasks	54
8.3.3.2	Cycling Tasks	55
8.3.3.2.1	How Multiple Graph Strings Combine	55
8.3.3.3	Combined Asynchronous And Synchronous Graphs	55
8.3.3.3.1	Synchronous Start-up vs One-off Asynchronous Tasks	56
8.3.3.4	Repeating Asynchronous Tasks	56
8.3.4	Trigger Types	57
8.3.4.1	Success Triggers	57
8.3.4.2	Failure Triggers	58
8.3.4.3	Internal Triggers	58
8.3.4.4	Intercycle Triggers	58
8.3.4.5	Conditional Triggers	58
8.3.4.6	Suicide Triggers	59
8.3.4.7	Family Triggers	59
8.3.5	Handling Intercycle Dependence At Start-Up	61
8.3.5.1	Cold-Start Tasks	61
8.3.5.2	Warm-Starting A Suite	61
8.3.6	Model Restart Dependencies	61
8.4	Runtime Properties - Task Execution	63

8.4.1	Namespace Names	64
8.4.2	Root - Runtime Defaults	64
8.4.3	Defining Multiple Namespaces At Once	64
8.4.4	Task Execution Environment	65
8.4.4.1	User Environment Variables	65
8.4.4.2	Overriding Environment Variables	65
8.4.4.3	Suite And Task Identity Variables	66
8.4.4.4	Suite Share And Task Work Directories	66
8.4.4.5	Environment Variable Evaluation	66
8.4.5	Remote Task Hosting	67
8.4.5.1	Remote Log Directories	67
8.5	Visualization	68
8.5.1	Collapsible Task Families In Suite Graphing And GUIs	68
8.6	Jinja2 Suite Templates	68
8.7	Special Placeholder Variables	71
8.8	Omitting Defined Tasks At Runtime	72
9	Task Implementation	72
9.1	Most Tasks Require No Modification For Cylc	72
9.2	Suite.rc Inlined Tasks	73
9.3	Voluntary Messaging Modifications	73
9.4	Tasks Requiring Modification For Cylc	74
9.4.1	Returning Non-zero Exit Status On Error	74
9.4.2	Reporting Internal Outputs Completed	74
9.4.3	Reconnecting Detaching Processes	74
9.4.4	Tasks That Detach And Exit Early	75
9.4.5	A Custom Task Wrapper Example	75
10	Task Job Submission	77
10.1	Task Job Scripts	77
10.2	Built-in Job Submission Methods	79
10.2.1	background	79
10.2.2	at	79
10.2.3	loadleveler	79
10.2.4	pbs	80
10.2.5	sge	80
10.2.6	Default Directives Provided	80
10.2.7	PBS and SGE Cylc Quirks	80
10.3	Whither Task stdout And stderr?	81
10.4	Overriding the Job Submission Command	81
10.5	Defining New Job Submission Methods	81
10.5.1	An Example	83
10.5.2	Where To Put New Job Submission Modules	83
11	Running Suites	84
11.1	Task States	84
11.2	Limiting Suite Activity	85
11.2.1	The Suite Runahead Limit	85
11.2.2	Internal Queues	85
11.3	Other Topics In Brief	86

12 Suite Design Principles	87
12.1 Make Fine-Grained Suites	87
12.2 Make Tasks Rerunnable	87
12.3 Make Models Rerunnable	87
12.4 Limit Previous-Instance Dependence	88
12.5 Put Task Cycle Time In All Output File Paths	88
12.5.1 Use Cylc's Cycle Time Filename Template Utility	88
12.6 How To Manage Input/Output File Dependencies	88
12.7 Use Generic Task Scripts	89
12.8 Make Suites Portable	89
12.9 Make Tasks As Self-Contained As Possible	89
12.10 Make Suites As Self-Contained As Possible	90
12.11 Orderly Product Generation?	90
12.12 Clock-triggered Tasks Wait On External Data	90
12.13 Do Not Treat Real Time Operation As Special	90
A Suite.rc Reference	91
A.1 Top Level Items	91
A.1.1 title	91
A.1.2 description	91
A.2 [cylc]	91
A.2.1 UTC mode	91
A.2.2 simulation mode only	91
A.2.3 use secure passphrase	92
A.2.4 [[logging]]	92
A.2.4.1 directory	92
A.2.4.2 roll over at start-up	92
A.2.5 [[state dumps]]	92
A.2.5.1 directory	92
A.2.5.2 number of backups	93
A.2.6 [[event hooks]]	93
A.2.6.1 events	93
A.2.6.2 script	93
A.2.7 [[lockserver]]	94
A.2.7.1 enable	94
A.2.7.2 simultaneous instances	94
A.2.8 [[environment]]	94
A.2.8.1 Warnings	94
A.2.8.2 __VARIABLE__	94
A.2.9 [[simulation mode]]	95
A.2.9.1 clock rate	95
A.2.9.2 clock offset	95
A.2.9.3 command scripting	95
A.2.9.4 [[[job submission]]]	95
A.2.9.4.1 method	95
A.2.9.5 [[[event hooks]]]	95
A.2.9.5.1 enable	96
A.3 [scheduling]	96
A.3.1 initial cycle time	96

A.3.2	final cycle time	96
A.3.3	runahead limit	96
A.3.4	[[queues]]	96
	A.3.4.1 [[[_QUEUE_]]]	97
	A.3.4.2 limit	97
	A.3.4.3 members	97
A.3.5	[[special tasks]]	97
	A.3.5.1 clock-triggered	97
	A.3.5.2 start-up	97
	A.3.5.3 cold-start	98
	A.3.5.4 sequential	98
	A.3.5.5 one-off	98
	A.3.5.6 explicit restart outputs	98
	A.3.5.7 exclude at start-up	99
	A.3.5.8 include at start-up	99
A.3.6	[[dependencies]]	99
	A.3.6.1 graph	99
	A.3.6.2 [[[_VALIDITY_]]]	99
	A.3.6.2.1 graph	100
	A.3.6.2.2 daemon	100
A.4	[runtime]	100
A.4.1	[[_NAME_]]	100
	A.4.1.1 inherit	101
	A.4.1.2 description	101
	A.4.1.3 initial scripting	101
	A.4.1.4 command scripting	101
	A.4.1.5 pre-command scripting	102
	A.4.1.6 post-command scripting	102
	A.4.1.7 manual completion	102
	A.4.1.8 [[job submission]]	102
	A.4.1.8.1 method	102
	A.4.1.8.2 command template	103
	A.4.1.8.3 shell	103
	A.4.1.8.4 log directory	103
	A.4.1.8.5 work directory	103
	A.4.1.8.6 share directory	103
	A.4.1.8.9 [[remote]]	104
	A.4.1.9.1 host	104
	A.4.1.9.2 owner	104
	A.4.1.9.3 cylc directory	104
	A.4.1.9.4 suite definition directory	104
	A.4.1.9.5 remote shell template	104
	A.4.1.9.6 log directory	105
	A.4.1.9.7 work directory	105
	A.4.1.9.8 share directory	105
	A.4.1.10 [[event hooks]]	105
	A.4.1.10.1 events	106
	A.4.1.10.2 script	106
	A.4.1.10.3 submission timeout	106

A.4.1.10.4	execution timeout	106
A.4.1.10.5	reset timer	106
A.4.1.11	[[environment]]	107
A.4.1.11.1	_VARIABLE_	107
A.4.1.12	[[directives]]	107
A.4.1.12.1	_DIRECTIVE_	107
A.4.1.13	[[outputs]]	107
A.4.1.13.1	_OUTPUT_	108
A.5	[visualization]	108
A.5.1	initial cycle time	108
A.5.2	final cycle time	108
A.5.3	collapsed families	108
A.5.4	use node color for edges	109
A.5.5	default node attributes	109
A.5.6	default edge attributes	109
A.5.7	[[node groups]]	109
A.5.7.1	_GROUP_	109
A.5.8	[[node attributes]]	109
A.5.8.1	_NAME_	109
A.5.9	[[run time graph]]	110
A.5.9.1	enable	110
A.5.9.2	cutoff	110
A.5.9.3	directory	110
A.6	Special Placeholder Variables	110
A.7	Default Suite Configuration	111
B	Command Reference	113
B.1	Command Categories	113
B.1.1	admin	113
B.1.2	all	113
B.1.3	control	115
B.1.4	database	115
B.1.5	discovery	115
B.1.6	hook	115
B.1.7	information	116
B.1.8	license	116
B.1.9	preparation	116
B.1.10	task	116
B.1.11	utility	117
B.2	Commands	117
B.2.1	alias	117
B.2.2	block	117
B.2.3	browse	117
B.2.4	check-examples	118
B.2.5	checkvars	118
B.2.6	conditions	118
B.2.7	copy	118
B.2.8	create-cdb	119
B.2.9	cycletime	120

B.2.10 depend	120
B.2.11 diff	120
B.2.12 dump	121
B.2.13 edit	121
B.2.14 email-suite	122
B.2.15 email-task	123
B.2.16 export	123
B.2.17 failed	123
B.2.18 gcyclc	123
B.2.19 get-config	124
B.2.20 get-directory	124
B.2.21 graph	125
B.2.22 hold	125
B.2.23 housekeeping	126
B.2.24 import	127
B.2.25 insert	127
B.2.26 jobscontrol	127
B.2.27 list	128
B.2.28 lockclient	128
B.2.29 lockserver	129
B.2.30 log	129
B.2.31 maxrunahead	129
B.2.32 message	130
B.2.33 monitor	130
B.2.34 nudge	130
B.2.35 ping	131
B.2.36 print	131
B.2.37 purge	131
B.2.38 refresh	132
B.2.39 register	132
B.2.40 release	133
B.2.41 remove	133
B.2.42 reregister	134
B.2.43 reset	134
B.2.44 restart	134
B.2.45 run	135
B.2.46 scan	136
B.2.47 scp-transfer	136
B.2.48 search	137
B.2.49 show	137
B.2.50 started	137
B.2.51 stop	137
B.2.52 submit	138
B.2.53 succeeded	138
B.2.54 template	138
B.2.55 test-db	139
B.2.56 test-suite	139
B.2.57 trigger	140
B.2.58 unblock	140

B.2.59	unregister	140
B.2.60	validate	141
B.2.61	verbosity	141
B.2.62	view	141
B.2.63	warranty	142
C	The Cyc Lockserver	142
D	The Graph-Based Suite Control GUI	143
E	Simulation Mode	143
E.1	Clock Rate and Offset	144
E.2	Switching A Suite Between Simulation And Live Modes?	144
F	Cyc Development History	144
F.1	Pre-3.0	144
F.2	Version 3.0	145
F.3	Version 4.0	145
G	Pyro	145
H	Troubleshooting	145
H.1	pre-3.9.1 Pyro Incompatibility	145
I	Acknowledgements	146
J	GNU GENERAL PUBLIC LICENSE v3.0	146

List of Figures

1	A single cycle dependency graph for a simple suite	11
2	A single cycle job schedule for real time operation	11
3	What if the external driving data is available early?	12
4	Attempted overlap of consecutive single-cycle job schedules	12
5	The only safe multicycle job schedule?	12
6	The complete multicycle dependency graph	14
7	The optimal two-cycle job schedule	14
8	Comparison of job schedules after a delay	14
9	Optimal job schedule when all external data is available	15
10	The cyc task pool	15
11	The cyc command line interface	17
12	gycyc showing a private suite database	18
13	gycyc showing the central suite database	18
14	A simple cyc suite definition edited in <i>vim</i>	19
15	A cyc suite definition with embedded Jinja2 expressions <i>vim</i>	19
16	The suite definition of Figure 14 graphed by cyc.	20
17	The suite control GUI, treeview interface.	20
18	The suite control GUI, graph interface.	21
19	A large operational suite graphed by cyc.	21
20	The large suite of Figure 19 running	22
21	The <i>QuickStart.a</i> dependency graph	32

22	Suite <i>QuickStart.a</i> at start-up, 06 or 18 hours	33
23	Suite <i>QuickStart.a</i> at start-up, 00 or 12 hours	34
24	Suite <i>QuickStart.a</i> running	35
25	Suite <i>QuickStart.a</i> stalled	36
26	Viewing current task state in gcylc	37
27	The <i>QuickStart.b</i> graph with model cold-start task.	38
28	Cylc suite stdout/stderr example.	41
29	A cylc suite log viewed via gcylc.	42
30	Example Suite SDepG	52
31	One-off Asynchronous Tasks	54
32	Cycling Tasks	55
33	One-off Asynchronous and Cycling Tasks	56
34	One-off Synchronous and Cycling Tasks	57
35	Repeating Asynchronous Tasks	57
36	Conditional Triggers	59
37	Automated failure recovery via suicide triggers	60
38	<i>namespaces</i> example suite graphs	69
39	The Ninja2 ensemble example suite graph.	70
40	Jinja2 cities example suite graph.	71

1 INTRODUCTION: HOW CYLC WORKS

1 Introduction: How Cylc Works

1.1 Scheduling Forecast Suites

Environmental forecasting suites generate forecast products from a potentially large group of interdependent scientific models and associated data processing tasks. They are constrained by availability of external driving data: typically one or more tasks will wait on real time observations and/or model data from an external system, and these will drive other downstream tasks, and so on. The dependency diagram for a single forecast cycle in such a system is a *Directed Acyclic Graph* as shown in Figure 1 (in our terminology, a *forecast cycle* is comprised of all tasks with a common *cycle time*, which is the nominal analysis time or start time of the forecast models in the group). In real time operation processing will consist of a series of distinct forecast cycles that are each initiated, after a gap, by arrival of the new cycle's external driving data.

From a job scheduling perspective task execution order in such a system must be carefully controlled in order to avoid dependency violations. Ideally, each task should be queued for execution at the instant its last prerequisite is satisfied; this is the best that can be done even if queued tasks are not able to execute immediately because of resource contention.

1.2 EcoConnect

Cylc was developed for the EcoConnect Forecasting System at NIWA (National Institute of Water and Atmospheric Research, New Zealand). EcoConnect takes real time atmospheric and stream flow observations, and operational global weather forecasts from the Met Office (UK), and uses these to drive global sea state and regional data assimilating weather models, which in turn drive regional sea state, storm surge, and catchment river models, plus tide prediction, and a large number of associated data collection, quality control, preprocessing, postprocessing, product generation, and archiving tasks.⁴ The global sea state forecast runs once daily. The regional weather forecast runs four times daily but it supplies surface winds and pressure to several downstream models that run only twice daily, and precipitation accumulations to catchment river models that run on an hourly cycle assimilating real time stream flow observations and using the most recently available regional weather forecast. EcoConnect runs on heterogenous distributed hardware, including a massively parallel supercomputer and several Linux servers.

1.3 Dependence Between Tasks

1.3.1 Intracycle Dependence

Most inter-task dependence exist within a single forecast cycle. Figure 1 shows the dependency diagram for a single forecast cycle of a simple example suite of three forecast models (*a*, *b*, and *c*) and three post processing or product generation tasks (*d*, *e* and *f*). A scheduler capable of handling this must manage, within a single forecast cycle, multiple parallel streams of execution that branch when one task generates output for several downstream tasks, and merge when one task takes input from several upstream tasks.

Figure 2 shows the optimal job schedule for two consecutive cycles of the example suite in real time operation, given execution times represented by the horizontal extent of the task bars. There is a time gap between cycles as the suite waits on new external driving data. Each

⁴Future plans for EcoConnect include additional deterministic regional weather forecasts and a statistical ensemble.

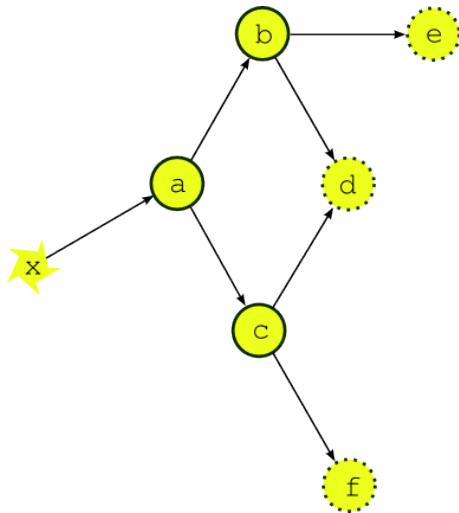


Figure 1: The dependency graph for a single forecast cycle of a simple example suite. Tasks a , b , and c represent forecast models, d , e and f are post processing or product generation tasks, and x represents external data that the upstream forecast model depends on.

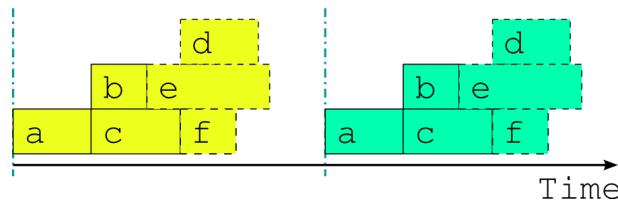


Figure 2: The optimal job schedule for two consecutive cycles of our example suite during real time operation, assuming that all tasks trigger off upstream tasks finishing completely. The horizontal extent of a task bar represents its execution time, and the vertical blue lines show when the external driving data becomes available.

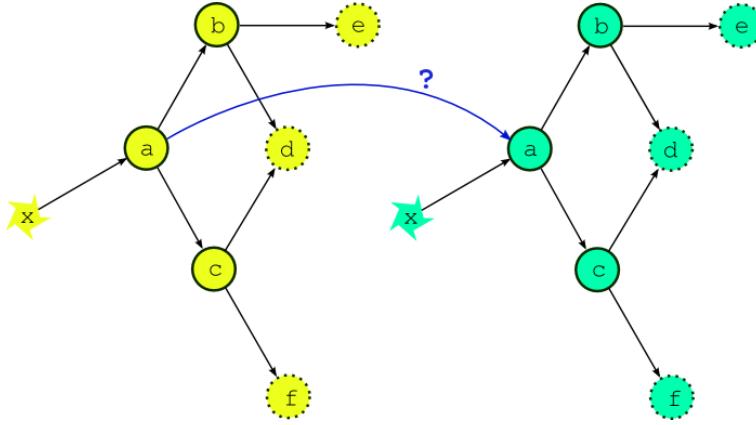


Figure 3: If the external driving data is available in advance, can we start running the next cycle early?

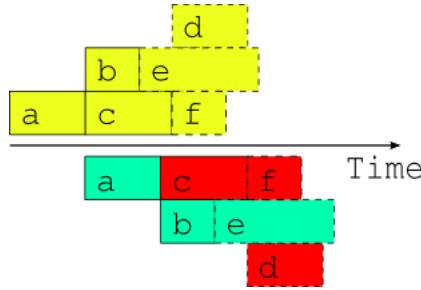


Figure 4: A naive attempt to overlap two consecutive cycles using the single-cycle dependency graph. The red shaded tasks will fail because of dependency violations (or will not be able to run because of upstream dependency violations).

task in the example suite happens to trigger off upstream tasks *finishing*, rather than off any intermediate output or event; this is merely a simplification that makes for clearer diagrams.

Now the question arises, what happens if the external driving data for upcoming cycles is available in advance, as it would be after a significant delay in operations, or when running a historical case study? While the forecast model *a* appears to depend only on the external data *x* at this stage of the discussion, in fact it would typically also depend on its own previous instance for the model *background state* used in initializing the new forecast. Thus, as alluded to in Figure 3, task *a* could in principle start as soon as its predecessor has finished. Figure 4 shows, however, that starting a whole new cycle at this point is dangerous - it results in dependency violations in half of the tasks in the example suite. In fact the situation is even worse than this - imagine that task *b* in the first cycle is delayed for any reason *after* the second cycle has been

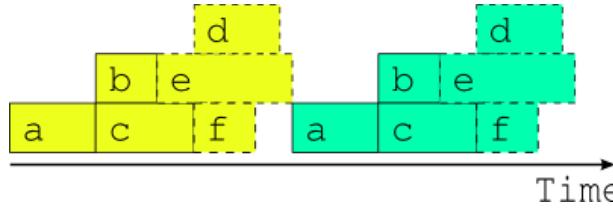


Figure 5: The best that can be done *in general* when intercycle dependence is ignored.

launched? Clearly we must consider handling intercycle dependence explicitly or else agree not to start the next cycle early, as is illustrated in Figure 5.

1.3.2 Intercycle Dependence

In most suites dependence between tasks in different cycles exist. Forecast models, as above, typically depend on their own most recent previous forecast for a background state, and different types of tasks in different forecast cycles can also be linked (in an atmospheric forecast analysis suite, for instance, the weather model may also generate background states for use by the observation processing and data-assimilation systems in the next cycle). In real time operation this intercycle dependence can be ignored because it is automatically satisfied when each cycle finishes before the next one begins. This is just as well because they dramatically increase the complexity of the dependency graph of even the simplest suites, by destroying the clean boundary between forecast cycles. Figure 6 illustrates the problem for our simple example suite assuming the minimal likely intercycle dependence: the forecast models (*a*, *b*, and *c*) each depend on their own previous instances.

For this reason, and because we tend to imagine that forecasting suites always run in distinct cycles, existing metaschedulers (as far as the author is aware!) ignore intercycle dependence and therefore *require* a series of distinct cycles at all times. While this does not affect normal real time operation it can be a serious impediment when advance availability of external driving data makes it possible, in principle, to run some tasks from upcoming cycles before the current cycle is finished - as suggested at the end of the previous section. This occurs after delays (late arrival of external data, system maintenance, etc.) and, to an even greater extent, in historical case studies, and parallel test suites that are delayed with respect to the main operation. It is a serious problem, in particular, for suites that have little downtime between forecast cycles and therefore take many cycles to catch up after a delay. Without taking account of intercycle dependence, the best that can be done, in general, is to reduce the gap between cycles to zero as shown in Figure 5. A limited crude overlap of the single cycle job schedule may be possible for specific task sets but the allowable overlap may change if new tasks are added, and it is still dangerous: it amounts to running different parts of a dependent system as if they were not dependent and as such it cannot be guaranteed that some unforeseen delay in one cycle, after the next cycle has begun, (e.g. due to resource contention or task failures) won't result in dependency violations.

Figure 7 shows, in contrast to Figure 4, the optimal two cycle job schedule obtained by respecting all intercycle dependence. This assumes no delays due to resource contention or otherwise - i.e. every task runs as soon as it is ready to run. The scheduler running this suite must be able to adapt dynamically to external conditions that impact on multicycle scheduling in the presence of intercycle dependence or else, again, risk bringing the system down with dependency violations.

To further illustrate the potential benefits of proper intercycle dependency handling, Figure 8 shows an operational delay of almost one whole cycle in a suite with little downtime between cycles. Above the time axis is the optimal schedule that is possible, in principle, when intercycle dependence is taken into account, and below is the only safe schedule possible *in general* when they are ignored. In the former case, even the cycle immediately after the delay is hardly affected, and subsequent cycles are all on time, whilst in the latter case it takes five full cycles to catch up to normal real time operation.

Similarly, Figure 9 shows example suite job schedules for an historical case study, or when catching up after a very long delay; i.e. when the external driving data are available many cycles in advance. Task *a*, which as the most upstream forecast model is likely to be a resource

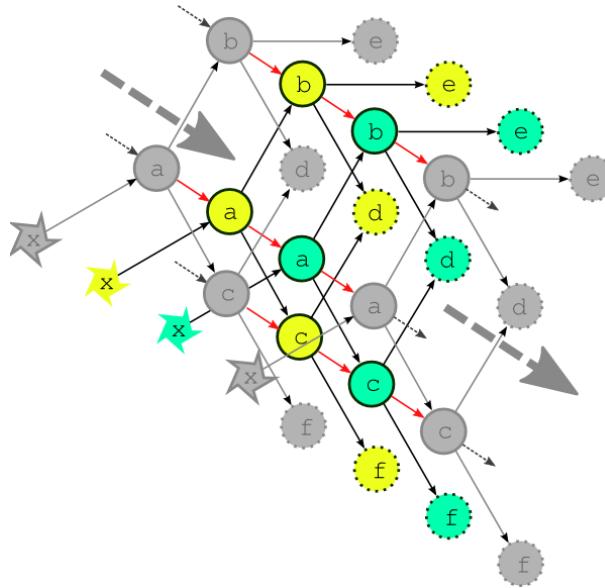


Figure 6: The complete dependency graph for the example suite, assuming the least possible intercycle dependence: the forecast models (a , b , and c) depend on their own previous instances. The dashed arrows show connections to previous and subsequent forecast cycles.

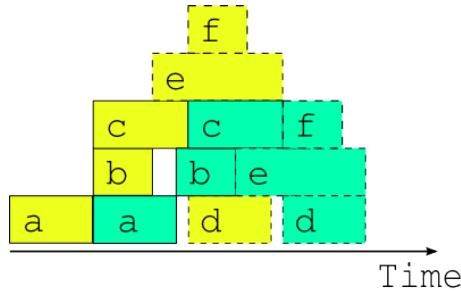


Figure 7: The optimal two cycle job schedule when the next cycle's driving data is available in advance, possible in principle when intercycle dependence is handled explicitly.

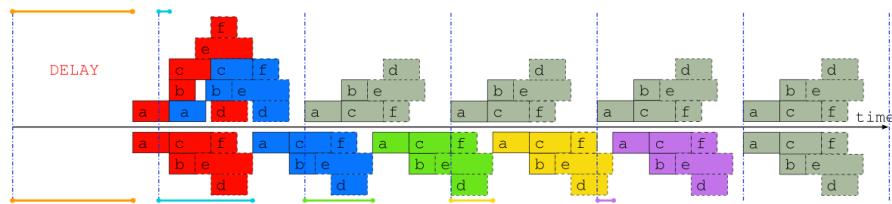


Figure 8: Job schedules for the example suite after a delay of almost one whole forecast cycle, when intercycle dependence is taken into account (above the time axis), and when it is not (below the time axis). The colored lines indicate the time that each cycle is delayed, and normal ‘caught up’ cycles are shaded gray.

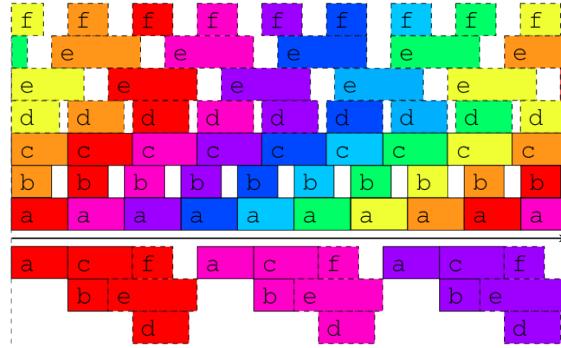


Figure 9: Job schedules for the example suite in case study mode, or after a long delay, when the external driving data are available many cycles in advance. Above the time axis is the optimal schedule obtained when the suite is constrained only by its true dependencies, as in Figure 3, and underneath is the best that can be done, in general, when intercycle dependence is ignored.

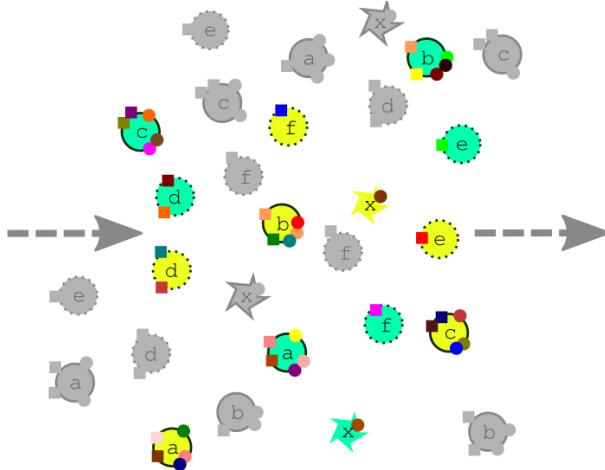


Figure 10: How cylc sees a suite, in contrast to the multicycle dependency graph of Figure 6. Task colors represent different cycle times, and the small squares and circles represent different prerequisites and outputs. A task can run when its prerequisites are satisfied by the outputs of other tasks in the pool.

intensive atmosphere or ocean model, has no upstream dependence on cotemporal tasks and can therefore run continuously, regardless of how much downstream processing is yet to be completed in its own, or any previous, forecast cycle (actually, task *a* does depend on cotemporal task *x* which waits on the external driving data, but that returns immediately when the external data is available in advance, so the result stands). The other forecast models can also cycle continuously or with short gap between, and some post processing tasks, which have no previous-instance dependence, can run continuously or even overlap (e.g. *e* in this case). Thus, even for this very simple example suite, tasks from three or four different cycles can in principle run simultaneously at any given time. In fact, if our tasks are able to trigger off internal outputs of upstream tasks, rather than waiting on full completion, successive instances of the forecast models could overlap as well (because model restart outputs are generally completed early in the forecast) for an even more efficient job schedule.

1.4 The Cylc Scheduling Algorithm

Cylc manages a pool of proxy objects that represent real tasks in the forecasting suite. A task proxy can run the real task that it represents when its prerequisites are satisfied, and can receive reports of completed outputs from the real task as it runs. There is no global cycling mechanism to advance the suite in time; instead each individual task proxy has a private cycle time and spawns its own successor. Task proxies are self-contained - they just know their own prerequisites and outputs and are not aware of the wider suite context. Intercycle dependencies are not treated as special, and the task pool can be populated with tasks from many different cycle times. The cylc task pool is illustrated in Figure 10. Now, *whenever any task changes state due to completion of an output, every task checks to see if its own prerequisites are now satisfied.*⁵ Moreover, this matching of prerequisites and outputs involves the entire task pool, regardless of individual cycle times, so that inter- and intra-cycle dependence is handled with ease.

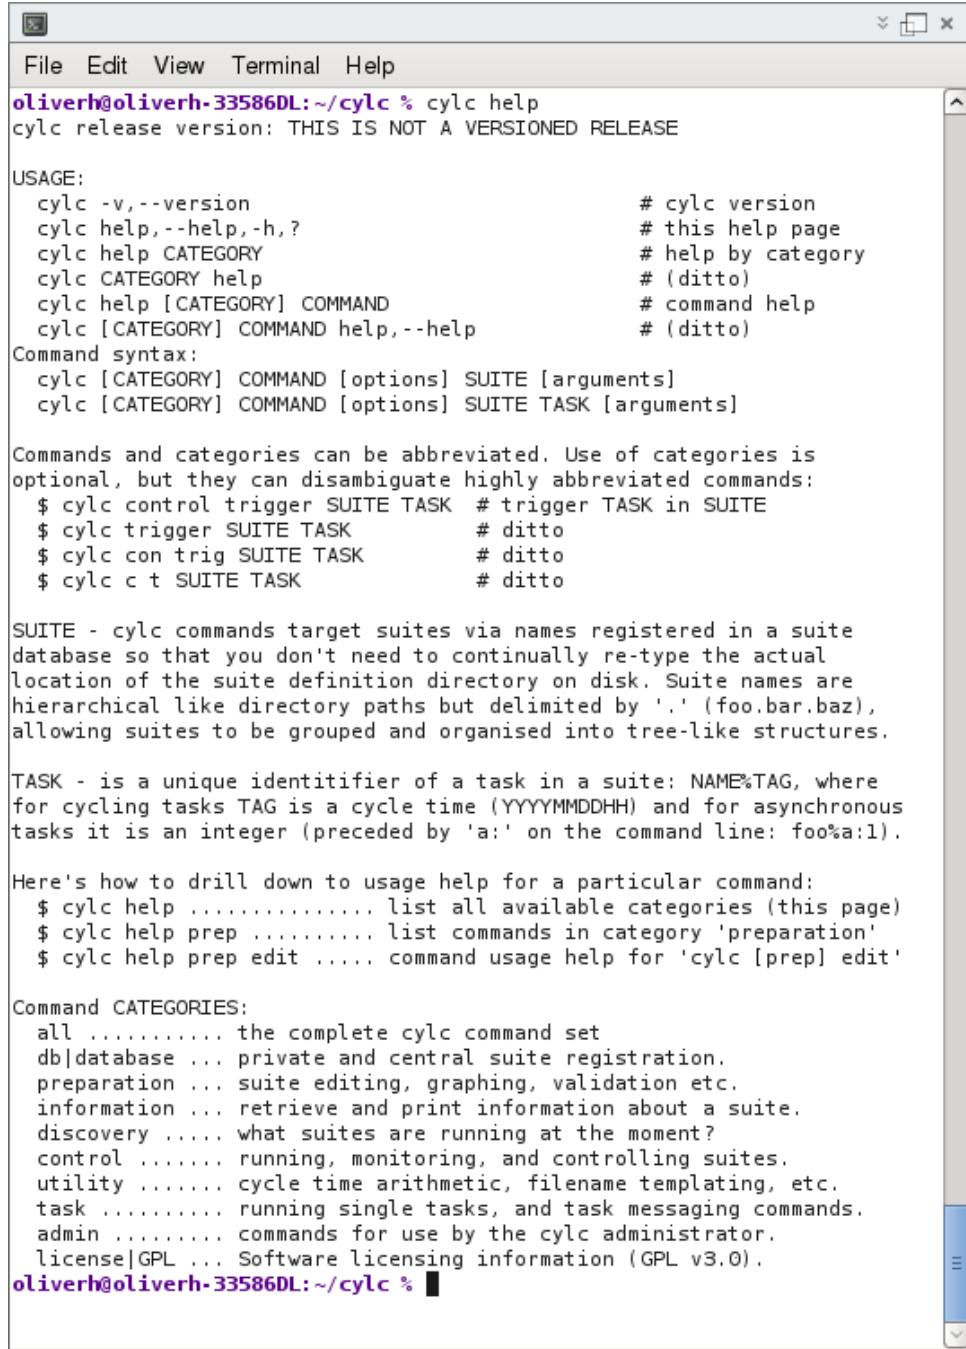
Thus without using global cycling mechanisms, and treating all inter-task dependence equally, cylc in effect gets a pool of tasks to self-organise by negotiating their own dependencies so that optimal scheduling, as described in the previous section, emerges naturally at run time.

2 Cylc Screenshots

- Figure 11 - the cylc command line interface.
- Figure 12 - viewing a private suite database in gcylc.
- Figure 13 - viewing the central suite database in gcylc.
- Figure 14 - a cylc suite definition in the vim editor.
- Figure 15 - a cylc suite definition with embedded Jinja2 expressions, in vim.
- Figure 16 - the suite definition of Figure 14 graphed by cylc.
- Figure 17 - suite control GUI, treeview interface.
- Figure 18 - suite control GUI, graph-based interface.
- Figure 19 - a large operational suite graphed by cylc.
- Figure 20 - another view of the large operational suite of Figure 19.

⁵In fact this dependency negotiation goes through a broker object (rather than every task literally checking every other task) which scales as n (rather than n^2) where n is the number of task proxies in the pool.

2 CYLC SCREENSHOTS



The screenshot shows a terminal window with a light gray background and a dark gray title bar. The title bar contains the text "File Edit View Terminal Help". Below the title bar, the terminal prompt is "oliverh@oliverh-33586DL:~/cylc %". The main area of the terminal displays the help output for the "cylc" command. The output includes:

- cylc release version:** THIS IS NOT A VERSIONED RELEASE
- USAGE:**
 - cylc -v,--version # cylc version
 - cylc help,--help,-h,? # this help page
 - cylc help CATEGORY # help by category
 - cylc CATEGORY help # (ditto)
 - cylc help [CATEGORY] COMMAND # command help
 - cylc [CATEGORY] COMMAND help,--help # (ditto)
- Command syntax:**
 - cylc [CATEGORY] COMMAND [options] SUITE [arguments]
 - cylc [CATEGORY] COMMAND [options] SUITE TASK [arguments]
- Commands and categories can be abbreviated.** Use of categories is optional, but they can disambiguate highly abbreviated commands:
 - \$ cylc control trigger SUITE TASK # trigger TASK in SUITE
 - \$ cylc trigger SUITE TASK # ditto
 - \$ cylc con trig SUITE TASK # ditto
 - \$ cylc c t SUITE TASK # ditto
- SUITE** - cylc commands target suites via names registered in a suite database so that you don't need to continually re-type the actual location of the suite definition directory on disk. Suite names are hierarchical like directory paths but delimited by '.' (foo.bar.baz), allowing suites to be grouped and organised into tree-like structures.
- TASK** - is a unique identifier of a task in a suite: NAME%TAG, where for cycling tasks TAG is a cycle time (YYYYMMDDHH) and for asynchronous tasks it is an integer (preceded by 'a:' on the command line: foo%a:1).
- Here's how to drill down to usage help for a particular command:**
 - \$ cylc help list all available categories (this page)
 - \$ cylc help prep list commands in category 'preparation'
 - \$ cylc help prep edit command usage help for 'cylc [prep] edit'
- Command CATEGORIES:**
 - all the complete cylc command set
 - db|database ... private and central suite registration.
 - preparation ... suite editing, graphing, validation etc.
 - information ... retrieve and print information about a suite.
 - discovery what suites are running at the moment?
 - control running, monitoring, and controlling suites.
 - utility cycle time arithmetic, filename templating, etc.
 - task running single tasks, and task messaging commands.
 - admin commands for use by the cylc administrator.
 - license|GPL ... Software licensing information (GPL v3.0).

Figure 11: The cylc command line interface

2 CYLC SCREENSHOTS

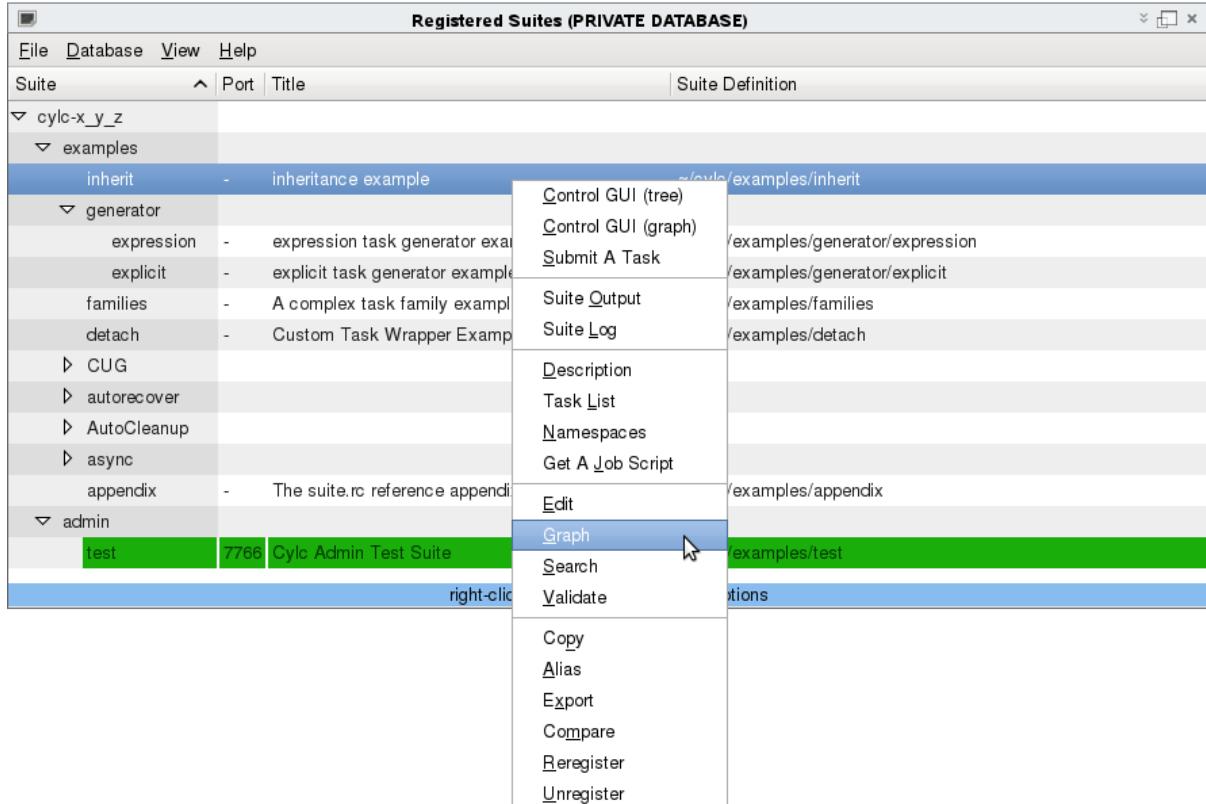


Figure 12: g cylc showing a private suite database, with one suite running on port 7766.

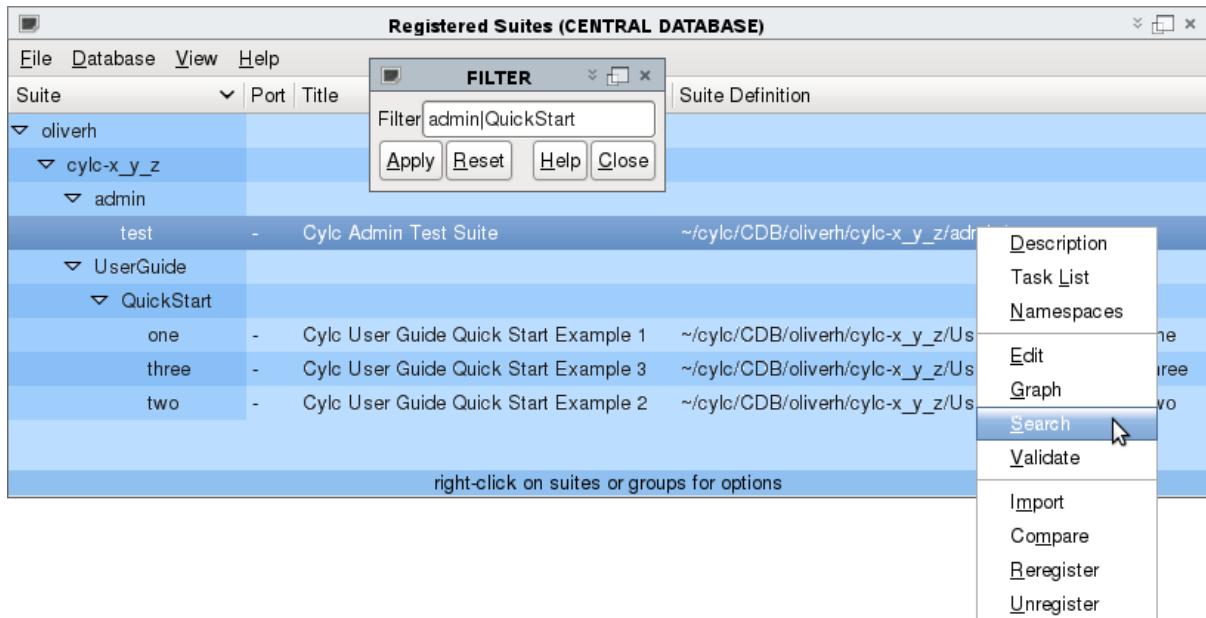


Figure 13: g cylc showing the central suite database.

2 CYLC SCREENSHOTS

```
suite.rc + (tmp/oliver/cyclic-4X-1322609804/QuickStart/c) - GVIM
File Edit Tools Syntax Buffers Window Help

1 title = 'Quick Start Example C'
2 description = '(Quick Start b plus real tasks)*'
3
4 # A comment ...
5
6 [scheduling]
7   runnable_limit = 12
8   initial_cycle_time = 2011010106
9   [[special_tasks]]
10    start-up           = Prep
11    cold-start         = ColdModel
12    [[cold-start]]     = GetData(1)
13   [[dependencies]]
14    [[[6,12,18]]]
15    graph = """Prep => GetData & ColdModel
16          GetData => Model => PostA
17          ColdModel | Model[T-6] => Model***"""
18 3 lines:   [[[6,18]]] .....
19 [runtime]
20 [[root]]
21 [[environment]]
22   TASK_EXE_SECONDS = 5
23   WORKSPACE = /tmp/$USER/$CYLC_SUITE_REG_NAME/common
24
25 4 lines:   [[Prep]] .....
26 6 lines:   [[GetData]] .....
27 [[Models]]
28 [[environment]]
29   MODEL_INPUT_DIR = WORKSPACE
30   MODEL_OUTPUT_DIR = $WORKSPACE
31   MODEL_RUNNING_DIR = $WORKSPACE/Model
32
33 4 lines:   [[ColdModel]] .....
34 [[Model]]
35   inherit = Models
36   description = "the forecast model"
37   command_scripting = Model.sh
38
39 5 lines:   [[Post]] .....
40 4 lines:   [[PostA,PostB]] .....
41 12 lines:[visualization] .....
-
-
```

Figure 14: A simple cyc suite definition edited in *vim*

```
suite.rc ([tmp/oliver/cyc4-0-0-rc4X-1322609804]injia2) - GVIM
File Edit Icons Syntax Buffers Window Help

2 % set HOST = "SuperComputer" %
3 % set CITIES = "NewYork", 'Philadelphia', 'Newark', 'Houston', 'SantaFe', 'Chic
4 % set CITYS = 'one', 'two', 'three', 'four' %
5 % set WALL_CLOCK_LIMIT_MINS = 20 %
6 %
7 %{ set CLEANUP = True %
8 %
9 [scheduler]
10   [[ dependencies ]]
11   (% if CLEANUP %
12     [[23]]!
13     graph = "clean"
14   endif %)
15   [[{0,1,2}]]!
16   graph = ***
17   setup => get_lbc & get_ic # foo
18   (% for CITY in CITIES %) (# comment #)
19     get_lbc => {{ CITY }}_one
20     get_ic => {{ CITY }}_two
21     {{ CITY }}_one & {{ CITY }}_two => {{ CITY }}_three & {{ CITY }}_four
22   (% if CLEANUP %)
23     {{ CITY }}_three & {{ CITY }}_four => cleanup
24   (% endif %)
25   (% endifor %)
26   ***
27
28 [runtime]
29   [[{0..{ HOST } }]]
30   [[{remote}!]]
31   host = {{ HOST }}
32   cyclic_directory = /path/to/cyclic
33   [[{directives}!]]
34   wall_clock_limit = *00:{ WALL_CLOCK_LIMIT_MINS:int() + 2 }:00,0:0:0
35 %
36 (% for CITY in CITIES %)
37   %% ({ CITY }) %
38   inherit = on,{{ CITY }}
39 (% for CITY in CITIES %)
40   [[({ CITY })_{{ 300 }}!]]
41   inherit = { CITY }
42   (% endifor %)
43   (% endifor %)
44
45 [visualization]
46   initial_cycle_time = 2011080812
47   final_cycle_time = 2011080823
48
```

Figure 15: A cyc suite definition with embedded Jinja2 expressions *vim*

2 CYLC SCREENSHOTS

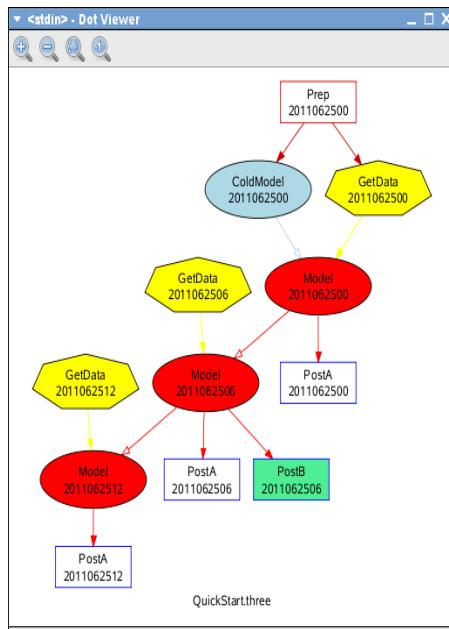


Figure 16: The suite definition of Figure 14 graphed by cylc.

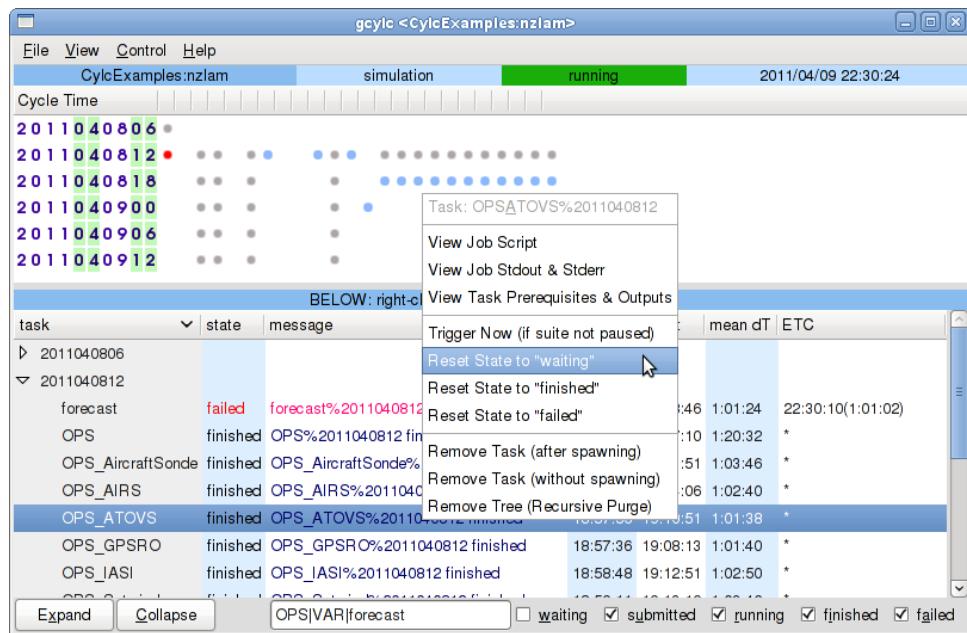


Figure 17: The suite control GUI, treeview interface.

2 CYLC SCREENSHOTS

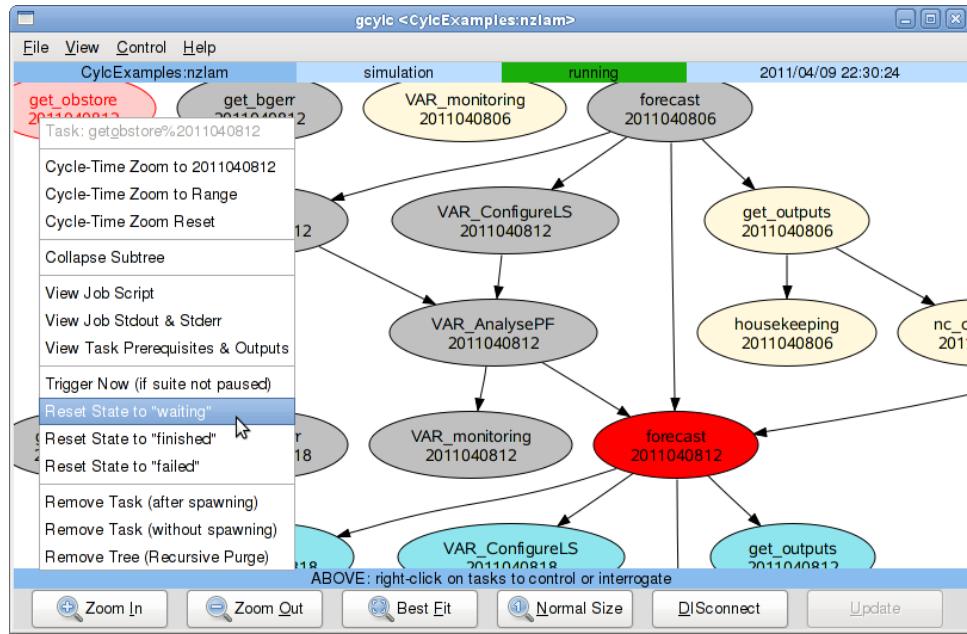


Figure 18: The suite control GUI, graph interface.

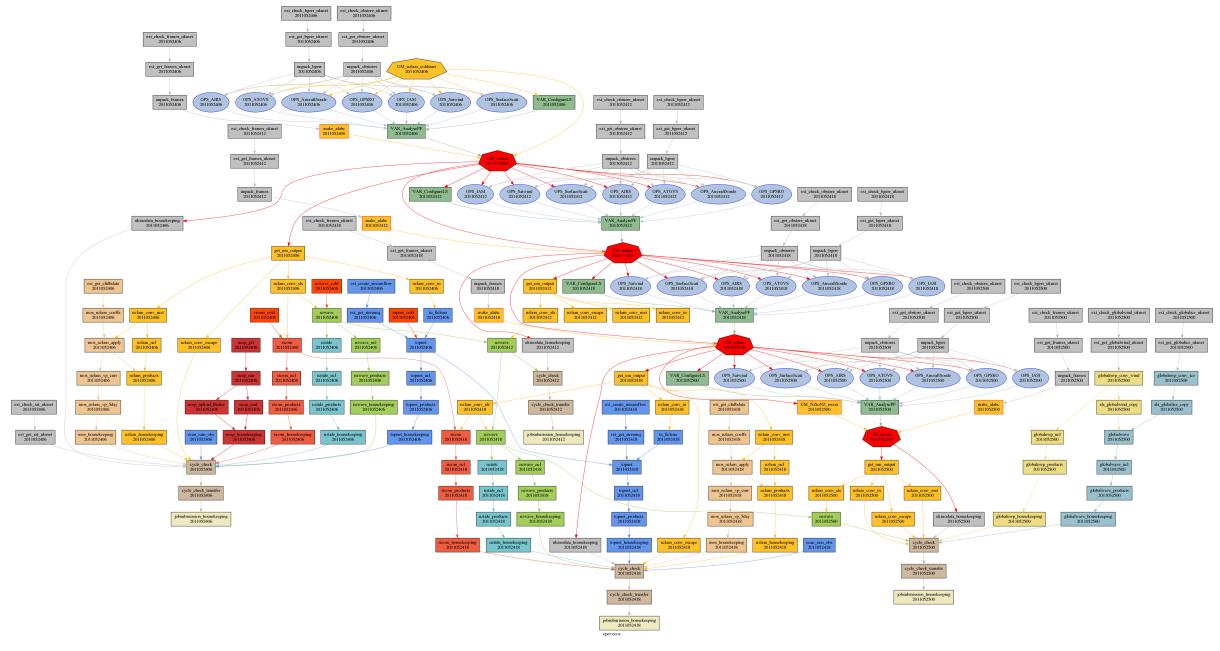


Figure 19: A large operational suite graphed by cylc.

2 CYLC SCREENSHOTS

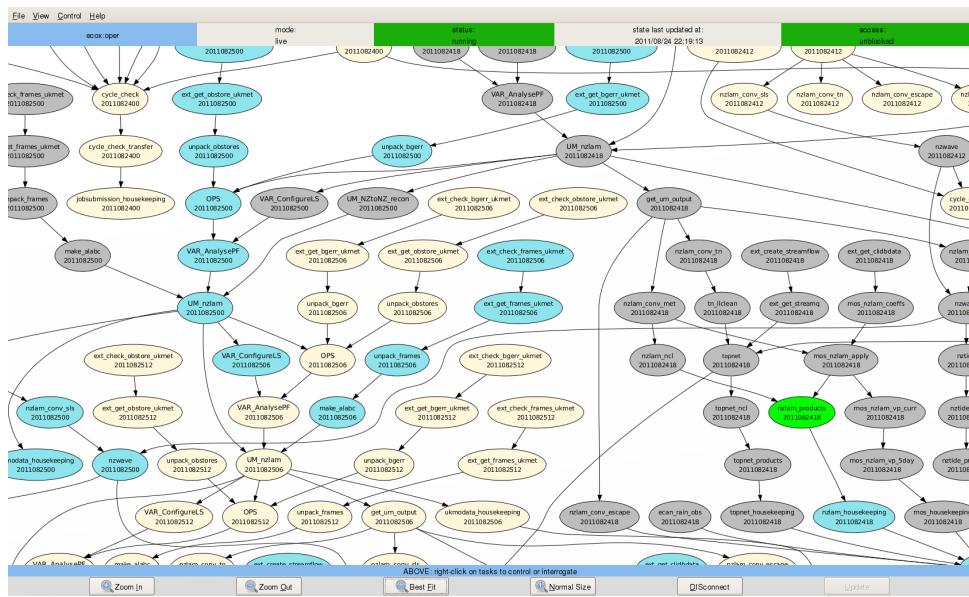


Figure 20: The large suite of Figure 19 running.

3 REQUIRED SOFTWARE

3 Required Software

- **A Linux or Unix OS.** Cylc assumes Unix-style file paths in places, and a few command scripts are written in the Bash shell. In principle it could be made more portable without much effort. Cylc has run successfully on Mac OSX.
- **The Python Language, version 2.4+, latest tested 2.7.2.** Not Python 3.x yet; as of mid 2011 version 2.7 is the standard for new Linux distributions.
<http://python.org>
- **Pyro (Python Remote Objects), version 3.9.1+, latest tested 3.15.** Not Pyro 4.x yet; as of mid 2011 version 3.x is still recommended for production use. Pyro is used by cylc for network communication between server processes (cylc suites) and client programs (running tasks, control GUIs and commands).
<http://irmen.home.xs4all.nl/pyro3>
- **Cylc**, this version: 4.2.0 .
<http://hjoliver.github.com/cylc>

The following packages are technically optional as you can construct and control cylc suites without dependency graphing, the cylc GUIs, and template processing:

- **PyGTK**, a Python wrapper for the GTK+ GUI toolkit, required for the g cylc GUI (but you can prepare and control cylc suites entirely from the command line if you like). PyGTK is included in most Linux Distributions.
<http://www.pygtk.org>
- **Graphviz** (latest tested 2.28.0) and **Pygraphviz** (latest tested 1.1), a graph layout engine and a Python interface to it. These are required for suite dependency graphing and the graph suite control GUI (but you can also run cylc without them).
<http://www.graphviz.org>
<http://networkx.lanl.gov/pygraphviz>
- **Jinja2**, a template processor for Python (latest tested: 2.6). Jinja2 expressions allow use of variables, mathematical expressions, loop control and conditional expressions in suite definitions, to generate the final suite seen by cylc.
<http://jinja.pocoo.org/docs>

3.1 Known Compatibility Issues

Cylc should run out of the box on reasonably recent Linux distributions.

Beware of old Linux distributions that come packaged with Pyro 3.9 or earlier, which is not compatible with the new-style Python classes used in cylc. It has been reported that Ubuntu 10.04 (Lucid Lynx), released in September 2009, suffers from this problem.

Note also that for distributed suites the Pyro versions installed on each machine must be mutually compatible. Using identical versions obviously guarantees compatibility, but this may not be necessary.

3.2 Other Software Used Internally By Cylc

Cylc has absorbed the following in modified form (no need to install these separately):

- xdot, a graph viewer (<http://code.google.com/p/jrfonseca/wiki/XDot>, LGPL license)
- ConfigObj and Validate (<http://www.voidspace.org.uk/python>, BSD license)

4 INSTALLATION

4 Installation

4.1 Install The External Packages

First install Pyro, graphviz, Pygraphviz, and Jinja2 on your system, following the instructions provided with each package. If you do not have root access on your intended cylc host machine and cannot get a sysadmin to do this at system level, see Section 4.7 for local (user-specific) installation instructions.

4.2 Install The Cylc Release

Cylc typically installs into a normal user account; just unpack the release tarball in the desired location - which will be referred to below as `$CYLC_DIR`.

4.3 Configure Your Environment

The file `$CYLC_DIR/cylc.profile` documents the environment configuration required for interactive cylc usage:

```
#!/bin/bash

# CYLC USER LOGIN SCRIPT EXAMPLE. Copy the following to your .profile
# and adapt according to your preferences and local cylc installation.

# These environment variables are required for interactive cylc usage.
# Access to cylc by running tasks is automatically configured by cylc.
#
# Add the cylc bin directory to $PATH. The 'cylc' and 'gcylc' commands
# configure access to cylc sub-commands and python modules at run time.
export PATH=/path/to/cylc/bin:$PATH
#
# For 'cylc edit' or gcylc -> Edit, set terminal and GUI editors:
export EDITOR=vim
export GEDITOR='gvim -f'
# (See 'cylc edit --help' for examples of other editors).
#
# To access the Cylc User Guide via the gcylc GUI menus:
export PDF_READER=evince
export HTML_READER=firefox
# (The HTML guide is opened via file path, not http URL).
#
# Some cylc commands require $TMPDIR to be set and writeable.
export TMPDIR=/path/to/my/temporary/directory
#
# For a local user install of one or more of Pyro, Graphviz, Pygraphviz,
# and Jinja2 (if you can't easily get them installed at system level on
# the cylc host) follow Cylc User Guide Installation instructions on how
# to install them locally, and modify your PYTHONPATH accordingly, e.g.:
## PYTHONPATH=$HOME/external/lib64/python2.6/site-packages:$PYTHONPATH
## export PYTHONPATH
```

Copy this into your `.profile` login script and adapt appropriately for your environment. After sourcing your modified login script, or logging in again, you should be able to run cylc:

```
% cylc --version
x.y.z
```

4.4 Create The Central Suite Database

Cylc has a central suite database, visible to all users on the cylc host, to facilitate sharing of suites. Run the following command immediately after installation to create the central database and export cylc's example suites to it:

```
% cylc admin create-cdb
```

To view the content of the resulting central database, run gcylc and switch to the central database using the Database menu, or use `cylc db print` on the command line:

```
% cylc db print --central --tree -x 'Auto|Quick'
<AdminUserName>
`-cylc-x-y-z
  |-AutoCleanup
  | `|-FamilyFailHook  family failure hook script example
  | `|-FamilyFailTask  family failure cleanup task example
  |-AutoRecover
  | `|-async           asynchronous automated failure recovery example
  | `|-cycling         cycling automated failure recovery example
  `-QuickStart
    |-a               Quick Start Example A
    |-b               Quick Start Example B
    |-c               Quick Start Example C
    `|-z              Quick Start Example Z
```

where `<AdminUserName>` should be replaced with the username of the cylc installation account on your system (the suite owner's username is always the first registration name component for suites in the central database); replace this as appropriate for your local cylc installation. Note that the dots in the cylc release version number are replaced, in central database suite registration names, with hyphens (because '.' is the registration name delimiter). Type `cylc db print --help` to see what the command options mean.

4.5 Automated Database Test

The command `cylc admin test-db` gives suite registration database functionality a work out - it copies one of the cylc example suites, registers it under a new name and then manipulates it by recopying the suite in various ways, exporting it to the central database, and so on, before finally deleting the test registrations. This process should complete without error in a few seconds.

4.6 Automated Scheduler Test

The command `cylc admin test-suite` tests the cylc scheduler itself by running a suite temporarily registered under `testsuite`, configuring it to fail out a specific task, and then doing some advanced failure recovery intervention on it (recursive purge plus insertion of cold-start tasks). This process should complete in 2-3 minutes and can be watched in real time by right-clicking on the temporary test suite when it appears in the gcylc GUI and opening up a suite control GUI.

4.7 Complete Non-System-Level Installation

If you do not have root access to your host machine and cannot easily get Pyro, graphviz, Pygraphviz, and Jinja2 installed at system level, here's how to install everything under your home directory.

First, cylc is already designed to be installed into a normal user account - just unpack the release tarball into `$CYLC_DIR`. If you invoke cylc commands at this stage you will get a warning that Pyro is not installed.

Next, create a new sub-directory in the cylc source tree, `$CYLC_DIR/external`, and download the Pyro, Graphviz, and Pygraphviz source distributions to it (the URLs are given at the beginning of Section 3).

4.7.1 Pyro

Install Pyro under `$HOME/external/installed` as follows:

```
% cd $HOME/external
% tar xzf Pyro-3.14.tar.gz
% cd Pyro-3.14
% python setup.py install --prefix=$HOME/external/installed
```

Take note of the resulting Python `site-packages` directory under `external/installed/`, e.g.:

```
$HOME/external/installed/lib64/python2.6/site-packages/
```

The exact path will depend on your local Python environment. Configure your login script for the cylc environment as described in Section 4.3 above and add in this new installation path; for example:

```
# .profile
PYTHONPATH=$HOME/external/installed/lib64/python2.6/site-packages:$PYTHONPATH
PATH=$HOME/external/installed/bin:$PATH
```

Now you should be able to get cylc to print its release version:

```
% . $HOME/.profile  # (or log in again)
% cylc -v
x.y.z
```

If this command aborts and says that Pyro is not installed or is not available, then you have either not installed Pyro (check the output of the installation command carefully) *or* you have not pointed to the installed Pyro modules in your PYTHONPATH, *or* you have not sourced the cylc environment since updating PYTHONPATH.

Note that Pyro can also be installed with `easy_install`, which downloads and installs python packages in one shot.

4.7.2 Create The Central Suite Database

Now create the suite database and upload the example suites as described in Section 4.4. At this point you should have access to all cylc functionality except for suite graphing and the graph based suite control GUI. For example (change your admin username and cylc version number appropriately, after inspecting the content of the new central database):

```
% cylc info list -tp --central <AdminUserName>.cylc-x-y-z-1322609804.admin.test
root
|-A      model A - weather forecast
|-B      model B - sea state forecast
|-C      model C - river flow forecast
|-ColdA cold start model A
|-ColdB model B
|-ColdC model C
|-D      combined post processing for models B and C
|-E      post processing for model B
|-F      post processing for model F
|-X      retrieve real time data for model A
`-prep  clean out the suite workspace for a new run
```

4.7.3 Graphviz

Install Graphviz under `$CYLC_DIR/external/installed` as follows:

```
% cd $CYLC_DIR/external
% tar xzf graphviz-2.28.0.tar.gz
% cd graphviz-2.28.0
% ./configure --prefix=$CYLC_DIR/external/installed --with-qt=no
% make
% make install
```

This installs graphviz files into the bin, include, and lib sub-directories of your local installation directory. The local installation section of `cylc.profile` (above) provides access to the graphviz executables in this bin directory, although you will probably not need to use them. The graphviz lib and include locations are required when installing Pygraphviz (next).

Note that the graphviz build may fail on systems that do not have QT installed, hence the `./configure --with-qt=no` option above.

4.7.4 Pygraphviz

Install Pygraphviz under `$CYLC_DIR/external/installed` as follows:

```
% cd $CYLC_DIR/external
% tar xzf pygraphviz-1.1.tar.gz
% cd pygraphviz-1.1
```

Now edit setup.py lines 31 and 32 to specify the graphviz lib and include directories:

```
library_path=os.environ['CYLC_DIR'] + '/external/installed/lib'
include_path=os.environ['CYLC_DIR'] + '/external/installed/include/graphviz'
```

Or you can just specify the absolute paths if you like, instead of using the `$CYLC_DIR` environment variable. Check that these are the correct library and include paths by inspecting the contents of the specified directories, and adjust them if necessary. Finally, install pygraphviz:

```
% export CYLC_DIR=/path/to/cylc
% python setup.py install --prefix=$CYLC_DIR/external/installed
```

This may or may not, depending on your local Python setup, install the Pygraphviz modules into the same place as the Pyro modules, e.g.:

```
% ls $CYLC_DIR/external/installed/lib64/python2.6/site-packages/
pygraphviz  pygraphviz-1.1-py2.6.egg-info  Pyro  Pyro-3.14-py2.6.egg-info
```

If not, add the correct Pyraphviz installation path to your PYTHONPATH.

The easiest way to check that pygraphviz has been installed properly is to start an interactive Python session (type `python` after sourcing the cylc environment to configure your PYTHONPATH) then type `import pygraphviz` at the interpreter prompt. If this results in an error message `ImportError: No module named pygraphviz` then either you have not installed pygraphviz properly, *or* you have not configured your PYTHONPATH to point to the installed pygraphviz modules, *or* you have not sourced the cylc environment since updating PYTHONPATH. Finally, if you have installed pygraphviz and configured your PYTHONPATH but graphviz itself has not been installed properly (or if the graphviz libraries have been deleted since you installed pygraphviz) then the initial pygraphiz import will be successful but a lower level import will fail when the pygraphviz modules cannot load the underlying graphviz libraries - in that case, reinstall graphviz.

4.7.5 Jinja2

You can download Jinja2 from the project web site and install it with `python setup.py install --prefix/path/to/` or use the `easy_install` command to do it all in one step. Either way Jinja2 installation requires the final installed package location to be present in the `PYTHONPATH` variable, so you have to arrange for this first. You made also need to create the installed package directory yourself first (if so, the install will abort and print the name of the missing directory in the error message). Here's how to easy_install Jinja2 into your new private python site packages directory:

```
% LOCALPREFIX=$CYLC_DIR/external/installed
% LOCALPACKAGES=$CYLC_DIR/external/installed/lib64/python2.6/site-packages
% export PYTHONPATH=$LOCALPACKAGES:$PYTHONPATH
% easy_install --prefix=$LOCALPREFIX Jinja2
```

Adapt the site-packages path according to your actual path, as above.

4.8 What Next?

You should now have access to all cylc functionality. Create the central suite database and upload the example suites, if you have not done so already (Section 4.4), then test your cylc installation by running the automated suite database test (Section 4.5) and the automated scheduler test (Section 4.6), then go on to the *Quick Start Guide* (Section 6).

4.9 Upgrading To New Cylc Versions

Upgrading is just a matter of unpacking the new cylc release, copying the old central database into the new installation (`$CYLC_DIR/CDB/`),⁶ and running the new `cylc admin create-cdb` to load the new example suites. Finally, change the newly registered central suite group to something sensible using `cylc db reregister` - it will have an ungainly string of digits attached to ensure uniqueness.

4.10 Cross-Version Suite Compatibility

It may not be convenient for users to upgrade all of their suites at once when a new cylc release is installed. As of cylc-4.2.0 any installed version of cylc can act as a transparent front end to any other installed version. This allows construction of new-version suites while older existing ones are still in use, without having to reconfigure your environment or explicitly invoke different cylc versions for different suites.

The required cylc version must be specified on the *first line* of a non-Jinja2 suite definition:

```
#!cylc-4.2.0
# suite definition follows ...
```

And on the *second line* for Jinja2 suites:

```
#!Jinja2
#!cylc-4.2.0
# suite definition follows ...
```

If cylc releases are installed in parallel under the same top level directory then the location of the required cylc installation is automatically computed from that of the invoked one. Otherwise, for unrelated install locations, the full path can be specified:

```
#!/path/to/cylc-4.2.0
# suite definition follows ...
```

This results in the right cylc version being used for any suite-referencing command, e.g.:

```
% cylc validate foo
-----
Cylc cross-version suite compatibility:
  Invoked version: cylc-4.2.0 (/home/oliverh/cylc-4.2.0)
  Suite requires: cylc-4.2.1
Path not given, assuming parallel cylc installations
=> Re-issuing command using /home/oliverh/cylc-4.2.1
-----
Suite foo is valid for cylc-4.2.1
```

Note that non suite-referencing commands necessarily execute exactly as invoked. This includes all command line help so it is advisable, although not strictly necessary, to keep your login environment configured for the latest cylc version.

If the required cylc version is not installed the command will abort and you will have to upgrade the suite definition for compatibility with an installed version: use `cylc validate` and the *Suite.rc Reference* (Appendix A) to identify problems, and consult the cylc change log to see what changed between versions.

⁶The central database is currently associated with the cylc installation; in future we will separate the two.

6 QUICK START GUIDE

4.10.1 Version Compatibility Limitations

The top level command interface, `$CYLC_DIR/bin/cylc`, determines the legality of requested sub-commands prior to version compatibility checking (which is in fact done by the sub-command⁷) so the correct cylc version must be explicitly invoked if the sub-command has undergone a name change between the versions concerned.

5 On The Meaning Of *Cycle Time* In Cylc

From using other schedulers you may be accustomed to the idea that a forecasting suite has a “current cycle time”, which is typically the analysis time or nominal start time of the main forecast model(s) in the suite, and that the whole suite advances to the next forecast cycle when all tasks in the current cycle have finished (or even when a particular wall clock time is reached, in real time operation). As is explained in the Introduction, this is not how cylc works.

Cylc suites advance by means of individual tasks with private cycle times independently spawning successors at the next valid cycle time for the task, not by incrementing a suite-wide forecast cycle. Each task will be submitted when its own prerequisites are satisfied, regardless of other tasks with other cycle times running, or not, at the time. It may still be convenient at times, however, to refer to the “current cycle”, the “previous cycle”, or the “next cycle” and so forth, with reference to a particular task, or in the sense of all tasks that “belong to” a particular forecast cycle. But keep in mind that the members of these groups may not be present simultaneously in the running suite - i.e. different tasks may pass through the “current cycle” (etc.) at different times as the suite evolves, particularly in delayed (catch up) operation.

6 Quick Start Guide

This section works through some basic cylc functionality using the “QuickStart” example suites, which should have been registered in the central suite database, under the cylc admin username (replace `<AdminUserName>` below with the correct username for your system) during cylc installation,

```
% cylc db print --central --tree QuickStart
<AdminUserName>
`-cylc-x-y-z
  `-QuickStart
    |-a      Quick Start Example A | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/a
    |-b      Quick Start Example B | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/b
    |-c      Quick Start Example C | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/c
    `-z      Quick Start Example Z | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/z
```

6.1 Configure Your Environment

To get access to cylc you just need to set a few environment variables in your login script, as described in Section 4.3.

6.2 Starting The g cylc GUI

At the command prompt:

```
% g cylc &
```

⁷This is because the top level command interface just passes options and arguments, which vary according to the sub-command, on to the sub-command; it does not understand them itself and so it cannot determine which one, if any, is the suite reference.

and use the Database menu to switch to the central database, which displays suites accessible to all users (see Figure 13).

6.3 Import The QuickStart Suites

You can register new suites directly in your private database, via `gcylc` or `cylc db register`, or you can import suites from the central database that is visible to all users. Suites in the central database can be viewed, validated, and graphed, etc., but you have to import them in order to run them. Section 7 explains exactly what actions can be performed on suites in central and private databases.

In `gcylc`, find the Quick Start suites (use View → Filter to filter for “QuickStart” if you like), right click on the *QuickStart* group and choose ‘Import’ to copy the whole group to your private database. In the dialog box that pops up, enter ‘QuickStart’ as the TARGET suite name, and `$TMPDIR` (or some other directory of your choice) as a destination for the suite definition directories in the group. Equivalently, on the command line:

```
% cylc db import <AdminUserName>.cylc-x-y-z.QuickStart Quickstart $TMPDIR
COPY /home/<AdminUserName>/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/a
  TO /tmp/oliver/QuickStart/a
REGISTER QuickStart.a: /tmp/oliver/QuickStart/a
COPY /home/<AdminUserName>/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/b
  TO /tmp/oliver/QuickStart/b
REGISTER QuickStart.b: /tmp/oliver/QuickStart/b
COPY /home/<AdminUserName>/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/c
  TO /tmp/oliver/QuickStart/c
REGISTER QuickStart.c: /tmp/oliver/QuickStart/c
```

Note that you can also import individual suites, and that registration groups do not need to be created explicitly - they are automatically created and deleted as required.

Now switch `gcylc` back to your private database and confirm that you have a copy of the example suites; on the command line:

```
% cylc db pr -t Quick
QuickStart
|-a      Quick Start Example A | /tmp/oliver/QuickStart/a
|-b      Quick Start Example B | /tmp/oliver/QuickStart/b
|-c      Quick Start Example C | /tmp/oliver/QuickStart/c
`-z      Quick Start Example Z | /tmp/oliver/QuickStart/z
```

6.4 View The QuickStart.a Suite Definition

Cylc suites are defined by *suite.rc files*, discussed at length in *Suite Definition* (Section 8) and the *Suite.rc Reference* (Appendix A). To view the `Quickstart.a` suite definition right-click on the suite name and choose ‘Edit’; or use the edit command:

```
% cylc edit QuickStart.a
```

This opens the suite definition in your editor (`$EDITOR`, or `$GEDITOR` from `gcylc`) *from the suite definition directory so that you can easily open other suite files in the editor*. You can of course do this manually, but by using the `cylc` interface you don’t have to remember suite definition directory locations. For the rare occasions that you do need to move to a suite definition directory, you can do this:

```
% cd $( cylc db get-dir QuickStart.a )
```

Suites that use include-files can optionally be edited in a temporarily inlined state - the inlined file gets split back into its constituent include-files when you save it and exit the editor. While editing, the inlined file becomes the official suite definition so that changes take effect whenever you save the file.

Anyhow, you should now see the following `suite.rc` file in your editor:

```

title = "Quick Start Example A"
description = "(see the CycL User Guide)"

[scheduling]
runahead limit = 12
initial cycle time = 2011010106
final cycle time = 2011010200
[[special tasks]]
    start-up      = Prep
    clock-triggered = GetData(1)
[[dependencies]]
    [[[0,6,12,18]]]
        graph = """Prep => GetData => Model => PostA
                    Model[T-6] => Model"""
    [[[6,18]]]
        graph = "Model => PostB"

[visualization] # optional
[[node groups]]
    post = PostA, PostB
[[node attributes]]
    post   = "style=unfilled", "color=blue", "shape=rectangle"
    PostB = "style=filled", "fillcolor=seagreen2"
    Model  = "style=filled", "fillcolor=red"
    GetData = "style=filled", "fillcolor=yellow3", "shape=septagon"
    Prep   = "shape=box", "style=bold", "color=red3"

```

(CycL comes with syntax highlighting and section folding for the *vim* editor - see Section 8.2.3).

This defines a complete, valid, runnable suite. Here's how to interpret it: At 0, 6, 12, and 18 hours each day a clock-triggered task called GetData triggers 1 hour after the wall clock reaches its (GetData's) nominal cycle time; then a task called Model triggers when GetData finishes; and a task called PostA triggers when Model is finished. Additionally, Model depends on its own previous instance from 6 hours earlier; and twice per day at 6 and 18 hours another task called PostB also triggers off Model.

All the tasks in this suite can run in parallel with their own previous instances if the opportunity arises (i.e. if their prerequisites are satisfied before the previous instance is finished). Most tasks should be capable of this (see Section 12.4) but if necessary you can force particular tasks to run sequentially like this:

```

# SUITE.RC
[scheduling]
[[special tasks]]
    sequential = GetData, PostB

```

Finally, when the suite is *cold-started* (started from scratch) it is made to wait on a special *synchronous start-up task* called Prep. Start-up tasks are one-off (non-spawning) tasks that are only used at suite start-up, and any dependence on them only applies at suite start-up. Start-up tasks are *synchronous* because they have a defined cycle time even though they are not cycling tasks. CycL also has *asynchronous one-off tasks*, which have no cycle time:

```

# SUITE.RC
[scheduling]
[[dependencies]]
    graph = "prep"      # an asynchronous one-off task (no cycle time)
    [[[ 0,6,12,18 ]]]
        graph = "prep => foo => bar"  # followed by cycling tasks

```

The optional visualization section configures graph plotting.

6.5 Plotting The QuickStart.a Dependency Graph

Right-click on the *QuickStart.a* suite in gcyclc and choose Graph; or by command line,

```
% cycl graph QuickStart.a 2011052300 2011052318 &
```

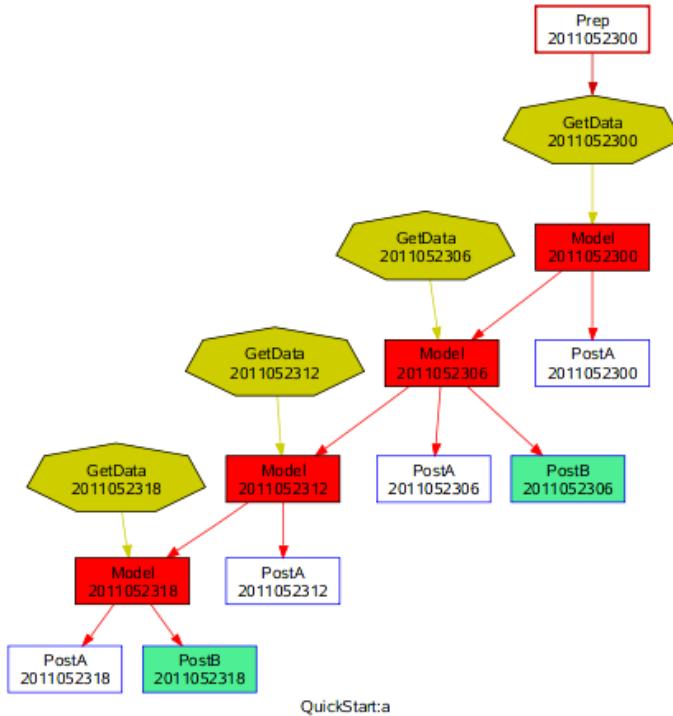


Figure 21: The *QuickStart.a* dependency graph, plotted by cylc.

This will pop up a zoomable, pannable, graph viewer showing the graph of Figure 21. If you edit the suite.rc file the viewer will update in real time whenever you save the file.

6.6 Run The QuickStart.a Suite

Each cylc task defines command scripting to invoke the right external processing when the task is ready to run. This has not been explicitly configured in the example suite, so it defaults, for all tasks, to the *dummy task* scripting inherited from the *root namespace* (see Section 8):

```
% cylc get-config QuickStart.a runtime GetData 'command scripting'
['echo DUMMY $CYLC_TASK_ID; sleep 10; echo BYE']
```

where `$CYLC_TASK_ID` is the unique identifier by which cylc knows the task (`NAME%CYCLE`) exported to each task's execution environment by cylc. The command arguments reflect the suite definition section nesting. If it were explicitly configured in the example suite it would look like this:

```
[runtime]
  [[GetData]]
    command scripting = "echo DUMMY $CYLC_TASK_ID; sleep 10; echo BYE"
```

Now start a suite control GUI by right-clicking on the suite in gcylc and choosing ‘Control (graph)’. You can also open a text treeview control GUI for the same suite, if you like. Multiple GUIs running at the same time will automatically connect to the same running suite (they won't try to run separate instances). Note also that if you shut down a suite control GUI, the suite will keep running. You can reconnect to it later by opening another control GUI.

In the control GUI click on Control → Run, enter an initial cold-start cycle time (e.g. 2011052306), and select “Hold (pause) on start-up” so that the suite will start in the held state (tasks will not be submitted even if they are ready to run).

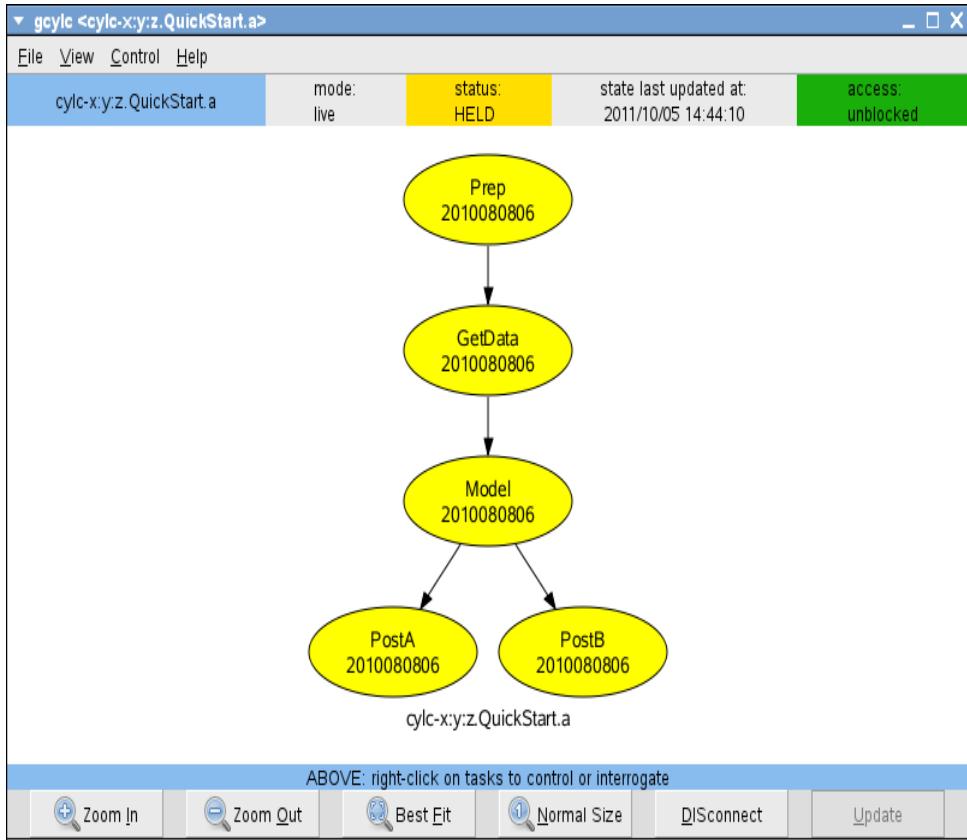


Figure 22: Suite *QuickStart:one* at start-up with an initial cycle time ending in 06 or 18 hours. Yellow nodes represent waiting tasks in the held state.

Do not choose an initial cycle time in the future unless you’re running in simulation mode, or nothing much will happen until that time.

If the initial cycle time ends in 06 or 18 the suite controller should look like Figure 22, or otherwise (00 or 12) like Figure 23.

The reason for the difference in graph structure between the two figures is this: cylc starts up with every task present in the waiting state (blue) at the initial cycle time *or* at the first subsequent valid cycle time for the task - and PostB does not run at 00 or 12. The off-white tasks are from the base graph, defined in the suite.rc file, and aren’t actually present in the suite as yet (they are shown in the graph in order to put the live tasks in context).

Now, click on Control → Release in the suite control GUI to *release the hold on the suite*, and observe what happens: the GetData tasks will rapidly go off in parallel out to a few cycles ahead (how far ahead depends on the suite runahead limit as explained below and in *The Suite Runahead Limit*, Section 11.2.1) and then the suite will stall, as shown in Figures 24 and 25.

The Prep task runs immediately because it has no prerequisites and is not clock-triggered. The clock-triggered GetData tasks then all go off at once because they have no prerequisites (i.e. they do not have to wait on any upstream tasks), their trigger time has long passed (the initial cycle time was in the past), and they are not sequential tasks (so they are able to run in parallel - try declaring GetData sequential to see the difference). Beyond the suite *runahead limit* of 12 hours (set in the suite.rc file), however, GetData is put into a special ‘runahead’ held state indicated by the darker blue graph node. The task will be released from this state once the slower tasks in the suite have caught up sufficiently. The runahead limit is designed

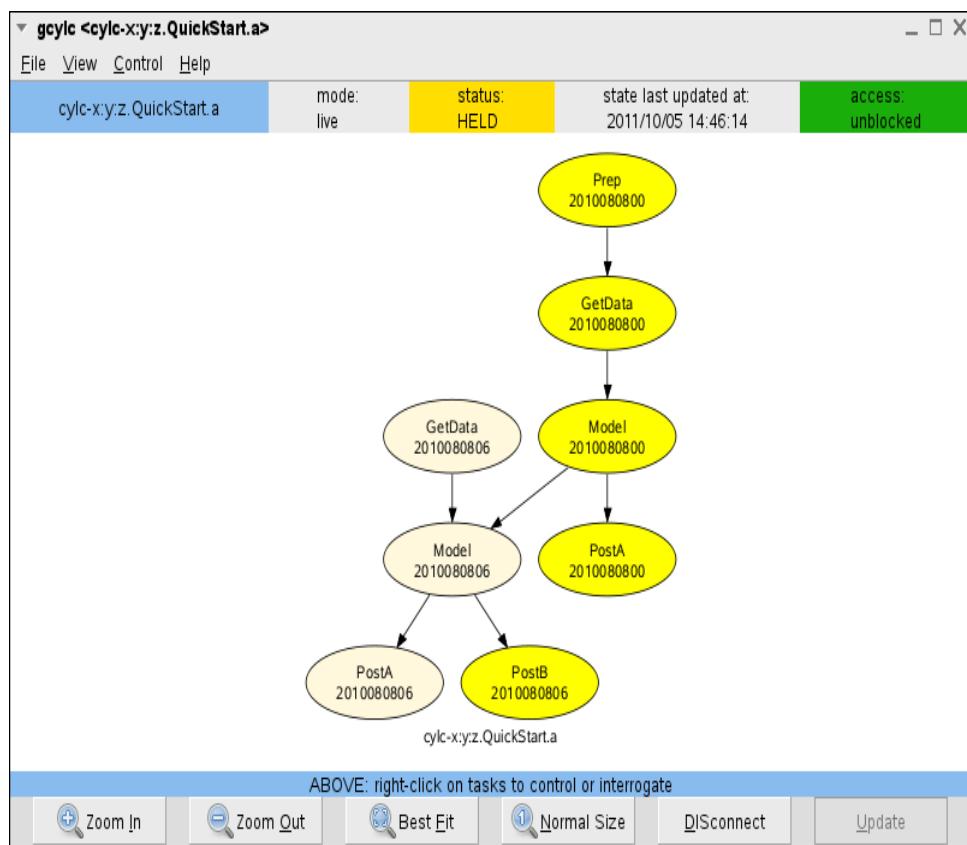


Figure 23: Suite *QuickStart.a* at start-up with an initial cycle time ending in 00 or 12 hours. Yellow nodes represent waiting tasks in the held state and off-white nodes are tasks from the base graph, defined in the suite.rc file, that aren't currently live in the suite.

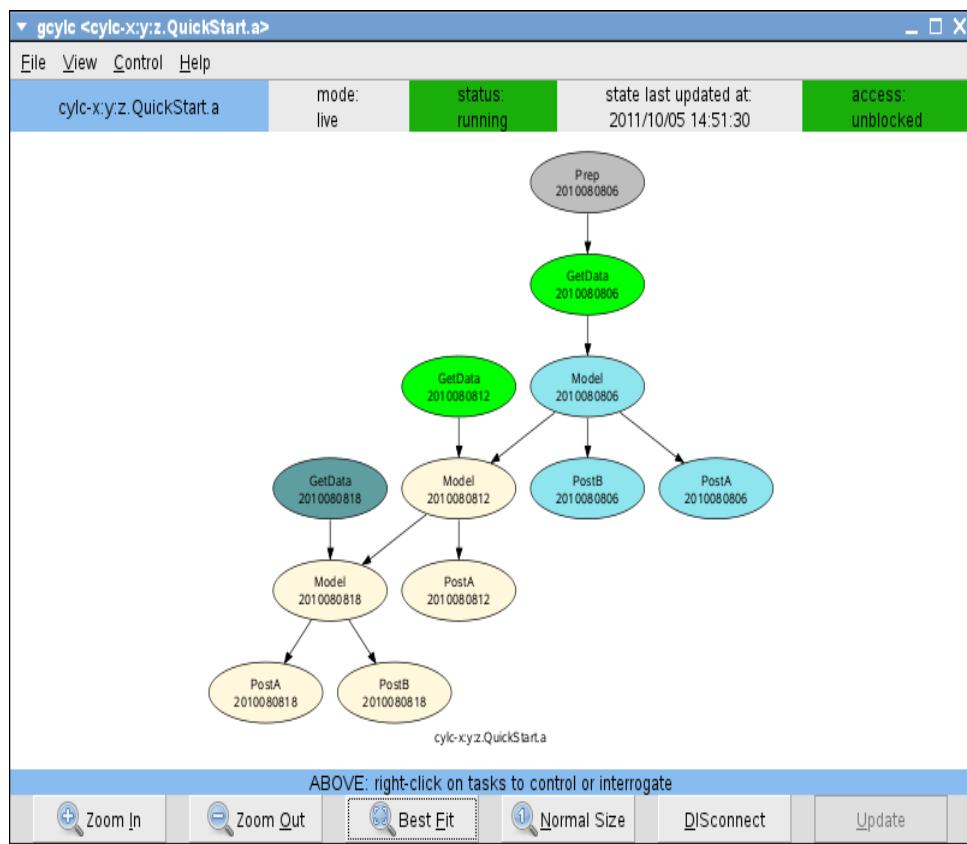


Figure 24: Suite *QuickStart.a* running, showing several consecutive instances of the clock-triggered *GetData* task running at once, out to the suite runahead limit of 12 hours.

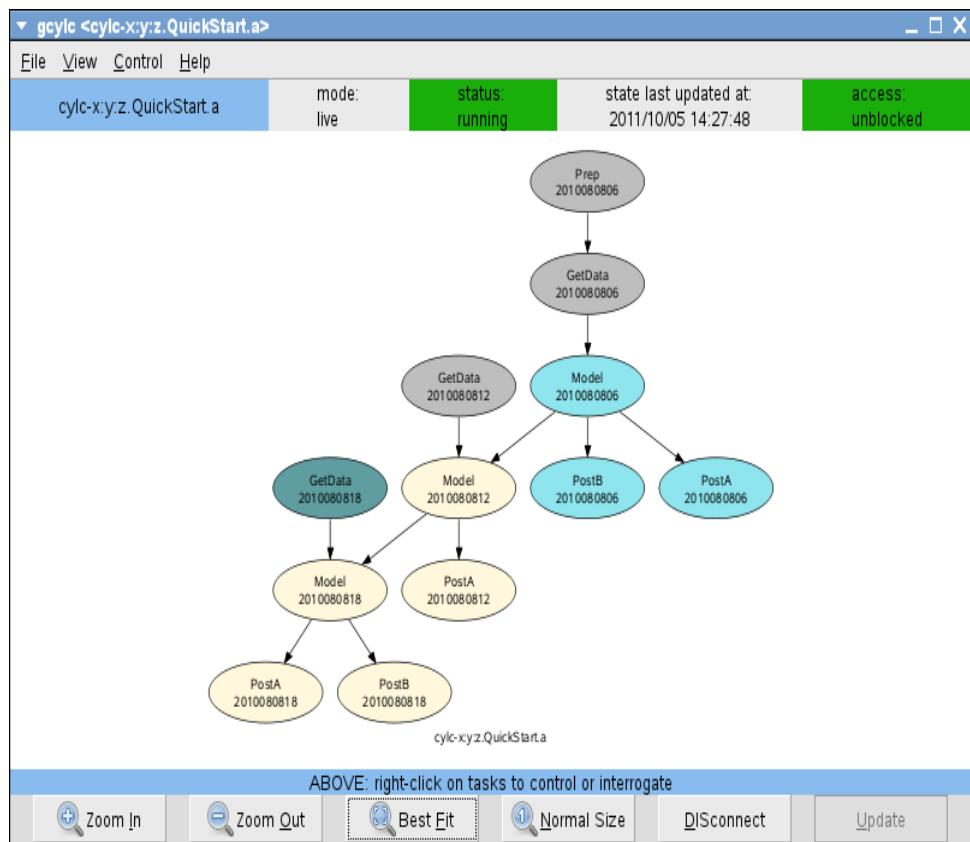


Figure 25: Suite *QuickStart.a* stalled after the clock-triggered GetData tasks have finished, because of Model's previous-cycle dependence and the suite runahead limit (see main text).

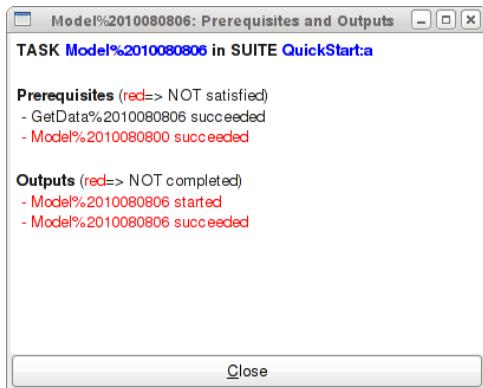


Figure 26: Viewing current task state after right-clicking on a task in cylc. The same information is available from the `cylc show` command.

to stop free tasks like this from running off too far into the future in delayed operation. It is of little consequence in real time operation⁸ because clock triggered tasks are then constrained by the wall clock, and other tasks have to wait on them, generally speaking. See Section 11.2.1 for more on this.

6.6.1 Viewing The State Of Tasks

If you're wondering why a particular task has not triggered yet in a running suite you can view the current state of its prerequisites by right-clicking on the task and choosing ‘View State’, or using `cylc show`. Do this for the first Model task, which appears to be stuck in the waiting state; it will pop up a small window as in Figure 26.

It is clear that the reason the task is not running, and consequently, by virtue of the runahead limit, why the suite has stalled, is that Model[T] is waiting on Model[T-6] which does not exist at suite start-up. Model represents a warm-cycled forecast model that depends on a model background state or restart file(s) generated by its own previous run.

6.6.2 Triggering Tasks Manually

Right-click on the waiting Model task and choose Trigger, or use `cylc trigger`, to force the task to trigger, and thereby get the suite up and running. In a real suite this would not be sufficient: the real forecast model that Model represents would fail for lack of the real restart files that it requires as input. Well see how to handle this properly shortly.

6.6.3 Suite Shut-Down And Restart

Having watched the *QuickStart.a* suite run for a while, choose Stop from the Control menu, or `cylc stop`, to shut it down. The default stop method waits for any tasks that are currently running to finish before shutting the suite down, so that the final recorded suite state is perfectly consistent with what actually happened.

You can restart the suite from where it left off by choosing Control → Run and selecting the ‘restart’ option, or using `cylc restart`. Note that cylc always writes a special state dump, and logs its name, prior to actioning any intervention, and you can also restart a suite from one of these states, rather than the default most recent state.

⁸So long as the runahead limit is sufficient to span the normal range of cycle times present in the suite - task that only run once per day, for example, have to spawn a successor that is 24 hours ahead.

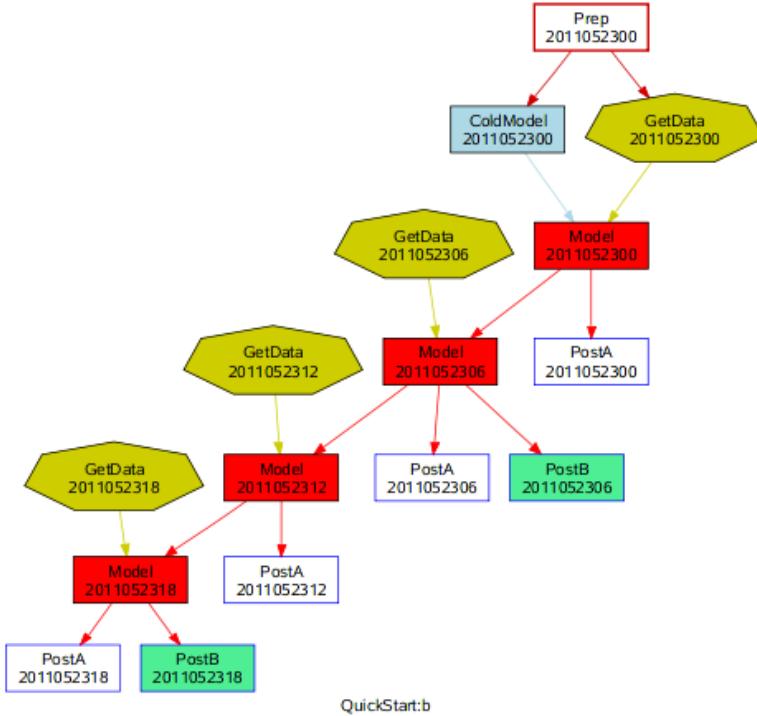


Figure 27: The *QuickStart.b* dependency graph showing a model cold start task.

6.7 QuickStart.b - Handling Cold-Starts Properly

Now take a look at *QuickStart.b*, which is a minor modification of *QuickStart.a*. Its suite.rc file has a new *cold-start* task called *ColdModel*,

```
# SUITE.RC
[scheduling]
[[special tasks]]
cold-start = ColdModel
```

and the dependency graph (see also Figure 27) looks like this:

```
# SUITE.RC
[scheduling]
[[dependencies]]
[[[ 0, 6, 12, 18 ]]]
graph = """Prep => GetData & ColdModel
          GetData => Model => PostA
          ColdModel | Model[T-6] => Model"""
[[[ 6, 18 ]]]
graph = "Model => PostB"
```

In other words, *Model*[T] can trigger off *either* *Model*[T-6] *or* *ColdModel*[T].

A cold-start task is a one-off task used to satisfy the previous-cycle dependence of a cotemporal task whose previous-cycle trigger is not available. The obvious use for this is to cold-start warm-cycled forecast models at suite start-up, when there is no previous cycle. Unlike start-up tasks though, cold-start dependencies are not restricted to suite start-up because it is sometimes useful to be able to insert a cold-start task into a running suite, to get a model restarted after it had to skip one or more cycles due to problems, without having to restart the whole suite.

A model cold-start task in a real suite may submit a real “cold start forecast” to generate the previous-cycle input files required by its associated model, or it may just stand in for some

external spinup process, or similar, that has to be completed before the suite is started (in the latter case the cold-start task would be a dummy task that just reports successful completion in order to satisfy the initial previous-cycle dependence of the model).

Run *QuickStart.b* to confirm that no manual triggering is required to get the suite started now.

6.8 QuickStart.c - Real Task Implementations

The suite *QuickStart.c* is the same as *QuickStart.b* except that it has real task implementations (scripts located in the suite bin directory) that generate and consume files in such a way that they have to run according to the graph of Figure 27. The suite gets them to run together out of a common I/O workspace, configured via the suite.rc file.

By studying this suite and its tasks, and by making quick copies of it to modify and run, you should be able to learn a lot about how to build real cylc suites. Here's the complete suite definition

```
title = "Quick Start Example C"
description = "(Quick Start b plus real tasks)"

# A clock-triggered data-gathering task, a warm-cycled model, and two
# post-processing tasks (one runs every second cycle). The tasks are not
# cylc-aware, have independently configured I/O directories, and abort
# if their input files do not exist. This suite gets them all to run out
# of a common I/O workspace (although the warm-cycled model uses a
# private running directory for its restart files).

[scheduling]
runahead limit = 12
initial cycle time = 2011010106
final cycle time = 2011010200
[[special tasks]]
    start-up      = Prep
    cold-start    = ColdModel
    clock-triggered = GetData(1)
[[dependencies]]
    [[[0,6,12,18]]]
        graph = """Prep => GetData & ColdModel
                  GetData => Model => PostA
                  ColdModel | Model[T-6] => Model"""
    [[[6,18]]]
        graph = "Model => PostB"

[runtime]
[[root]]
    [[[environment]]]
        TASK_EXE_SECONDS = 5
        WORKSPACE = /tmp/$USER/$CYLC_SUITE_REG_NAME/common

    [[Prep]]
        description = "prepare the suite workspace for a new run"
        command scripting = clean-workspace.sh $WORKSPACE

    [[GetData]]
        description = "retrieve data for the current cycle time"
        command scripting = GetData.sh
    [[[environment]]]
        GETDATA_OUTPUT_DIR = $WORKSPACE

    [[Models]]
    [[[environment]]]
        MODEL_INPUT_DIR = $WORKSPACE
        MODEL_OUTPUT_DIR = $WORKSPACE
        MODEL_RUNNING_DIR = $WORKSPACE/Model
    [[ColdModel]]
        inherit = Models
```

```

description = "cold start the forecast model"
command scripting = Model.sh --coldstart
[[Model]]
inherit = Models
description = "the forecast model"
command scripting = Model.sh

[[Post]]
description = "post processing for model"
[[[environment]]]
INPUT_DIR = $WORKSPACE
OUTPUT_DIR = $WORKSPACE
[[PostA,PostB]]
inherit = Post
command scripting = <TASK>.sh

[visualization]
default node attributes = "shape=ellipse"
[[node groups]]
post = PostA, PostB
models = ColdModel, Model
[[node attributes]]
post = "style=unfilled", "color=blue", "shape=rectangle"
PostB = "style=filled", "fillcolor=seagreen2"
models = "style=filled", "fillcolor=red"
ColdModel = "fillcolor=lightblue"
GetData = "style=filled", "fillcolor=yellow", "shape=septagon"
Prep = "shape=box", "style=bold", "color=red3"

```

Here's the namespace hierarchy defined by this suite:

```
% cyc list --tree QuickStart.c
root
|-GetData      retrieve data for the current cycle time
|-Models
| |-ColdModel cold start the forecast model
| '-Model      the forecast model
|-Post
| |-PostA     post processing for model
| '-PostB     post processing for model
`-Prep        prepare the suite workspace for a new run
```

And here, for example, is the complete implementation for the PostA task (located with the other task scripts in the suite bin directory):

```
#!/bin/bash

set -e

cyc checkvars TASK_EXE_SECONDS
cyc checkvars -d INPUT_DIR
cyc checkvars -c OUTPUT_DIR

# CHECK INPUT FILES EXIST
PRE=${INPUT_DIR}/surface-winds-${CYLC_TASK_CYCLE_TIME}.nc
if [[ ! -f $PRE ]]; then
    echo "ERROR, file not found $PRE" >&2
    exit 1
fi

echo "Hello from ${CYLC_TASK_NAME} at ${CYLC_TASK_CYCLE_TIME} in ${CYLC_SUITE_REG_NAME}"

sleep $TASK_EXE_SECONDS

# generate outputs
touch ${OUTPUT_DIR}/surface-wind.products
```

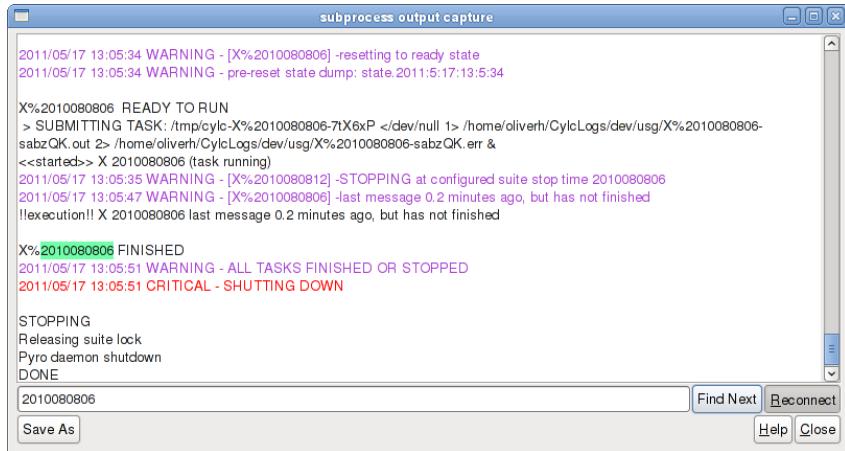


Figure 28: Cylc suite stdout/stderr example.

6.9 Monitoring Running Suites

6.9.1 Suite stdout And stderr

When cylc runs a suite it writes warnings and other informative messages, such as when and how each task is submitted, to the stdout stream. If you start a suite at the command line you can direct this output wherever you like. If you start a suite via gcylc, however, the output is directed to a special file, `$HOME/.cylc/[SUITE].out` that can be accessed again later if you reconnect to the suite with a new control GUI.

6.9.2 Suite Logs

The cylc suite log records every event that occurs (incoming messages from tasks and so on) along with the time of the event. The top level logging directory, under which a suite-specific log is written, is configurable in the suite.rc file. Suite logs are (optionally) rolled over at start-up and the five most recent back ups are kept.

Figure 29 shows a suite log viewed from within gcylc. The `cylc log` command also enables viewing and filtering of suite logs without having to remember the actual log file location. But if you want to know the location:

```
% cylc get-config --directories QuickStart.a
SUITE LOG DIRECTORY:
  /home/oliverh/cylc-run/QuickStart.a/log/suite
SUITE STATE DUMP DIRECTORY:
  /home/oliverh/cylc-run/QuickStart.a/state
JOB SUBMISSION LOG DIRECTORIES:
  + root: /home/oliverh/cylc-run/QuickStart.a/log/job
```

(There would be other job submission log directories listed here if any of the tasks in the suite had overridden the default location set by the root namespace).

6.9.3 Task stdout And stderr Logs

The stdout and stderr logs generated when a task is submitted end up in the suite.rc *job submission log directory*, default location `$HOME/cylc-run/$CYLC_SUITE_REG_NAME/log/job/` (where the variable `$CYLC_SUITE_REG_NAME` evaluates to the registered suite name hierarchy, e.g. `foo.bar.baz`):

```
% get-config QuickStart.a runtime GetData 'job submission' 'log directory'
/home/oliverh/cylc-run/QuickStart.a/log/job
```

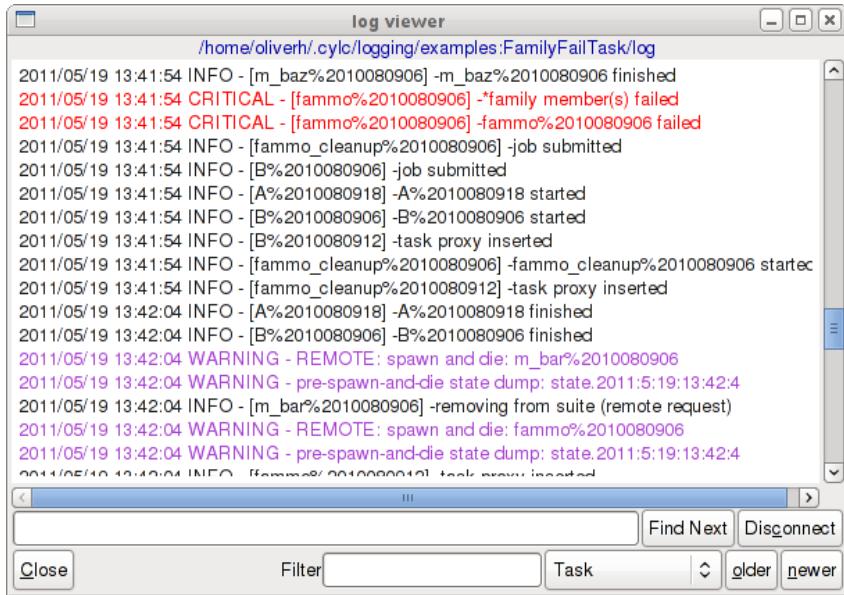


Figure 29: A cylc suite log viewed via gcylc.

(or use the `--directories` argument as shown just above).

These files will contain the complete stdout and stderr record for tasks whose initiating processes do not detach and exit before all task processing is finished. The location of output generated by secondary processes, if the original process detaches and exits early, is up to the task implementation (of said secondary processes, specifically).

6.10 Searching A Suite

The cylc suite search tool reports matches in the suite.rc file by line number, suite section, and file, even if include-files are used (and even if they are nested), and by file and line number for matches in the suite bin directory. The following output listing is from a search of the *QuickStart.c* suite.

```
% cylc grep OUTPUT_DIR QuickStart.c

SUITE: QuickStart.c /tmp/oliverh/QuickStart/c/suite.rc
PATTERN: OUTPUT_DIR

FILE: /tmp/oliverh/QuickStart/c/suite.rc
SECTION: [runtime]->[[GetData]]->[[[environment]]]
(40):           GETDATA_OUTPUT_DIR = $WORKSPACE
SECTION: [runtime]->[[Models]]->[[[environment]]]
(45):           MODEL_OUTPUT_DIR = $WORKSPACE
SECTION: [runtime]->[[Post]]->[[[environment]]]
(60):           OUTPUT_DIR = $WORKSPACE

FILE: /tmp/oliverh/QuickStart/c/bin/PostB.sh
(7): cylc checkvars -c OUTPUT_DIR
(21): touch $OUTPUT_DIR/precip.products

FILE: /tmp/oliverh/QuickStart/c/bin/Model.sh
(11): cylc checkvars -c MODEL_OUTPUT_DIR MODEL_RUNNING_DIR
(54): touch $MODEL_OUTPUT_DIR/surface-winds-${CYLC_TASK_CYCLE_TIME}.nc
(55): touch $MODEL_OUTPUT_DIR/precipitation-${CYLC_TASK_CYCLE_TIME}.nc

FILE: /tmp/oliverh/QuickStart/c/bin/PostA.sh
(7): cylc checkvars -c OUTPUT_DIR
```

```
(21): touch $OUTPUT_DIR/surface-wind.products

FILE: /tmp/oliverh/QuickStart/c/bin/GetData.sh
(6): cylc checkvars -c GETDATA_OUTPUT_DIR
(11): touch ${GETDATA_OUTPUT_DIR}/obs-$CYLC_TASK_CYCLE_TIME.nc
```

(Suite search is also available from the gcylc right-click menu).

6.11 Comparing Suites

Cylc can also compare suites and report differences by suite.rc section and item. For instance, comparing the example suites `QuickStart.a` and `QuickStart.b` by GUI or command line results in:

```
% cylc diff QuickStart.a QuickStart.b

Parsing QuickStart.a
Parsing QuickStart.b
Suite definitions QuickStart.a and QuickStart.b differ.
!SNIP!
13 common items differ QuickStart.a(<) QuickStart.b(>)
  (top)
<   description = (see the Cylc User Guide)
>   description = (Quick Start a plus a cold-start task)
<   title = Quick Start Example A
>   title = Quick Start Example B

  [cylc] [[logging]]
<   directory = /home/oliverh/cylc-run/QuickStart.a/log/suite
>   directory = /home/oliverh/cylc-run/QuickStart.b/log/suite

  [scheduling] [[dependencies]] [[[0,6,12,18]]]
<   graph = Prep => GetData => Model => PostA
          Model(T-6) => Model
>   graph = Prep => GetData & ColdModel
          GetData => Model => PostA
          ColdModel | Model(T-6) => Model
!SNIP!
```

(much of the diff output has been omitted here for the sake of brevity). Note that suite log directories and the like may differ even though they are not explicitly configured in either suite, because their default values (see the *Suite.rc Reference*, Appendix A) are suite-registration-specific.

6.12 Validating A Suite

Suite validation checks for errors by parsing the suite definition, comparing all items against the suite.rc specification file, and then parsing the suite graph and attempting to instantiate all task proxy objects. This can be done using gcylc or `cylc validate`:

```
% cylc validate -v foo.bar
Parsing Suite Definition
LOADING suite.rc
VALIDATING against the suite.rc specification.
PARSING clock-triggered tasks
PARSING runtime generator expressions
PARSING runtime hierarchies
PARSING SUITE GRAPH
Instantiating Task Proxies:
root
|-GEN
| |-OPS
| | |-aircraft ... OK
| | |-atovs ... OK
| | `-'atovs_post ... OK
```

```

| `--VAR
|   |-AnPF      ... OK
|   `--ConLS    ... OK
|-baz
| |-bar1       ... OK
| `--bar2      ... OK
|-foo          ... OK
`--prepobs    ... OK
Suite foo.bar is valid for cylc-4.2.0

```

For more information on suite validation see Section [8.2.5](#).

6.13 Other Example Suites

Cylc has been designed from the ground up to make prototyping and testing new suites very easy. Simply drawing (in text) a dependency graph in a new suite definition creates a valid suite that you can run: the tasks will be *dummy tasks* that default to emitting an identifying message, sleeping for a few seconds, and then exiting; but you can then arrange for particular tasks to do more complex things by configuring their runtime properties appropriately.

Cylc has example suites intended to illustrate most facets of suite construction. These are held centrally under `$CYLC_DIR/examples` and are exported to the central suite database, accessible to all users, when cylc is installed. They all run “out the box” and can be copied and modified by users to test almost anything. Some of them just configure a suite dependency graph, in which case cylc will run dummy tasks according to the graph; some also configure task runtime properties (e.g. command scripting and environment variables) within the suite definition; and some have real task implementations that generate and consume real files and which will fail if they are not executed in the right order. All of the example suites are portable in the sense that all suite and task I/O directory paths incorporate the suite registration name (this is the default situation for any suite in fact), so you can run multiple copies of the same suite at once without any interference between them.

Use `gcylc` or `cylc db print --central <AdminUserName>` to see the example suites that are currently available. Now use `cylc import` to copy an example suite (or some subset of them, or all of them at once) to your private database so that you can modify it and run it:

```
% cylc db print --central --tree <AdminUserName>.*\QuickStart
# or just this:
% cylc db print --central --tree QuickStart
<AdminUserName>
`--cylc-x-y-z
`--QuickStart
  |-a  Quick Start Example A | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/a
  |-b  Quick Start Example B | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/b
  |-c  Quick Start Example C | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/c
  `--z  Quick Start Example Z | ~/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/z
% cylc db import <AdminUserName>.cylc-x-y-z.QuickStart.z eg.z $TMPDIR
COPY /home/<AdminUserName>/cylc/CDB/<AdminUserName>/cylc-x-y-z/QuickStart/z
  TO /tmp/oliver/eg/z
REGISTER eg.z: /tmp/oliver/eg/z

% cylc db print eg
eg.z | Quick Start Example Z | /tmp/oliver/eg/z

% cylc prep edit eg.z
```

```
title = "Quick Start Example Z"
description = "(Example A without the visualization config)"

[scheduling]
  runahead limit = 12
  initial cycle time = 2011010106
  final cycle time = 2011010200
```

7 SUITE REGISTRATION

```
[[special tasks]]
  start-up      = Prep
  clock-triggered = GetData(1)
[[dependencies]]
  [[[0,6,12,18]]]
    graph = """Prep => GetData => Model => PostA
              Model[T-6] => Model"""
  [[[6,18]]]
    graph = "Model => PostB"
```

(This suite is explained in the *Quick Start Guide*, Section 6).

6.13.1 Choosing The Initial Cycle Time

When running any suite in live mode (or suites that depend on clock-triggered tasks at least) do not give an initial cycle time in the future unless you want nothing to happen until that time. However, you can also run any suite in *simulation mode* in which case a future start time is fine (see Appendix E).

7 Suite Registration

Cyc commands target particular suites via names registered in a *suite database*, so that you don't need to remember and continually refer to the actual location of the suite definition on disk. A suite registration name is a hierarchical name akin to a directory path but delimited by the '.' character; this allows suites to be organised in nested tree-like structures:

```
% cylc db print -tp nwp
nwp
|-oper
| |-region1 Local Model Region1      /oper/nwp/suite_defs/LocalModel/nested/Region1
| |-region2 Local Model Region2      /oper/nwp/suite_defs/LocalModel/nested/Region2
`-test
  '-region1 Local Model TEST Region1 /home/oliverh/nwp_suites/Regional/TESTING/Region1
```

Note that registration groups are entirely virtual, they do not need to be explicitly created before use, and they automatically disappear if all tasks are removed from them. From the listing above, for example, to move the suite `nwp.oper.region2` into the `nwp.test` group:

```
% cylc db rereg nwp.oper.region2 nwp.test.region2
REREREGISTER nwp.oper.region2 to nwp.test.region2
% cylc db print -tpx nwp
nwp
|-oper
| |-region1 Local Model Region1
`-test
  |-region1 Local Model TEST Region1
  '-region2 Local Model Region2
```

And to move `nwp.test.region2` into a new group `nwp.para`:

```
% cylc db rereg nwp.test.region2 nwp.para.region2
REREREGISTER nwp.test.region2 to nwp.para.region2
% cylc db print -tpx nwp
nwp
|-oper
| |-region1 Local Model Region1
`-test
  |-region1 Local Model TEST Region1
`-para
  '-region2 Local Model Region2
```

Currently you cannot explicitly indicate a group name on the command line by appending a dot character. Rather, in database operations such as copy, reregister, or unregister, the identity of the source item (group or suite) is inferred from the content of the database; and if

the source item is a group, so must the target be a group (or it will be, in the case of an item that will be created by the operation). This means that you cannot copy a single suite into a group that does not exist yet unless you specify the entire target registration name.

`cylc db register --help` shows a number of other examples.

7.1 Private Suite Database

Each cylc user has a private suite database for working with his/her own suites. By right-clicking on a suite in your private database (or using cylc commands) you can:

1. start a suite control GUI to run the suite (or connect to a running suite),
2. submit a single task to run, just as it would be submitted by its suite
3. view the suite stdout and stderr streams,
4. view the suite log (which records all events and messages from tasks),
5. retrieve the suite description,
6. list tasks in the suite,
7. view the suite namespace hierarchy,
8. edit the suite definition in your editor,
9. plot the suite dependency graph,
10. search the suite definition and bin scripts,
11. validate the suite definition,
12. copy the suite or group,
13. alias the suite name to another name,
14. compare (difference) the suite with another suite,
15. export the suite or group to the central database,
16. unreregister the suite or group,
17. reregister the suite or group.

Note that the suite title shown in gcylc is parsed from the suite.rc file at the time of initial registration; if you change the title (by editing the suite.rc file) use `cylc db refresh` or gcylc View → Refresh to update the database.

7.1.1 Private Database Location

Your private suite database, currently, will be `$HOME/.cylc/LDB/db`.

7.2 Central Suite Database

The central suite database facilitates sharing of suites among cylc users. It works similarly to a private suite database except that central suite registration names always begin with the suite owner's username. Users can export suites to the central database to make them available to others.

Suites in the central database can be inspected by various means, but you can't run them without *importing* them to your private database.

By right-clicking on a suite in the central database, or using the cylc command line, you can:

1. retrieve the suite description,
2. list tasks in the suite,
3. view the suite namespace hierarchy,

4. view the suite definition in your editor,
5. plot the suite dependency graph,
6. search the suite definition and bin scripts,
7. validate the suite definition,
8. compare (difference) the suite with another suite,
9. import the suite or group to your private database,
10. unreregister or delete the suite or group (if you own it),
11. reregister the suite or group (if you own it).

7.2.1 Central Database Location

The central suite database is currently located here: `$CYLC_DIR/CDB/db`, and is accessed through the filesystem on the cylc host machine. Consequently it has to be world, or at least group, writeable, which is not very secure.

We intend to develop a client/server interface to the central suite database so that users can easily share suites across the network. In the meantime, “importing” a suite manually from another user on another host is simply a matter of copying the suite definition directory over, registering the local copy in your private database, and then exporting it to the local central database.

7.3 Database Operations

On the command line, the ‘database’ (or ‘db’) command category contains commands to implement the aforementioned operations.

```
CATEGORY: db|database - private and central suite registration.

HELP: cylc [db|database] COMMAND help,--help
      You can abbreviate db|database and COMMAND.
      The category db|database may be omitted.

COMMANDS:
  alias ..... Register an alias for another registered suite
  browse ..... Browse cylc documentation.
  copy|cp ..... Copy a suite or group of suites.
  export ..... Export private registrations to the central database
  g cylc ..... The cylc Graphical User Interface
  get-directory ..... Print a suite definition directory path
  import ..... Import central registrations to your private database
  print ..... Print private or central suite registrations
  refresh ..... Report invalid registrations and update suite titles
  register ..... Add a suite to your private registration database
  reregister|rename ... Change a suite's registered name.
  unregister ..... Remove suites from the private or central database
```

Groups of suites (at any level in the registration hierarchy) can be deleted, copied, imported, and exported; as well as individual suites. To do this, just use suite group names as source and/or target for operations, as appropriate. For instance, if a group `foo.bar` contains the suites `foo.bar.baz` and `foo.bar.waz`, you can copy a single suite like this:

```
% cylc copy foo.bar.baz boo $TMPDIR
```

(resulting in a new suite `boo`); or the group like this:

```
% cylc copy foo.bar boo $TMPDIR
```

(resulting in new suites `boo.baz` and `boo.waz`); or the group like this:

```
% cylc copy foo boo $TMPDIR
```

8 SUITE DEFINITION

(resulting in new suites `boo.bar.baz` and `boo.bar.waz`). When suites are copied, the suite definition directories are copied into a directory tree, under the target directory, that reflects the registration name hierarchy. `cylc copy --help` has some explicit examples.

The same functionality is also available by right-clicking on suites or groups in the cylc GUI, as shown in Figure 12.

8 Suite Definition

Cylc suites are defined in structured, validated, `suite.rc` files that concisely specify the properties of, and the relationships between, the various tasks managed by the suite. This section of the User Guide deals with the format and content of the `suite.rc` file, including task definition. Task implementation - what's required of the real commands, scripts, or programs that do the processing that the tasks represent - is covered in Section 9; and task job submission - how tasks are submitted to run - is in Section 10.

8.1 Suite Definition Directories

A cylc *suite definition directory* contains:

- **A `suite.rc` file:** this is the suite definition.
 - And any include-files used in it (see below; may be kept in sub-directories).
- **A `bin/` sub-directory.**
 - *Optional.*
 - For scripts and executables that implement, or are used by, suite tasks.
 - Automatically added to `$PATH` in task execution environments.
 - Alternatively, tasks can call external commands, scripts, or programs; or they can be scripted entirely within the `suite.rc` file.
- **A `python/` sub-directory.**
 - *Optional.*
 - For user-defined job-submission methods, if needed (see Section 10).
 - Alternatively, new job submission methods can be installed into the cylc source tree, or in any directory in your `$PYTHONPATH`.
- **Any other sub-directories and files** - documentation, control files, etc.
 - *Optional.*
 - Holding everything in one place makes proper suite revision control possible.
 - Portable access to files here, for running tasks, is provided through `$CYLC_SUITE_DEF_PATH` (see Section 8.4.4).
 - Ignored by cylc, but the entire suite definition directory tree is copied when you copy a suite using cylc commands.

A typical example:

```
/path/to/my/suite    # suite definition directory
suite.rc             # THE SUITE DEFINITION FILE
bin/                 # scripts and executables used by tasks
  foo.sh
  bar.sh
  ...
# (OPTIONAL) any other suite-related files, for example:
inc/                 # suite.rc include-files
```

```

nwp-tasks.rc
globals.rc
...
doc/           # documentation
control/       # control files
ancil/         # ancillary files
...

```

8.2 Suite.rc File Overview

Suite.rc files conform to the ConfigObj extended INI format (<http://www.voidspace.org.uk/python/configobj.html>) with several modifications to allow continuation lines and include-files, and to make it legal to redefine environment variables and scheduler directives (duplicate config item definitions are normally flagged as an error).

Additionally, embedded template processor expressions may be used in the file, to programmatically generate the final suite definition seen by cylc. Currently the Jinja2 template engine is supported (<http://jinja.pocoo.org/docs>). In the future cylc may provide a plug-in interface to allow use of other template engines too. See *Jinja2 Suite Templates* (Section 8.6) for some examples.

8.2.1 Syntax

The following list shows legal raw suite.rc syntax. Suites using the Jinja2 template processor (Section 8.6) can of course use Jinja2 syntax as well (it must generate raw syntax on processing).

- **Entries** are of the form `item = value`, and item names may contain spaces.
- **Comments** `#` follow a hash character.
- **List Values** are comma, separated.
- **Strings** must be quoted “if, they, contain, commas”
- **Multiline Strings** must be “““triple-quoted”””.
- **Boolean Values** are written as True or False (capitalized).
- **White Space** is ignored; indentation can be used for clarity.
- **Continuation Lines** follow a trailing backslash.\
- **[Section Headings]** are enclosed in square brackets.
- **[[Subsection Nesting]]** is indicated by the number of square brackets.⁹
- **Duplicate Items** are illegal, except in `environment` and `directives` sections.¹⁰
- **Include-files** can be used: `%include inc/foo.rc`.

The following pseudo-listing illustrates suite.rc syntax:

```

# a full line comment
an item = value # a trailing comment
a boolean item = True # or False
one string item = the quick brown fox # string quotes optional ...
two string item = "the quick, brown fox" # ... unless internal commas
a multiline string item = """the quick brown fox
jumped over the lazy dog"""\ # triple quoted
a list item = foo, bar, baz # comma separated
a list item with continuation = a, b, c, \
                                d, e, f

```

⁹Sections are closed by the next section heading, so items within a section must be defined before any subsequent subsection headings.

¹⁰The exceptions were designed to allow tasks to override environment variables defined in include-files that could be included in multiple tasks, to assist in factoring out common task configuration. However, namespace inheritance now provides a better way to do this in most cases.

```
[section]
  item = value
%include inc/vars/foo.inc # include file
  [[subsection]]
    item = value
    [[[subsubsection]]]
      item = value
[another section]
  [[another subsection]]
    #
  #
# ...
```

8.2.2 Include-Files

Cylc has native support for suite.rc include-files, which may help to organize large suites. Inclusion boundaries are completely arbitrary - you can think of include-files as chunks of the suite.rc file simply cut-and-pasted into another file. Include-files may be included multiple times in the same file, and even nested. Include-file paths can be specified portably relative to the suite definition directory, e.g.:

```
# SUITE.RC
# include the file $CYLC_SUITE_DEF_PATH/inc/foo.rc:
%include inc/foo.rc
```

8.2.2.1 Editing Temporarily Inlined Suites

Cylc's native file inclusion mechanism supports optional inlined editing:

```
% cylc edit --inline SUITE
```

The suite will be split back into its constituent include-files when you exit the edit session. While editing, the inlined file becomes the official suite definition so that changes take effect whenever you save the file. See `cylc prep edit --help` for more information.

8.2.2.2 Include-Files via Jinja2

Jinja2 (Section 8.6) provides template inclusion functionality although this is more akin to Python module import than simple text inclusion, and the implications of this for suite design have not yet been explored.

8.2.3 Syntax Highlighting In Vim

Cylc comes with a syntax file to configure suite.rc syntax highlighting and section folding in the *vim* editor, as shown in Figure 14. To use this, copy `$CYLC_DIR/conf/cylc.vim` to your `$HOME/.vim/syntax/` directory and make some minor modifications, as described in the syntax file, to your `$HOME/.vimrc` file.

8.2.4 Gross File Structure

Cylc suite.rc files consist of a suite title and description followed by configuration items grouped under several top level section headings:

- **[cylc]** - *non task-related suite configuration*
 - suite logging directories, simulation mode, UTC mode, etc.
- **[scheduling]** - *determines when tasks are ready to run*

- tasks with special behaviour, e.g. clock-triggered tasks
- the dependency graph, which defines the relationships between tasks
- [runtime] - determines how, where, and what to execute when tasks are ready
 - command scripting, environment, job submission, remote hosting, etc.
 - suite-wide defaults in the *root* namespace
 - a nested family hierarchy with common properties inherited by related tasks
- [visualization] - configures suite graphing and the graph suite control GUI.

8.2.5 Validation

Cylc suite.rc files are automatically validated against a specification file that defines all legal entries, values, options, and defaults (`$_CYLC_DIR/conf/suiterc.spec`). This detects any formatting errors, typographic errors, illegal items and illegal values prior to run time. Some values are complex strings that require further parsing by cylc to determine their correctness (this is also done during validation). All legal entries are documented in the *Suite.rc Reference* (Appendix A).

The validator reports the line numbers of detected errors. Here's an example showing a subsection heading with a missing right bracket.

```
% cylc validate foo.bar
Parsing Suite Config File
ERROR: [[special tasks]
NestingError('Cannot compute the section depth at line 19.',)
_validate foo.bar failed: 1
```

If the suite.rc file contains include-files you can use `cylc view` to view an inlined copy with correct line numbers (you can also edit suites in a temporarily inlined state with `cylc edit --inline`).

Validation does not check the validity of chosen job submission methods; this is to allow users to extend cylc with their own job submission methods, which are by definition unknown to the suiterc.spec file.

8.3 Scheduling - Dependency Graphs

The [scheduling] section of a suite.rc file defines the relationships between tasks in a suite - the information that allows cylc to determine when tasks are ready to run. The most important component of this is the suite dependency graph. Cylc graph notation makes clear textual graph representations that are very concise because sections of the graph that repeat at different hours of the day only have to be defined once. Here's an example showing a simple task tree with dependencies that vary depending on cycle time:

```
# SUITE.RC
[scheduling]
[[dependencies]]
  [[[0,6,12,18]]] # validity (hours of the day)
    graph = """
A => B & C  # B and C trigger off A
A[T-6] => A  # Model A restart trigger
"""
  [[[6,18]]]
    graph = "C => X"
```

Figure 30 shows the complete suite.rc listing alongside the suite graph, plotted with,

```
% cylc graph SDepG 2011052200 2011052206
# (or use right-click Graph in gcylc)
```

This is actually a complete, valid, runnable suite (it will use default runtime properties such as dummy command scripting, because no runtime properties are explicitly defined, and you'll

```
# SUITE.RC
title = "Dependency Graph Example"
[scheduling]
  [[dependencies]]
    [[[0,6,12,18]]] # validity (hours)
    graph = """
A => B & C  # B and C trigger off A
A[T-6] => A  # Model A restart trigger
"""
    [[[6,18]]] # hours
    graph = "C => X"
[visualization]
  [[node_attributes]]
    X = "color=red"
```

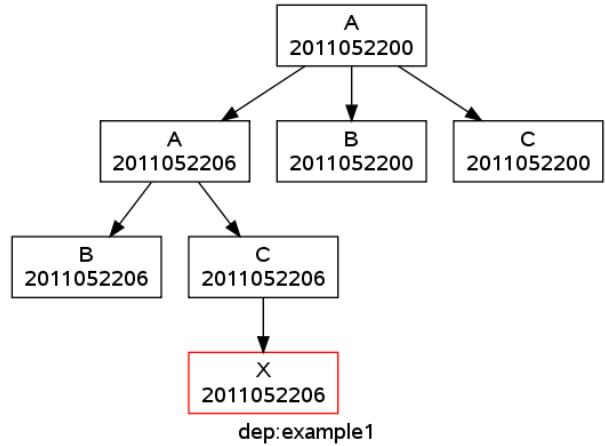


Figure 30: Example Suite SDepG

need to trigger task A manually to get the suite started because A[T] depends on A[T-6] and at start-up there is no previous cycle to satisfy that dependence - how to handle this properly is described in *Handling Intercycle Dependencies At Start-Up* (Section 8.3.5) and in the *Quick Start Guide* (Section 6).

8.3.1 Graph String Syntax

Multiline graph strings may contain:

- blank lines
- arbitrary white space
- task names
- internal comments: following the # character
- cycle time offsets: A[T-6]
- success triggers: foo => bar
- failure triggers: foo:fail => bar
- suicide triggers: foo => !bar
- explicit output triggers: foo:out1 => bar
- conditional triggers: (A | B) & C => D

8.3.2 Interpreting Graph Strings

Suite dependency graphs can be broken down into pairs in which the left side (which may be a single task or family, or several that are conditionally related) defines a trigger for the task or family on the right. For instance the “word graph” *C triggers off B which triggers off A* can be deconstructed into pairs *C triggers off B* and *B triggers off A*.

In the case of cycling tasks, the triggers defined by a graph string are valid for cycle times matching the list of hours specified for the graph section. For example this graph,

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[0,12]]]
    graph = "A => B"
```

implies that B triggers off A for cycle times in which the hour matches 0 or 12.

To define intercycle dependencies, attach an offset indicator to the left side of a pair:

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[0,12]]]
      graph = "A[T-12] => B"
```

This means $B[T]$ triggers off $A[T-12]$ for cycle times T with hours matching 0 or 12. Note that T must be left implicit unless there is a cycle time offset (this helps to keep graphs clean and concise because the majority of tasks in a typical suite will only depend on others with the same cycle time) and that cycle time offsets can only appear on the left (because each pair defines a trigger for the right task at cycle time T).

Now, having explained that dependency graphs are interpreted pairwise, you can optionally chain pairs together to “follow a path” through the graph. So this,

```
# SUITE.RC
graph = """A => B # B triggers off A
          B => C # C triggers off B"""
```

is equivalent to this:

```
# SUITE.RC
graph = "A => B => C"
```

Cycle time offsets, if they appear in a chain of triggers, must be leftmost (because, as explained previously they can't appear on the right of any pair). So this is legal:

```
# SUITE.RC
graph = "A[T-6] => B => C" # OK
```

but this isn't:

```
# SUITE.RC
graph = "A => B[T-6] => C" # ERROR!
```

The trigger $A \Rightarrow B[T-6]$ does not make sense in any case - if this kind of relationship seems necessary it probably means that B should be “reassigned” to the next cycle (keep in mind that cycle time is really just a label used to define the relationships between tasks).

Each trigger in the graph must be unique but the same task can appear in multiple pairs or chains. Separately defined triggers for the same task have an AND relationship. So this:

```
# SUITE.RC
graph = """A => X # X triggers off A
          B => X # X also triggers off B"""
```

is equivalent to this:

```
# SUITE.RC
graph = "A & B => X" # X triggers off A AND B
```

In summary, the branching tree structure of a dependency graph can be partitioned into lines (in the suite.rc graph string) of pairs or chains, in any way you like, with liberal use of internal white space and comments to make the graph structure as clear as possible.

```
# SUITE.RC
# B triggers if A succeeds, then C and D trigger if B succeeds:
graph = "A => B => C & D"
# which is equivalent to this:
graph = """A => B => C
          B => D"""
# and to this:
graph = """A => B => D
          B => C"""
# and to this:
graph = """A => B
          B => C
          B => D"""
# and it can even be written like this:
```

```
# SUITE.RC
title = some one-off asynchronous tasks
[scheduling]
[[dependencies]]
graph = "foo => bar & baz => waz"
```

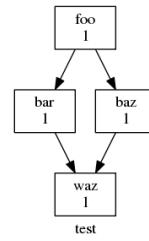


Figure 31: One-off Asynchronous Tasks.

```
graph = """A => B # blank line follows:
B => C # comment ...
B => D"""

```

8.3.2.1 Handling Long Graph Lines

Long chains of dependencies can be split into pairs:

```
# SUITE.RC
graph = "A => B => C"
# is equivalent to this:
graph = """A => B
          B => C"""
# BUT THIS IS AN ERROR:
graph = """A => B => # WRONG!
          C""""      # WRONG!
```

If you have very long task names, or long conditional trigger expressions (below) then you can use the suite.rc line continuation marker:

```
# SUITE.RC
graph = "A => B \
=> C" # OK
```

Note that a line continuation marker must be the final character on the line; they cannot be followed by trailing spaces or comments.

8.3.3 Graph Types (VALIDITY)

A suite definition can contain multiple graph strings that are combined to generate the final graph. There are different graph VALIDITY section headings (the heading of the suite.rc section that encloses the graph string) for cycling, one-off asynchronous, and repeating asynchronous tasks. Additionally, there may be multiple graph strings under different VALIDITY sections for cycling tasks with different dependencies at different cycle times.

The small suite definitions used in the following subsections can be found in the central database registered under the validity group name.

8.3.3.1 One-off Asynchronous Tasks

Figure 31 shows a small suite of one-off asynchronous tasks; these have no associated cycle time and don't spawn successors (once they're all finished the suite just exits). The integer 1 attached to each graph node is just an arbitrary label, akin to the task cycle time in cycling tasks; it increments when a repeating asynchronous task (below) spawns.

```
# SUITE.RC
title = some cycling tasks
# (no dependence between cycles here)
[scheduling]
  [[dependencies]]
    [[[0,12]]]
      graph = "foo => bar & baz => waz"
```

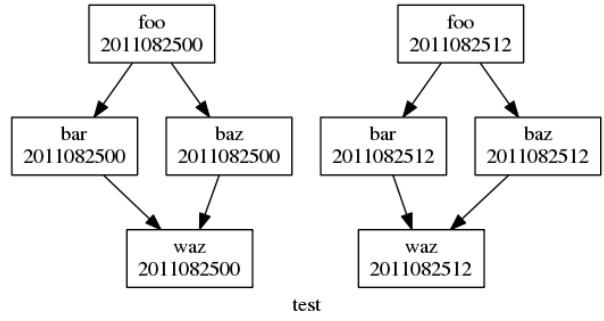


Figure 32: Cycling Tasks.

8.3.3.2 Cycling Tasks

Valid cycle times for cycling tasks are defined by the graph VALIDITY section headings - lists of hours in the day - for the graph strings in which the tasks appear. Figure 32 shows a small suite of cycling tasks.

8.3.3.2.1 How Multiple Graph Strings Combine

For a cycling graph with multiple validity sections for different hours of the day, the different sections *add* to generate the complete graph. Different graph sections can overlap (i.e. the same hours may appear in multiple section headings) and the same tasks may appear in multiple sections, but individual dependencies must be unique across the entire graph. Thus the following graph is valid:

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[0,6,12,18]]]
      graph = "A => B => C"
    [[[6,18]]]
      # OK (X triggers off C only at 6 and 18 hours)
      graph = "C => X"
```

But this is an error:

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[0,6,12,18]]]
      graph = "A => B => C"
    [[[6,18]]]
      # ERROR (B => C is already defined at 6 and 18 hours)
      graph = "B => C => X" # ERROR
```

Currently this is one of the few remaining errors not caught by suite validation. Instead the suite will abort on start-up with a “duplicate prerequisite” error:

```
cycl run foo 2012080806
# (some startup stdout...)
C%2012080806: Duplicate prerequisite: 'B%2012080806 succeeded'
/run foo 2012080806 failed: 1
```

8.3.3.3 Combined Asynchronous And Synchronous Graphs

Cycling tasks can be made to wait on one-off asynchronous tasks, as shown in Figure 33. Alternatively, they can be made to wait on one-off synchronous start-up tasks, which have an associated cycle time even though they are non-cycling - see Figure 34.

```
# SUITE.RC
title = one-off async and cycling tasks
# (with dependence between cycles too)
[scheduling]
  [[dependencies]
    graph = "prep1 => prep2"
    [[[0,12]]]
    graph = """
    prep2 => foo => bar & baz => waz
    foo[T-12] => foo
    """
  ]]
```

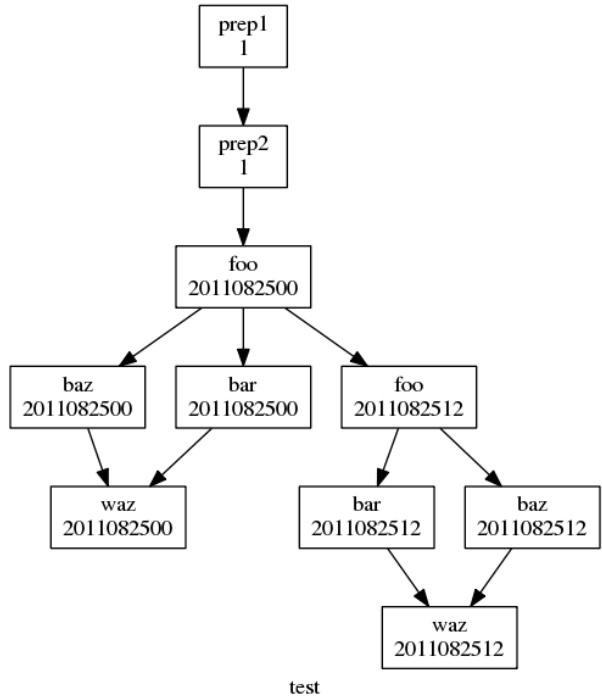


Figure 33: One-off asynchronous and cycling tasks in the same suite.

8.3.3.3.1 Synchronous Start-up vs One-off Asynchronous Tasks

One-off synchronous start-up tasks run only when a cycling suite is *cold-started* and they are often associated with subsequent one-off *cold-start tasks* used to bootstrap a cycling suite into existence.

The distinction between cold- and warm-start is only meaningful for cycling tasks, and one-off asynchronous tasks may be best used in constructing entirely non-cycling suites.

However, one-off asynchronous tasks can precede cycling tasks in the same suite, as shown above. It seems likely that, if used in this way, they will be intended as start-up tasks - so currently *one-off asynchronous tasks only run in a cold-start*.

8.3.3.4 Repeating Asynchronous Tasks

Repeating asynchronous tasks can be used to process satellite data that arrives asynchronously - i.e. at irregular time intervals. Each new dataset must have a unique “asynchronous ID” - if it doesn’t already have such an ID you could use some string representation of the data arrival time. The graph VALIDITY section heading must contain “`ASYNCCID:`” followed by a regular expression designed to match the actual IDs. Additionally, one task in the suite must be designated the “daemon” - it waits indefinitely on incoming data and reports each new dataset and its ID back to the suite by means of a special output message. When the task proxy receives that message it dynamically registers a new output (containing the asynchronous ID) that downstream tasks can then trigger off. The downstream tasks likewise have prerequisites containing the ID pattern (because they trigger off the aforementioned outputs) and when these get satisfied the actual ID is substituted into their own registered outputs. Additionally, each asynchronous task proxy passes the ID to its task execution environment as `$ASYNCCID` to allow identification of the correct dataset by task scripts. In this way the whole tree of tasks becomes dedicated to processing the new dataset, but if the data arrives quickly successive datasets can

```
# SUITE.RC
title = one-off start-up and cycling tasks
# (with dependence between cycles too)
[scheduling]
    [[special tasks]]
        start-up = prep1, prep2
    [[dependencies]]
        [[[0,12]]]
            graph = """
prep1 => prep2 => foo => bar & baz => waz
foo[T-12] => foo
"""

```

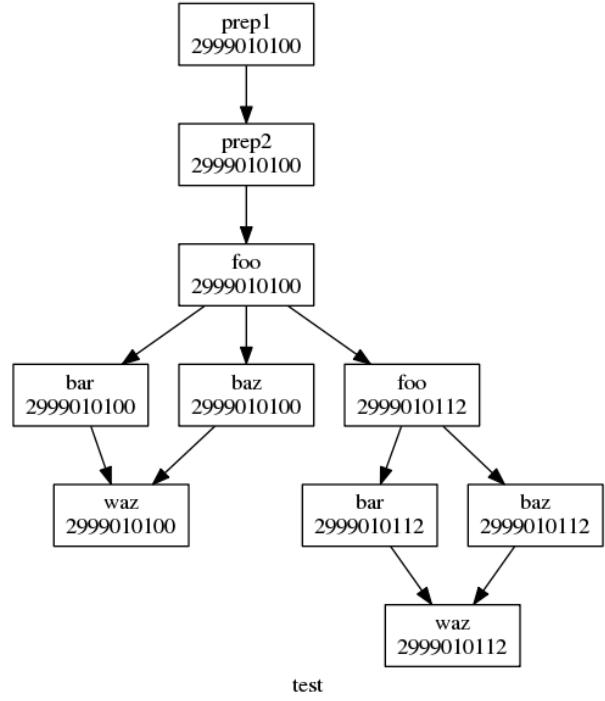


Figure 34: One-off synchronous and cycling tasks in the same suite.

be processed in parallel. As Figure 35 shows, a repeating asynchronous suite currently plots just like a one-off asynchronous suite. But at run time the daemon task stays put, while the others continually spawn successors to wait for new datasets to come in. The *asynchronous.repeating* example suite demonstrates how to do this in a real suite.

8.3.4 Trigger Types

8.3.4.1 Success Triggers

To trigger off a task finishing successfully:

```
# SUITE.RC
```

```
# SUITE.RC
title = a suite of repeating asynchronous tasks
# for processing real time satellite datasets
[scheduling]
    [[dependencies]]
        [[[ASYNCCID:satX-\d{6}]]]
            # match datasets satX-1424433 (e.g.)
            graph = "watcher:a => foo:a & bar:a => baz"
            daemon = watcher
[runtime]
    [[watcher]]
        [[[outputs]]]
            a = "New dataset <ASYNCCID> ready for processing"
    [[foo,bar]]
        [[[outputs]]]
            a = "Products generated from dataset <ASYNCCID>"
```

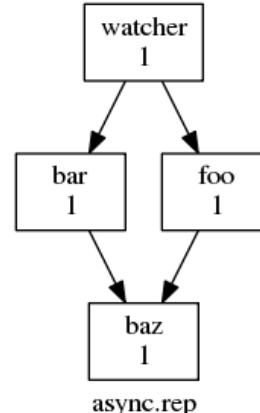


Figure 35: Repeating Asynchronous Tasks.

```
# B triggers if A SUCCEEDS:
graph = A => B
```

8.3.4.2 Failure Triggers

To trigger off a task that fails:

```
# SUITE.RC
# B triggers if A FAILS:
graph = A:fail => B
```

See *Suicide Triggers* (Section 8.3.4.6) for how to handle task B if A does not fail.

To trigger off a task finishing (success or failure):

```
# SUITE.RC
# B triggers if A either SUCCEEDS or FAILS:
graph = A | A:fail => B
```

8.3.4.3 Internal Triggers

Only required if you need to trigger before the upstream task finishes.

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[6,18]]]
      # B triggers off specific OUTPUT "out1" of task A:
      graph = A:out1 => B
[runtime]
  [[A]]
    [[[outputs]]]
      out1 = "NWP_products_uploaded_for <CYLC_TASK_CYCLE_TIME>"
```

Task A must emit this message when the actual output has been completed - see *Reporting Internal Outputs Completed*, Section 9.4.2.

8.3.4.4 Intercycle Triggers

Typically most tasks in a suite will trigger off other cotemporal tasks (same cycle time) but some may depend on tasks with earlier cycle times. This notably applies to warm-cycled forecast models, which depend on their own previous instances (see below); but other kinds of intercycle are possible too.¹¹ Here's how to express this kind of relationship in cylc:

```
# SUITE.RC
[dependencies]
  [[0,6,12,18]]
    # B triggers off A in the previous cycle
    graph = "A[T-6] => B"
```

This notation can be used with explicit outputs too:

```
# SUITE.RC
  # B triggers if A in the previous cycle fails:
  graph = "A[T-6]:fail => B"
```

8.3.4.5 Conditional Triggers

AND operators (&) can appear on both sides of an arrow. They provide a concise alternative to defining multiple triggers separately:

¹¹In NWP forecast analysis suites parts of the observation processing and data assimilation subsystem will typically also depend on model background fields generated by the previous forecast.

```
# SUITE.RC
graph = """
# D triggers if A or (B and C) succeed
A | B & C => D
# just to align the two graph sections
D => W
# Z triggers if (W or X) and Y succeed
(W|X) & Y => Z
"""

```

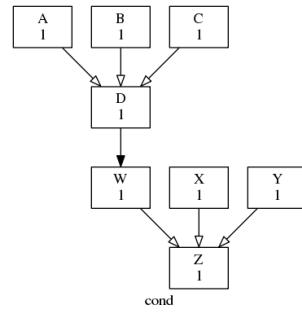


Figure 36: Conditional triggers are plotted with open arrow heads.

```
# SUITE.RC
# 1/ this:
graph = "A & B => C"
# is equivalent to:
graph = """A => C
          B => C"""
# 2/ this:
graph = "A => B & C"
# is equivalent to:
graph = """A => B
          A => C"""
# 3/ and this:
graph = "A & B => C & D"
# is equivalent to this:
graph = """A => C
          B => C
          A => D
          B => D"""

```

OR operators (`|`), for truly conditional triggers, can only appear on the left,¹²

```
# SUITE.RC
# C triggers when either A or B finishes:
graph = "A | B => C"
```

Forecasting suites typically have quite simple requirements for conditional triggers, but any valid conditional expression can be used, as shown in Figure 36 (conditional triggers are plotted with open arrowheads).

8.3.4.6 Suicide Triggers

Suicide triggers take tasks out of the suite. This can be used for automated failure recovery. The suite.rc listing and accompanying graph in Figure 37 show how to define a chain of failure recovery tasks that trigger if they're needed but otherwise remove themselves from the suite (you can run the `AutoRecover.async` example suite to see how this works). The dashed graph edges ending in solid dots indicate suicide triggers, and the open arrowheads indicate conditional triggers as usual.

8.3.4.7 Family Triggers

Family names defined by the namespace inheritance hierarchy (see Section 8.4) can be used in the suite graph as shorthand for the effective dependencies on member tasks. Member dependencies are substituted in such a way that downstream tasks will not trigger until all family members have either succeeded or failed. For example the graph in this suite:

¹²An OR operator on the right doesn't make much sense: if "B or C" triggers off A, what exactly should cyc do when A finishes?

```
# SUITE.RC
title = asynchronous automated recovery
description = """
Model task failure triggers diagnosis
and recovery tasks, which take themselves
out of the suite if model succeeds. Model
post processing triggers off model OR
recovery tasks.
"""

[scheduling]
[[dependencies]]
graph = """
pre => model
model:fail => diagnose => recover
model => !diagnose & !recover
model | recover => post
"""

[runtime]
[[model]]
# UNCOMMENT TO TEST FAILURE:
# command scripting = false
```

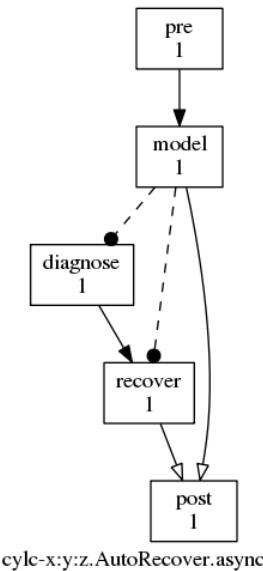


Figure 37: Automated failure recovery via suicide triggers.

```
# SUITE.RC
[scheduling]
[[dependencies]]
graph = "foo => fam => bar"
[runtime]
[[fam]] # a family (because others inherit from it)
[[a,b]] # family members (inherit namespace fam)
inherit = fam
```

is equivalent to this:

```
# SUITE.RC
[scheduling]
[[dependencies]]
graph = "foo => a & b => bar"
```

And the graph in this suite:

```
# SUITE.RC
[scheduling]
[[dependencies]]
graph = "fam:fail => bar"
[runtime]
[[fam]]
[[a,b]]
inherit = fam
```

is equivalent to this:

```
# SUITE.RC
[scheduling]
[[dependencies]]
graph = "( a:fail|b:fail ) & ( a|a:fail ) & ( b|b:fail ) => bar"
```

Task family triggers should be used sparingly as a convenient simplification for groups of tasks that would naturally trigger at the same time anyway (e.g. forecast ensembles and multiple tasks for processing different types of observations that are all made available at the same time) and whose outputs would all be put to use at a similar time too. Otherwise family triggers will unnecessarily constrain cyclc's ability to achieve maximum functional parallelism, because all

family members have to finish before downstream processing can continue.¹³

8.3.5 Handling Intercycle Dependence At Start-Up

In suites with intercycle dependence some kind of bootstrapping process is required to get the suite going initially. In the example shown in *Intercycle Triggers* (Section 8.3.4.4), for instance, in the very first cycle there is no previous instance of task A to satisfy B's prerequisites.

8.3.5.1 Cold-Start Tasks

A *cold-start* task is a special one-off task used to satisfy the initial previous-cycle dependence of another cotemporal task. In effect, the cold-start task masquerades as the previous-cycle trigger of its associated cycling task.

A cold-start task may invoke real processing, probably to generate the files that are normally produced by the associated cycling task; or it could be a dummy task that represents some external spinup process, presumably resulting in the same files, that has to be completed before the suite is started. In the latter case the cold-start task can just report itself successfully completed after checking that the required files are present.

This kind of relationship can easily be expressed with a conditional trigger:

```
# SUITE.RC
[scheduling]
  [[special_tasks]]
    cold-start = ColdFoo
  [[dependencies]]
    [[[0,6,12,18]]]
      graph = "ColdFoo | Bar[T-6] => Foo"
```

i.e. Foo[T] can trigger off *either* Bar[T-6] *or* ColdFoo[T]. At start-up ColdFoo will do the job, and thereafter Bar[T-6] will do it.

Cold-start tasks can also be inserted into the suite at run time to cold-start just their associated cycling tasks, if a problem of some kind prevents continued normal cycling.

8.3.5.2 Warm-Starting A Suite

Cold-start tasks have to be declared as such in the suite.rc "special tasks" section so that cylc knows they are one-off (non-spawning) tasks, but also because they play a critical role in suite warm-starts. A suite that has previously been running and was then shut down can be warm-started at a particular cycle time, an alternative to *restarting* from a previous state (although restarting is preferred because a warm start is likely to involve re-running some tasks). A warm-start assumes the existence of a previous cycle (i.e. that any files from the previous cycle required by the new cycle are in place already) so cold-start tasks do not need to run *but* cylc itself does not know the details of the previous cycle (it does in a restart, but not in a warm-start) so it still has to solve the bootstrapping problem to get the suite started. It does this by starting the suite with designated cold start tasks in the *succeeded* state - in other words finished cold start tasks stand in for the previous finished cycle, rather than pretending to be a running previous cycle as they do in a cold-start.

8.3.6 Model Restart Dependencies

Warm cycled forecast models generate *restart files*, e.g. model background fields, that are required to initialize the next forecast (this is essentially the definition of "warm cycling"). In

¹³Actually downstream tasks can also trigger off individual family members, rather than off the family as a whole, but extensive use of this would probably defeat the purpose of using a family in the first place.

fact restart files will often be written for a whole series of subsequent cycles in case the next cycle (or the next and the next-next, and so on) cycle has to be omitted:

```
# SUITE.RC
[scheduling]
  [[special tasks]]
    sequential = A
  [[dependencies]]
    [[[0,6,12,18]]]
      # Model A cold-start and restart dependencies:
      graph = "ColdA | A[T-6] | A[T-12] | A[T-18] | A[T-24] => A"
```

In other words, task A can trigger off a cotemporal cold-start task, *or* off its own previous instance, *or* off the instance before that, and so on. Restart dependencies are unusual because although A *could* trigger off A[T-12] we don't actually want it to do so unless A[T-6] fails and can't be fixed. *This is why Task A, above, is declared to be ‘sequential’.*¹⁴ Sequential tasks do not spawn a successor until they have succeeded (by default, tasks spawn as soon as they start running in order to get maximum functional parallelism in a suite) which means that A[T+6] will not be waiting around to trigger off an older predecessor while A[T] is still running. If A[T] fails though, the operator can force it, on removal, to spawn A[T+6], whose restart dependencies will then automatically be satisfied by the older instance, A[T-6].

Forcing a model to run sequentially means, of course, that its restart dependencies cannot be violated anyway, so we might just ignore them. This is certainly an option, but it should be noted that there are some benefits to having your suite reflect all of the real dependencies between the tasks that it is managing, particularly for complex multi-model operational suites in which the suite operator might not be an expert on the models. Consider such a suite in which a failure in a driving model (e.g. weather) precludes running one or more cycles of the downstream models (sea state, storm surge, river flow, ...). If the real restart dependencies of each model are known to the suite, the operator can just do a recursive purge to remove the subtree of all tasks that can never run due to the failure, and then cold-start the failed driving model after a gap (skipping as few cycles as possible until the new cold-start input data are available). After that the downstream models will kick off automatically so long as the gap is spanned by their respective restart files, because their restart dependencies will automatically be satisfied by the older pre-gap instances in the suite. Managing this kind of scenario manually in a complex suite can be quite difficult.

Finally, if a warm cycled model is declared to have explicit restart outputs, and is not declared to be sequential, and you define appropriate labeled restart outputs which *must contain the word ‘restart’*, then the task will spawn as soon its last restart output is completed so that successives instances of the task will be able to overlap (i.e. run in parallel) if the opportunity arises. Whether or not this is worth the effort depends on your needs.

```
# SUITE.RC
[scheduling]
  [[special tasks]]
    explicit restart outputs = A
  [[dependencies]]
    [[[0,6,12,18]]]
      graph = "ColdA | A[T-18]:res18 | A[T-12]:res12| A[T-6]:res6 => A"
[runtime]
  [[A]]
    [[[outputs]]]
      r6 = restart files completed for <CYLC_TASK_CYCLE_TIME+6>
      r12 = restart files completed for <CYLC_TASK_CYCLE_TIME+12>
      r18 = restart files completed for <CYLC_TASK_CYCLE_TIME+18>
```

¹⁴A warm cycling model that only writes out one set of restart files, for the very next cycle, does not need to be declared sequential because this early triggering problem cannot arise.

8.4 Runtime Properties - Task Execution

The `[runtime]` section of a suite.rc file configures what to execute, and where and how, when tasks are ready to run. The sections immediately below `[runtime]` are, for want of a better term, *namespaces*¹⁵ that define runtime properties for individual tasks and families of tasks.

Every namespace contains the same set of configuration items (see the *Suite.rc Reference*, Appendix A, for the complete list of items). Namespaces can inherit from other namespaces, overriding inherited items as required; *this allows configuration of related tasks without repetition*. A namespace represents a family if other namespaces inherit from it.

A namespace that does not explicitly inherit from another automatically inherits from the *root* namespace (below). Namespaces thus form a tree-like hierarchy of nested task families, rooted on the root namespace, in which the leaves are the individual tasks of the suite.

Nested families from the namespace inheritance hierarchy, even if they are not used as family triggers in the graph, can be expanded or collapsed in suite graphs, and in the graph-based suite control GUI (this will also be added to the text treeview control GUI in a future release). See *Visualization* (Section 8.5) for more on this.

The following listing of the *namespace.one* example suite (which you can import from the central database if you like) illustrates basic runtime property inheritance. How it works should be reasonable clear by inspection; if not read on.

```
# SUITE.RC
title = "User Guide [runtime] example."
[cyclc]
    simulation mode only = True # (no task implementations)
[scheduling]
    initial cycle time = 2011010106
    final cycle time = 2011010200
    [[dependencies]]
        graph = "foo => obs => bar"
            # (this is shorthand for "foo => land & ship => bar")
[runtime]
    [[root]] # base namespace for all tasks (defines suite-wide defaults)
        [[[job submission]]]
            method = at_now
        [[[environment]]]
            COLOR = red
    [[obs]] # family (inherited by land, ship); implicitly inherits root
        [[[environment]]]
            RUNNING_DIR = $HOME/running/$CYLC_TASK_NAME
    [[land]] # a task (a leaf on the inheritance tree) in the obs family
        inherit = obs
        description = land obs processing
        command scripting = run-land.sh
    [[ship]] # a task (a leaf on the inheritance tree) in the obs family
        inherit = obs
        description = ship obs processing
        command scripting = run-ship.sh
        [[[job submission]]]
            method = loadleveler
        [[[environment]]]
            RUNNING_DIR = $HOME/running/ship # override obs environment
            OUTPUT_DIR = $HOME/output/ship # add to obs environment
    [[foo]]
        # (just inherits from root)

    # The task [[bar]] is implicitly defined by its presence in the
    # graph. It is also a dummy task that just inherits from root.
```

¹⁵We use the term namespace in loose analogy with its meaning in modern programming languages. Possible future enhancements to cyclc, such as ability to import specific items from other namespaces rather than just wholesale inheritance, may tighten the analogy.

8.4.1 Namespace Names

Namespace names may contain letters, digits, underscores, and hyphens. They may not contain colons as that would preclude use of suite registration names in shell \$PATH variables. The ‘.’ character is the suite registration hierarchy delimiter (which separates suite registration groups and names, e.g. my_suites.test.foo). *Task names should not be hardwired into task implementations* because task and suite identity can be extracted portably from the task execution environment supplied by cylc (Section 8.4.4) - then to rename a task, can just change its name in the suite.rc file.

8.4.2 Root - Runtime Defaults

The root namespace, at the base of the inheritance hierarchy, provides default configuration for all tasks in the suite.

Most root items are unset by default, but some values are set, sufficient to allow simple test suites to be defined by dependency graph alone - as can be seen from many of the examples that illustrate this User Guide (command scripting, for example, defaults to printing a simple message, sleeping for ten seconds, and then exiting). These default values are documented with each configuration item in Appendix A, and Section A.7 shows them all in context. You can override them or provide your own defaults for other items by explicitly configuring the root namespace.

8.4.3 Defining Multiple Namespaces At Once

Groups of similar tasks or families that differ only in a few configuration details can be defined in single namespace sections by putting a list of names in the section heading. Potential applications include groups of similar obs processing tasks, and forecasting ensembles.

As the suite.rc file is parsed, any occurrence of <NAMESPACE> in any runtime configuration item is immediately replaced with the actual name of the namespace; and after inheritance processing any occurrence of <TASK> is replaced with the actual name of the task.

The following example illustrates this:

```
# SUITE.RC
[runtime]
  [[foo]]
    command scripting = "echo hello from <TASK> via <NAMESPACE>"
  [[bar, baz]]
    inherit = foo

% cylc get-config test runtime foo 'command scripting'
['echo Hello from foo via root']
% cylc get-config test runtime bar 'command scripting'
['echo Hello from bar via root']
% cylc get-config test runtime baz 'command scripting'
['echo Hello from baz via root']
```

As another example, consider a suite containing an ensemble of closely related forecast tasks in which each task invokes the same external script but with a unique argument that reflects the task name (which presumably results in the right initial conditions being used, or similar):

```
# SUITE.RC
[runtime]
  [[ensemble]]
    command scripting = "run-model.sh ic-<TASK>"
  [[m1, m2, m3]]
    # ensemble member tasks
    inherit = ensemble
```

This defines a unique command line for each ensemble member task:

```
% cylc get-config test runtime m2 'command scripting'
['run-model.sh ic-m2']
```

For large ensembles Jinja2 template processing can be used to automatically generate the member names and associated dependencies (see Section 8.6).

Note that in the namespace scripting and environment sections, which are evaluated in task job scripts at run time, <TASK> is functionally equivalent to `$CYLC_TASK_NAME`. But in other runtime config items that are never actually interpreted by the shell (e.g. `description`) <TASK> must be used to specialize from the list of names in the namespace section heading to a particular task name.

The multidef example suite illustrates use of these internal variables.

8.4.4 Task Execution Environment

The task execution environment contains suite and task identity variables provided by cylc, and user-defined environment variables. The environment is explicitly exported (by the task job script) prior to executing task command scripting (see *Task Job Submission*, Section 10).

Suite and task identity are exported first, so that user-defined variables can refer to them. Order of definition is preserved throughout so that variable assignment expressions can safely refer to previously defined variables.

Additionally, access to cylc itself is configured prior to the user-defined environment, so that variable assignment expressions can make use of cylc utility commands:

```
# SUITE.RC
[runtime]
[[foo]]
  [[[environment]]]
    REFERENCE_TIME = $( cylc util cycletime --add=6 )
```

8.4.4.1 User Environment Variables

The user-defined environment is the sum of a task's inherited `[[[environment]]]` sections.

```
# SUITE.RC
[runtime]
[[root]]
  [[[environment]]]
    COLOR = red
    SHAPE = circle
[[foo]]
  [[[environment]]]
    COLOR = blue # root override
    TEXTURE = rough # new variable
```

This results in a task `foo` with `SHAPE=circle`, `COLOR=blue`, and `TEXTURE=rough` in its environment.

8.4.4.2 Overriding Environment Variables

When you override inherited namespace items the original parent item definition is *replaced* by the new definition. This applies to all items including those in the environment sub-sections which, strictly speaking, are not “environment variables” until they are written, post inheritance processing, to the task job script that executes the associated task. Consequently, if you override an environment variable you cannot also access the original parent value:

```
# SUITE.RC
[runtime]
[[foo]]
  [[[environment]]]
    COLOR = red
[[bar]]
```

```

inherit = foo
[[[environment]]]
tmp = $COLOR          # !! ERROR: $COLOR is undefined here
COLOR = dark-$tmp    # !! as this overrides COLOR in foo.

```

The compressed variant of this, `COLOR = dark-$COLOR`, is also in error for the same reason. To achieve the desired result you must use a different name for the parent variable:

```

# SUITE.RC
[runtime]
  [[foo]]
    [[[environment]]]
      FOO_COLOR = red
  [[bar]]
    inherit = foo
    [[[environment]]]
      COLOR = dark-$FOO_COLOR # OK

```

8.4.4.3 Suite And Task Identity Variables

The task identity variables provided to tasks by cylc are:

```

$CYLC_TASK_ID           # X%2011051118 (e.g.)
$CYLC_TASK_NAME          # X
$CYLC_TASK_CYCLE_TIME    # 2011051118
$CYLC_TASK_LOG_ROOT      # ~/cylc-run/foo.bar.baz/log/job/X%2011051118-1317298378.191123
$CYLC_TASK_NAMESPACE_HIERARCHY # "root postproc X" (e.g.)

```

And the suite identity variables are:

```

$CYLC_SUITE_DEF_PATH   # $HOME/mysuites/baz (e.g.)
$CYLC_SUITE_REG_NAME    # foo.bar.baz (e.g.)
$CYLC_SUITE_REG_PATH    # foo/bar/baz
$CYLC_SUITE_HOST         # orca.niwa.co.nz (e.g.)
$CYLC_SUITE_PORT         # 7766 (e.g.)
$CYLC_SUITE_OWNER        # oliverh (e.g.)

```

The variable `$CYLC_SUITE_REG_PATH` is just `$CYLC_SUITE_REG_NAME` (the hierarchical name under which the suite definition is registered in your suite database) translated into a directory path. This can be used when configuring suite logging directories and the like to put suite output in a directory tree that reflects the suite registration hierarchy (as opposed to the namespace hierarchy).

Some of these variables are also used by cylc task messaging commands in order to automatically target the right task proxy object in the right suite.

8.4.4.4 Suite Share And Task Work Directories

The following variables, and the directories they represent, are also available to running tasks:

```

$CYLC_TASK_WORK_PATH     # task running directory
$CYLC_SUITE_SHARE_PATH    # suite shared directory

```

These two directories are configured in the suite.rc file, and they are created on the fly, if necessary, by task job scripts prior to executing the task command scripting. The job script changes directory to the work/running directory before executing the task command scripting. The shared directory is intended as a shared data directory for multiple tasks (however, the directory path can be different for different families of tasks or individual tasks if you like).

8.4.4.5 Environment Variable Evaluation

Variables in the task execution environment are not evaluated in the shell in which the suite is running prior to submitting the task. They are written in unevaluated form to the job script

that is submitted by cylc to run the task (Section 10.1) and are therefore evaluated when the task begins executing under the task owner account on the task host. Thus `$HOME`, for instance, evaluates at run time to the home directory of task owner on the task host.

8.4.5 Remote Task Hosting

If a task declares an owner other than the suite owner and/or a host other than the suit host, e.g.:

```
# SUITE.RC
[runtime]
[[foo]]
  [[[remote]]]
    host = orca.niwa.co.nz
    owner = bob
    cylc directory = /path/to/remote/cylc/installation/on/foo
    suite definition directory = /path/to/remote/suite/definition/on/foo
```

cylc will attempt to execute the task on the declared host, by the configured job submission method, as the declared owner, using passwordless ssh.

- passwordless ssh must be configured between the suite owner on the suite host and the task owner on the remote host.
- cylc and Pyro must be installed on the remote host so that the remote task can communicate with the suite (but graphviz is not needed).
- the suite definition directory must be installed on the remote host, if the task needs access to scripts in the suite bin directory or to any other files stored there.

A local task to run under another user account is treated as a remote task.

You may not need this functionality if you have a cross-platform resource manager, such as loadleveler, that allows you to submit a job locally to run on the remote host.

Remote host functionality, like other namespace properties, can be declared globally (in the root namespace) or per family, or per individual task. Use the global settings if all or most of your tasks need to run on the same remote host.

The remote cylc directory must be used to give remote tasks access to cylc commands if cylc is not in the default `$PATH` on the remote host. Similarly the remote suite directory gives task access to suite files via `$CYLC_SUITE_DEF_PATH` on the remote host and to the suite bin directory via `$PATH`. If a remote suite definition directory is not given the local path will be used with the local user's home directory, if present, substituted for the literal `$HOME` for evaluation on the remote host.

Note that you can easily run the cylc example suites on a remote host. For the example suites with task implementations that work with “real” files, you’ll have to run *all* tasks on the same remote host so that they can access their common input and output files. To distribute a suite across several hosts you must arrange (using additional tasks) to transfer files between the hosts as required to satisfy the real I/O dependencies.

8.4.5.1 Remote Log Directories

The stdout and stderr from local tasks is directed into files in the *job submission log directory* (specified in the suite.rc file) as explained in Section 10.3. The same goes for remotely hosted tasks, except that the task owner’s home directory is substituted as described above for the remote suite definition directory. Remote log directories are created on the fly by cylc, during job submission, if they do not already exist.

8.5 Visualization

This is the final major section in the suite.rc file. It is used to configure suite graph plotting - principally graph node (task) and edge (dependency arrow) style attributes. Tasks can be grouped for the purpose of applying common style attributes. See the suite.rc reference (Appendix A) for details.

8.5.1 Collapsible Task Families In Suite Graphing And GUIs

```
# SUITE.RC
[visualization]
# list namespace families to be shown in collapsed form
collapsed families = family1, family2
```

Nested families from the namespace inheritance hierarchy, even if they are not used as family triggers in the graph, can be expanded or collapsed in suite graphs and by menu options in the graph-based suite control GUI (this will also be added to the text treeview control GUI in a future release).

Family nodes in the graph-based suite control GUI are not currently color-coded in real time according to what state their members are in. However, any ungraphed tasks, which includes the members of collapsed families, are automatically plotted as rectangular nodes to the right of the main graph if they are doing anything interesting (submitted, running, or failed).

Note that family relationships can be defined purely for visualization purposes - you can group tasks at root level in the inheritance hierarchy prior to defining real properties at higher levels.

Figure 38 illustrates successive expansion of nested task families in the *namespaces* example suite, which has the following namespace hierarchy:

```
% cylc list --tree cylc-x-y-z.namespaces
root
|-GEN
| |-OPS
| | |-aircraft    OPS aircraft obs processing
| | |-atovs      OPS ATOVS obs processing
| | `-'atovs_post OPS ATOVS postprocessing
| |-VAR
| | |-AnPF       runs VAR AnalysePF
| | `-'ConLS     runs VAR ConfigureLS
|-baz
| |-bar1        Task bar1 of baz
| `-'bar2        Task bar2 of baz
|-foo          No description provided
`-prepobs     obs preprocessing
```

8.6 Jinja2 Suite Templates

Support for the Jinja2 template processor adds general variables, mathematical expressions, loop control structures, and conditional expressions to suite.rc files - which are automatically preprocessed to generate the final suite definition seen by cylc.

The need for Jinja2 processing must be declared with a hash-bang comment as the first line of the suite.rc file:

```
#!Jinja2
# ...
```

Potential uses for this include automatic generation of repeated groups of similar tasks and dependencies, and inclusion or exclusion of entire suite sections according to the value of a single flag. Consider a large complicated operational suite and several related parallel test suites with slightly different task content and structure (the parallel suites, for instance, might take certain

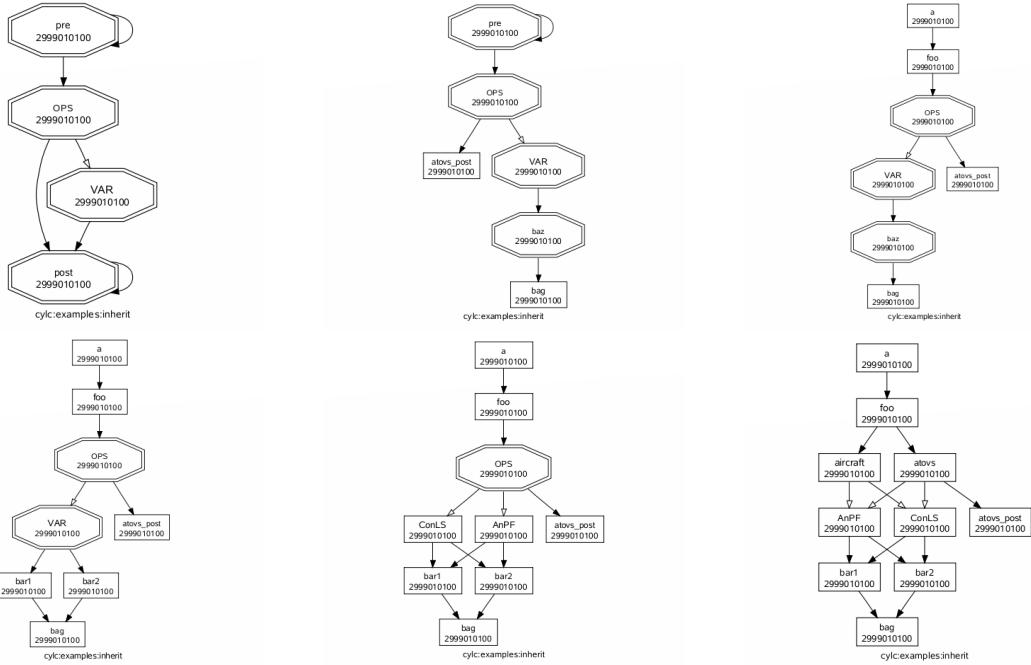


Figure 38: Graphs of the *namespaces* example suite showing various states of expansion of the nested namespace family hierarchy, from all families collapsed (top left) through to all expanded (bottom right). This can also be done by right-clicking on tasks in the graph-based suite control GUI.

large input files from the operation or the archive rather than downloading them again) - these can now be maintained as a single master suite definition that reconfigures itself according to the value of a flag variable indicating the intended use.

Template processing is the first thing done on parsing a suite definition so Jinja2 expressions can appear anywhere in the file (inside strings and namespace headings, for example).

Jinja2 is well documented at <http://jinja.pocoo.org/docs>, so here we just provide an example suite that uses it. The meaning of the embedded Jinja2 code should be reasonably self-evident to anyone familiar with standard programming techniques.

The `jinja2.ensemble` example, graphed in Figure 39, shows an ensemble of similar tasks generated using Jinja2:

```
% cylc view --stdout --central <AdminUserName>.cylc-4-0-0.jinja2.ensemble
#!jinja2
{% set N_MEMBERS = 5 %}
[scheduling]
  [[dependencies]]
    graph = """#{ generate ensemble dependencies #
      { for I in range( 0, N_MEMBERS ) %}
        foo => mem_{{ I }} => post_{{ I }} => bar
      { endfor }"""

```

Here is the generated suite definition, after Jinja2 processing:

```
% cylc view --stdout --processed --central <AdminUserName>.cylc-4-0-0.jinja2.ensemble
#!jinja2
[scheduling]
  [[dependencies]]
    graph =
      foo => mem_0 => post_0 => bar
      foo => mem_1 => post_1 => bar
      foo => mem_2 => post_2 => bar
      foo => mem_3 => post_3 => bar
      foo => mem_4 => post_4 => bar
      """

```

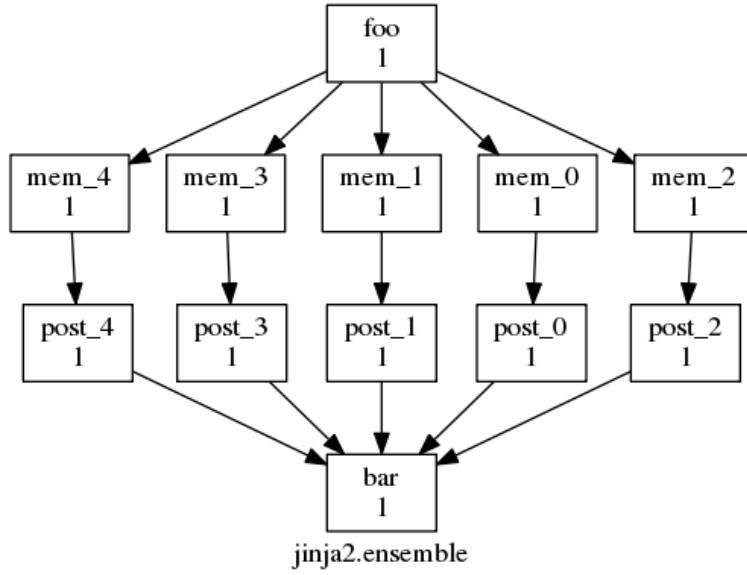


Figure 39: The Jinja2 ensemble example suite graph.

And finally, the `jinja2.cities` example uses variables, includes or excludes special cleanup tasks according to the value of a logical flag, and it automatically generates all dependencies and family relationships for a group of tasks that is repeated for each city in the suite. To add a new city and associated tasks and dependencies simply add the city name to list at the top of the file. The suite is graphed, with the New York City task family expanded, in Figure 40.

```

#!Jinja2

title = "Jinja2 city suite example."
description = """
Illustrates use of variables and math expressions, and programmatic
generation of groups of related dependencies and runtime properties."""

{% set HOST = "SuperComputer" %}
{% set CITIES = 'NewYork', 'Philadelphia', 'Newark', 'Houston', 'SantaFe', 'Chicago' %}
{% set CITYJOBS = 'one', 'two', 'three', 'four' %}
{% set LIMIT_MINS = 20 %}

{% set CLEANUP = True %}

[scheduling]
  [[ dependencies ]]
{% if CLEANUP %}
  [[[23]]]
    graph = "clean"
{% endif %}
  [[[0,12]]]
    graph = """
      setup => get_lbc & get_ic # foo
      {% for CITY in CITIES %} (# comment #
        get_lbc => {{ CITY }}_one
        get_ic => {{ CITY }}_two
        {{ CITY }}_one & {{ CITY }}_two => {{ CITY }}_three & {{ CITY }}_four
      {% if CLEANUP %}
        {{ CITY }}_three & {{ CITY }}_four => cleanup
      {% endif %}
      {% endfor %}
    """
  
```

[runtime]

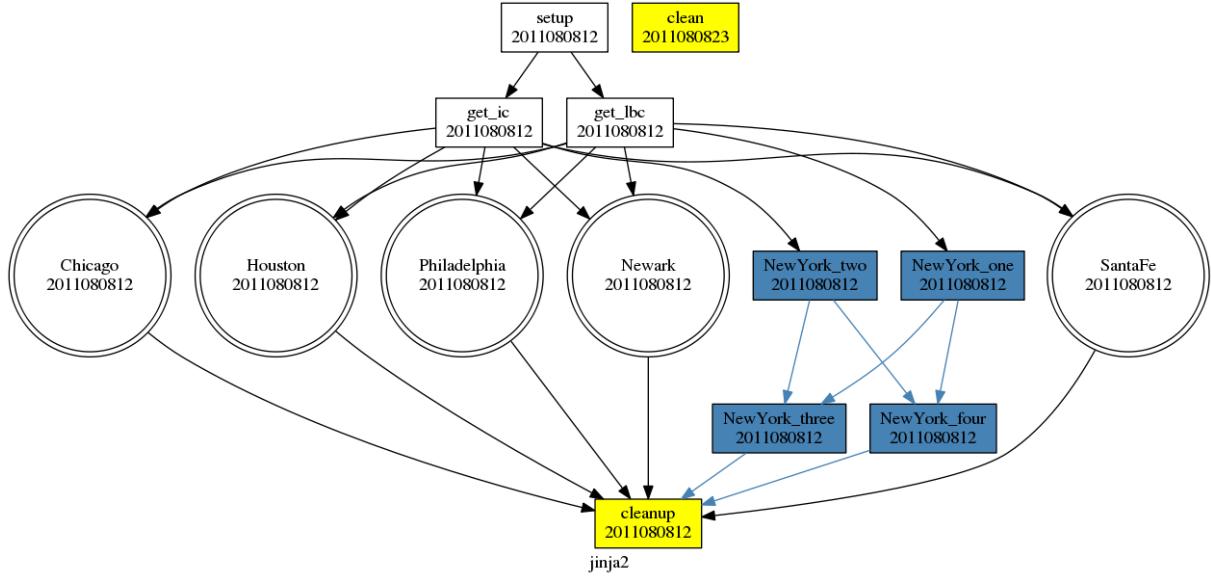


Figure 40: The Jinja2 cities example suite graph, with the New York City task family expanded.

```

[[on_{{ HOST }}]]
  [[[remote]]]
    host = {{ HOST }}
    cylc directory = /path/to/cylc
  [[[directives]]]
    wall_clock_limit = "00:{{ LIMIT_MINS|int() + 2 }}:00,00:{{ LIMIT_MINS }}:00"

{% for CITY in CITIES %}
  [[ {{ CITY }} ]]
    inherit = on_{{ HOST }}
{% for JOB in CITYJOBS %}
  [[ {{ CITY }}_{{ JOB }} ]]
    inherit = {{ CITY }}
{% endfor %}
{% endfor %}

[visualization]
  initial cycle time = 2011080812
  final cycle time = 2011080823
  [[node groups]]
    cleaning = clean, cleanup
  [[node attributes]]
    cleaning = 'style=filled', 'fillcolor=yellow'
    NewYork = 'style=filled', 'fillcolor=lightblue'
  {% for JOB in CITYJOBS %}
    NewYork_{{ JOB }} = 'style=filled', 'fillcolor=lightblue'
  {% endfor %}

```

8.7 Special Placeholder Variables

A small number of special variables are used as placeholders in cylc suite definitions:

- <CYLC_TASK_CYCLE_TIME> and <CYLC_TASK_CYCLE_TIME+offset>

These are replaced with the actual cycle time of the task (with optional computed positive offset) in the message strings registered for explicit internal (pre-completion) task outputs. They correspond to `$CYLC_TASK_CYCLE_TIME` in the task execution environment.

- `<TASK>` and `<NAMESPACE>`

These are replaced with actual task or namespace names in runtime settings. The difference between the two is explained in Section 8.4.3, *Defining Multiple Namespaces At Once*.

- `<ASYNCID>`

This is replaced with the actual “asynchronous ID” (e.g. satellite pass ID) for repeating asynchronous tasks.

To use proper variables (c.f. programming languages) in suite definitions, see the Jinja2 template processor (Section 8.6).

8.8 Omitting Defined Tasks At Runtime

It is sometimes convenient to omit certain tasks from the suite at runtime without actually deleting their definitions entirely from the suite definition.

Defining [runtime] properties for tasks that do not appear in the suite graph results in (verbose-mode) validation warnings that the tasks are disabled - they cannot be used, either for insertion into the running suite with `cylc insert` or outside of the suite with `cylc submit`, because if omitted from the graph cylc does not know their valid cycle times (or even if they are cycling tasks) or their dependencies (for insertion). Nevertheless, it is legal to leave these orphaned runtime sections in the suite definition because it allows you to temporarily remove tasks from the suite by simply commenting them out of the graph.

To omit a task from the suite at runtime but still leave it fully defined and available for insertion or lone submission, use one or both of the [scheduling][[special task]] lists, “tasks to include at start-up” and “tasks to exclude at start-up” (documented in Sections A.3.5.8 and A.3.5.7). In this case the graph still defines the validity of and dependencies involving the omitted tasks, but they are not actually instantiated in the suite at start-up. Other tasks that depend on the omitted ones will have to wait on their insertion at a later time or otherwise be triggered manually.

9 Task Implementation

This section lays out the minimal requirements on external commands, scripts, or executables invoked by cylc to carry out task processing.

9.1 Most Tasks Require No Modification For Cylc

Any existing command, script, or executable can function as a cylc task (or rather, perform the external processing that the task represents), with the following conditions exceptions:

- Tasks with internal (pre-completion) outputs that other tasks need to trigger off must be modified to report to the suite when those outputs have been completed.
- Tasks that spawn secondary internal processes which they then detach from and exit early must be modified so that the secondary processes report final success or failure to the suite.
- Tasks that do not return non-zero exit status on error must be made to do so, to allow cylc’s automatic error trapping to work. This is normal best-practice scripting anyway.

If these requirements are not met, see *Tasks Requiring Modification For Cylc*, Section 9.4.

The following suite runs a couple of external scripts that are not cylc-aware, but which meet the requirements above so that no special treatment is required at all:

```
# SUITE.RC
[runtime]
[[foo]]
  description = a task that runs foo.sh
  command scripting = foo.sh OPTIONS ARGUMENTS
[[bar]]
  description = a task that runs bar.sh
  command scripting = """echo HELLO
                        bar.sh
                        echo BYE"""
[[baz]]
  description = a task that runs baz.sh and retries on failure
  command scripting = """echo attempt No.1
                        baz.sh""",
    """echo attempt No.2 # only invoked if try No.1 fails
                        baz.sh --retry"""

```

9.2 Suite.rc Inlined Tasks

Simple tasks can be entirely implemented within the suite.rc file because the task *command scripting* string can contain as many lines of code as you like (or even a list of such strings, for retry on failure).

9.3 Voluntary Messaging Modifications

You can, if you like, modify task scripts to send explanatory or progress messages to the suite as the task runs. For example, a task can send a priority critical message before aborting on error:

```
#!/bin/bash
set -e # abort on error
if ! mkdir /illegal/dir; then
  # (use inline error checking to avoid triggering the above 'set -e')
  cylc task message -p CRITICAL "Failed to create directory /illegal/dir"
  exit 1 # now abort non-zero exit status to trigger the task failed message
fi
```

You can also use this syntax:

```
#!/bin/bash
set -e
mkdir /illegal/dir || { # inline error checking using OR operator
  cylc task message -p CRITICAL "Failed to create directory /illegal/dir"
  exit 1
}
```

But not this:

```
#!/bin/bash
set -e
mkdir /illegal/dir # aborted via 'set -e'
if [[ $? != 0 ]]; then # so this will never be reached.
  cylc task message -p CRITICAL "Failed to create directory /illegal/dir"
  exit 1
fi
```

You can also send warning messages, or general information:

```
#!/bin/bash
# a warning message (this will be logged by the suite):
cylc task message -p WARNING "oops, something's fishy here"
# information (this will also be logged by the suite):
cylc task message "Hello from task foo"
```

This may be useful - any message received from a task is logged by cylc - but it is not a requirement. If error messages are not reported, for instance, detected task failures will still

be registered, and task stdout and stderr logs can still be examined for evidence of what went wrong.

9.4 Tasks Requiring Modification For Cyc

There are two main categories of tasks that require some modification to work with cyc: those with internal (pre-completion) outputs that other tasks need to trigger off, and those that spawn detaching internal processes prior to exiting before the spawned processing is finished. Additionally, any task that fails to indicate a fatal error by return non-zero exit status must be corrected.

9.4.1 Returning Non-zero Exit Status On Error

The requirement to abort with non-zero exit status on error (which should be normal scripting practice in any case) allows the task job script to trap errors and send a `cyclc task failed` message to alert the suite. You can use `set -e` to avoid writing explicit error checks for every operation:

```
#!/bin/bash
set -e # abort on error
mkdir /illegal/dir # this error will abort the script with non-zero exit status
```

See Section 9.3, *Voluntary Messaging Modifications*, for more examples of error detection and cyc messaging.

9.4.2 Reporting Internal Outputs Completed

Tasks with internal outputs that allow downstream processing to trigger before they are finished must report when those internal outputs have been completed:

```
#!/bin/bash
# (task foo implementation)
#
# ...
# report an output completed:
cyclc task message "foo products uploaded for $CYLC_TASK_CYCLE_TIME"
```

This must match one of the task's registered outputs:

```
# SUITE.RC
[scheduling]
  [[dependencies]]
    [[[6,18]]]
      graph = foo:output1 => bar
[runtime]
  [[foo]]
    [[outputs]]
      output1 = "foo products uploaded for <CYLC_TASK_CYCLE_TIME>"
```

otherwise it will just be logged as a progress report or similar. Note the angle brackets around the cycle time variable in the output message as registered in the suite.rc file. *Do not write this as \$CYLC_TASK_CYCLE_TIME in the suite.rc file.* It is just a special tag that cyc replaces with the actual cycle time of the task; it is not an environment variable, although in this particular case it has the same value that is exported to `$CYLC_TASK_CYCLE_TIME` in the task execution environment.

9.4.3 Reconnecting Detaching Processes

You may be able to convert a detaching task to a non-detaching task very easily. If the detaching process is just a background shell process, for instance, run it in the foreground instead; for loadleveler the `-s` option prevents `llsubmit` from returning until the job has completed; for Sun Grid Engine, `qsub -sync yes` has the same effect. Section 10.4 shows how to override the

command template used by cylc in order to customize job submission command options like this.

9.4.4 Tasks That Detach And Exit Early

Tasks with processes that spawn jobs internally (e.g. to a batch queue scheduler or to another host) and then detach and exit without seeing the resulting processing through must arrange for the spawned processing to send its own “cylc task succeeded” or “cylc task failed” messages on completion, because the cylc-generated job script (Section 10.1) that otherwise arranges for automatic completion messaging cannot know when the task is really finished.

First check that you can’t easily “reconnect” the detaching internal processes, as described above in Section 9.4.3. If not, then start by disabling cylc’s automatic completion messaging:

```
# SUITE.RC
[runtime]
  [[root]]
    manual_completion = True    # global setting
  [[foo]]
    manual_completion = False   # task-specific setting
```

Now, reporting success or failure is just a matter of calling the cylc messaging commands:

```
#!/bin/bash
# ...
if $SUCCESS; then
    # release my task lock and report success
    cylc task succeeded
    exit 0
else
    # release my task lock and report failed
    cylc task failed "Input file X not found"
    exit 1
fi
```

Bear in mind, however, that cylc messaging commands read environment variables that identify the calling task and the target suite, so if your job submission method does not automatically copy its parent environment you must arrange for these variables, at the least, to be propagated through to your spawned sub-jobs.

One way to handle this is to write a *task wrapper* that modifies a copy of the detaching native job scripts, on the fly, to insert completion messaging in the appropriate places, and other variables if necessary, before invoking the (now modified) native process. A significant advantage of this method is that you don’t need to permanently modify the model or its associated native scripting for cylc. Another is that you can configure the native job setup for a single test case (running it without cylc) and then have your custom wrapper modify the test case on the fly with suite, task, and cycle-specific parameters as required.

To make this easier, for tasks that declare manual completion messaging, cylc makes non user-defined environment scripting available in a single variable called `$CYLC_SUITE_ENVIRONMENT` that can be inserted into the aforementioned native task scripts prior to calling the cylc messaging commands.¹⁶

9.4.5 A Custom Task Wrapper Example

The *detaching* example suite contains a script `model.sh` that runs a pseudo-model executable as follows:

¹⁶Note that `$CYLC_SUITE_ENVIRONMENT` is a string containing embedded newline characters and it has to be handled accordingly. In the bash shell, for instance, it should be echoed in quotes to avoid concatenation to a single line.

```
#!/bin/bash
set -e

MODEL="sleep 10; true"
#MODEL="sleep 10; false" # uncomment to test model failure

echo "model.sh: executing pseudo-executable"
eval $MODEL
echo "model.sh: done"
```

this is in turn executed by a script `run-model.sh` that detaches immediately after job submission (i.e. it exits before the model executable actually runs):

```
#!/bin/bash
set -e
echo "run-model.sh: submitting model.sh to 'at now'"
SCRIPT=model.sh # location of the model job to submit
OUT=$1; ERR=$2 # stdout and stderr log paths
# submit the job and detach
RES=$TMPDIR/atnow$$txt
( at now <<EOF
$SCRIPT 1> $OUT 2> $ERR
EOF
) > $RES 2>&1
if grep 'No atd running' $RES; then
    echo 'ERROR: atd is not running!'
    exit 1
fi
# model.sh should now be running at the behest of the 'at' scheduler.
echo "run-model.sh: done"
```

Note that your `at` scheduler daemon must be up if you want to test this suite.

Here's a cycl suite to run this unruly model:

```
title = "Custom Task Wrapper Example"

description = """This suite runs a single task that internally submits a
'model executable' before detaching and exiting immediately - so we have
to handle task completion messaging manually - see the Cyc User Guide."""

[scheduling]
    initial cycle time = 2011010106
    final cycle time = 2011010200
    [[special tasks]]
        sequential = model
    [[dependencies]]
        [[[0,6,12,18]]]
        graph = "model"

[runtime]
    [[model]]
        manual completion = True
        command scripting = model-wrapper.sh # invoke the task via a custom wrapper
    [[environment]]
        # location of native job scripts to modify for this suite:
        NATIVESCRIPTS = ${CYLC_SUITE_DEF_PATH}/native
        # output path prefix for detached model stdout and stderr:
        PREFIX = ${HOME}/detach
        FOO = "${HOME} bar ${PREFIX}"
```

The suite invokes the task by means of the custom wrapper `model-wrapper.sh` which modifies, on the fly, a temporary copy of the model's native job scripts as described above:

```
#!/bin/bash
set -e

# A custom wrapper for the 'model' task from the detaching example suite.
# See the Cyc User Guide for more information.

# Check inputs:
```

10 TASK JOB SUBMISSION

```
# location of pristine native job scripts:
cylc util checkvars -d NATIVESCRIPTS
# path prefix for model stdout and stderr:
cylc util checkvars PREFIX

# Get a temporary copy of the native job scripts:
TDIR=$TMPDIR/detach$$
mkdir -p $TDIR
cp $NATIVESCRIPTS/* $TDIR

# Insert task-specific execution environment in $TDIR/model.sh:
SRCH='echo "model.sh: executing pseudo-executable"'
perl -pi -e "s@^${SRCH}@${CYLC_SUITE_REG_NAME_ENVIRONMENT}\n${SRCH}@" $TDIR/model.sh

# Task completion message scripting. Use single quotes here - we don't
# want the $? variable to evaluate in this shell!
MSG='
if [[ $? != 0 ]]; then
    cylc task message -p CRITICAL "ERROR: model executable failed"
    exit 1
else
    cylc task succeeded
    exit 0
fi'
# Insert error detection and cylc messaging in $TDIR/model.sh:
SRCH='echo "model.sh: done"'
perl -pi -e "s@^${SRCH}@${MSG}\n${SRCH}@" $TDIR/model.sh

# Point to the temporary copy of model.sh, in run-model.sh:
SRCH='SCRIPT=model.sh'
perl -pi -e "s@^${SRCH}@$SCRIPT=$TDIR/model.sh@" $TDIR/run-model.sh

# Execute the (now modified) native process:
$TDIR/run-model.sh ${PREFIX}-${CYLC_TASK_CYCLE_TIME}-$$ .out ${PREFIX}-${CYLC_TASK_CYCLE_TIME}-$$ .err
echo "model-wrapper.sh: see modified job scripts under ${TDIR}!"'
# EOF
```

If you run this suite, or submit the model task alone with `cylc submit`, you'll find that the usual job submission log files for task stdout and stderr end before the task is finished. To see the "model" output and the final task completion message (success or failure), examine the log files generated by the job submitted internally to the *at* scheduler (their location is determined by the `$PREFIX` variable in the `suite.rc` file).

It should not be difficult to adapt this example to real tasks with detaching internal job submission. You will probably also need to replace other parameters, such as model input and output filenames, with suite- and cycle-appropriate values, but exactly the same technique can be used: identify which job script needs to be modified and use text processing tools (such as the single line `perl` search-and-replace expressions above) to do the job.

10 Task Job Submission

Task Implementation (Section 9) describes what requirements a command, script, or program, must fulfill in order to function as a cylc task. This section explains how tasks are submitted by cylc when they are ready to run, and how to define new task job submission methods.

10.1 Task Job Scripts

When a task is ready to run cylc generates a temporary *task job script* to configure the task's execution environment and call its command scripting. The job script is the embodiment of all `suite.rc` runtime settings for the task. It is submitted to run by the *job submission method* configured for the task. Different tasks can have different job submission methods. Like other

runtime properties, you can set a suite default job submission method and override it for specific tasks or families:

```
# SUITE.RC
[runtime]
  [[root]] # suite defaults
    [[[job submission]]]
      method = loadleveler
  [[foo]] # just task foo
    [[[job submission]]]
      method = at
```

The actual command line used to submit the job script is written to stdout by cylc. In the following shell transcript we generate a job script for a task in the QuickStart.c example suite and then examine it:

```
% cylc submit --dry-run QuickStart.c Model%2011080506
> JOB SCRIPT: ~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123
> THIS IS A DRY RUN. HERE'S HOW I WOULD SUBMIT THE TASK:
~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123 </dev/null
1> ~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123.out
2> ~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123.err &
```

And here is the generated job script:

```
#!/bin/bash

# +---+ THIS IS A CYLC TASK JOB SCRIPT +---+
# Task: Model%2011080506
# To be submitted by method: 'background'

echo "TASK JOB SCRIPT STARTING"

# CYLC LOCATION, SUITE LOCATION, SUITE IDENTITY:
export CYLC_DIR=/home/oliverh/cylc
export CYLC_MODE=submit
export CYLC_SUITE_HOST=oliverh-33586DL.greta.niwa.co.nz
export CYLC_SUITE_PORT=None
export CYLC_SUITE_DEF_PATH=/tmp/oliverh/QuickStart/c
export CYLC_SUITE_REG_NAME=QuickStart.c
export CYLC_SUITE_REG_PATH=QuickStart/c
export CYLC_SUITE_OWNER=oliverh
export CYLC_USE_LOCKSERVER=False
export CYLC_UTC=False

# TASK IDENTITY:
export CYLC_TASK_ID=Model%2011080506
export CYLC_TASK_NAME=Model
export CYLC_TASK_CYCLE_TIME=2011080506
export CYLC_TASK_LOG_ROOT=~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123
export CYLC_TASK_NAMESPACE_HIERARCHY="root Models Model"

# ACCESS TO CYLC:
PATH=$CYLC_DIR/bin:$PATH
# Access to the suite bin dir:
PATH=$CYLC_SUITE_DEF_PATH/bin:$PATH
export PATH

# SET ERROR TRAPPING:
set -u # Fail when using an undefined variable
# Define the trap handler
HANDLE_TRAP() {
  echo Received signal "$@"
  cylc task failed "Task job script received signal $@"
  trap "" EXIT
  exit 0
}
# Trap signals that could cause this script to exit:
trap "HANDLE_TRAP EXIT" EXIT
trap "HANDLE_TRAP ERR" ERR
```

```

trap "HANDLE_TRAP TERM" TERM
trap "HANDLE_TRAP XCPU" XCPU

# SEND TASK STARTED MESSAGE:
cylc task started || exit 1

# SHARE DIRECTORY CREATE:
CYLC_SUITE_SHARE_PATH=$CYLC_SUITE_DEF_PATH/share
export CYLC_SUITE_SHARE_PATH
mkdir -p $CYLC_SUITE_DEF_PATH/share || true

# WORK DIRECTORY CREATE:
CYLC_TASK_WORK_PATH=$CYLC_SUITE_DEF_PATH/work/$CYLC_TASK_ID
export CYLC_TASK_WORK_PATH
mkdir -p $(dirname $CYLC_TASK_WORK_PATH) || true
mkdir -p $CYLC_TASK_WORK_PATH
cd $CYLC_TASK_WORK_PATH

# ENVIRONMENT:
TASK_EXE_SECONDS="5"
WORKSPACE="/tmp/$USER/$CYLC_SUITE_REG_NAME/common"
MODEL_INPUT_DIR="$WORKSPACE"
MODEL_OUTPUT_DIR="$WORKSPACE"
MODEL_RUNNING_DIR="$WORKSPACE/Model"
export TASK_EXE_SECONDS WORKSPACE MODEL_INPUT_DIR MODEL_OUTPUT_DIR MODEL_RUNNING_DIR

# TASK COMMAND SCRIPTING:
Model.sh

# WORK DIRECTORY REMOVE:
cd
rmdir $CYLC_TASK_WORK_PATH 2>/dev/null || true

# SEND TASK SUCCEEDED MESSAGE:
cylc task succeeded

echo "JOB SCRIPT EXITING (TASK SUCCEEDED)"
trap "" EXIT

#EOF

```

You can also generate a job script and print it directly to stdout, with `cylc jobsript`.

10.2 Built-in Job Submission Methods

Cylc has a number of built-in job submission methods. If these do not suite your needs Section 10.5 shows how to extend cylc with new user-defined job submission methods.

10.2.1 background

This job submission method runs tasks directly in a background shell.

10.2.2 at

This job submission method submits tasks to the rudimentary Unix `at` scheduler. The `atd` daemon must be running.

10.2.3 loadleveler

This job submission method submits tasks to loadleveler using the `llsubmit` command. Loadleveler directives can be provided in the `suite.rc` file:

```
# SUITE.RC
[runtime]
```

```
[ [__NAME__]]
  [[[directives]]]
    foo = bar
    baz = waz
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
# @ foo = bar
# @ baz = waz
# @ queue
```

10.2.4 pbs

This job submission method submits tasks to PBS (or Torque) using the `qsub` command. PBS directives can be provided in the suite.rc file:

```
# SUITE.RC
[runtime]
  [ [__NAME__]]
    [[[directives]]]
      -q = foo
      -l = 'nodes=1,walltime=00:01:00'
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#PBS -q foo
#PBS -l nodes=1,walltime=00:01:00
```

10.2.5 sge

This job submission method submits tasks to Sun Grid Engine using the `qsub` command. SGE directives can be provided in the suite.rc file:

```
# SUITE.RC
[runtime]
  [ [__NAME__]]
    [[[directives]]]
      -cwd = ''
      -q = foo
      -l = 'h_data=1024M,h_rt=24:00:00'
```

These are written to the top of the task job script like this:

```
#!/bin/bash
# DIRECTIVES
#$ -cwd
#$ -q foo
#$ -l h_data=1024M,h_rt=24:00:00
```

10.2.6 Default Directives Provided

For loadleveler, pbs, and sge job submission, default directives are provided to set the job name and stdout and stderr file paths.

10.2.7 PBS and SGE Cyc Quirks

As shown in the example above, multiple entries for the same PBS or SGE directive option must be comma-separated on the same line, in the suite.rc file. Otherwise, repeating the option on another line will override the previous entry, not add to it. Also, the right-hand side must

be quoted to hide the comma from the suite.rc parser (commas indicate list values, whereas directives are treated as singular).

As also shown in the example above, to get a naked option flag such as `-cwd` in SGE you must give a quoted blank space as the directive value in the suite.rc file.

10.3 Whither Task stdout And stderr?

When a task is ready to run cylc generates task-specific stdout and stderr filenames containing the task name, cycle time, and a string of digits (seconds since epoch) to ensure uniqueness - rerunning the task won't overwrite any old output:

```
# task job script:
~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123
# task stdout:
~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123.out
# task stderr:
~/cylc-run/QuickStart.c/log/job/Model%2011080506-1317298378.191123.err
```

(the job submission log directory is configurable in the suite.rc file).

How the stdout and stderr streams are directed into these files depends on the job submission method. The `background` method just uses appropriate output redirection on the command line, as shown above. The `loadleveler` method writes appropriate directives to the job script that is submitted to loadleveler.

Cylc obviously has no control over the stdout and stderr output from tasks that do their own internal output management (e.g. tasks that submit internal jobs and direct their output to other files). For less internally complex tasks, however, these will be complete task job logs. *They can be viewed updating in real time in the suite control GUIs.*

10.4 Overriding the Job Submission Command

To change the form of the actual command used to submit a job you do not need to define a new job submission method; just override the `command template` in the relevant job submission sections of your suite.rc file:

```
# SUITE.RC
[runtime]
  [[root]]
    [[[job submission]]]
      method = loadleveler
      # Use '-s' to stop llsubmit returning until all job steps have completed:
      command template = llsubmit -s %s
```

As explained in the suite.rc reference (Appendix A), the template's first `%s` will be substituted by the job file path and, where applicable a second and third `%s` will be substituted by the paths to the job stdout and stderr files.

10.5 Defining New Job Submission Methods

Defining a new job submission method requires some minimal Python programming. You can derive (in the sense of object oriented programming inheritance) new methods from one of the existing ones, or directly from cylc's job submission base class,

```
$CYLC_DIR/lib/cylc/job_submission/job_submit.py
```

using the existing methods as examples. Most often this should merely be a matter of defining the command line used to execute the aforementioned job scripts and using the provided stdout and stderr file paths appropriately. For example, here is the entire class code for the `background` method:

```
#!/usr/bin/env python

from job_submit import job_submit

class background( job_submit ):
    """
Run the task job script directly in a background shell.
    """

    # stdin redirection (< /dev/null) allows background execution on
    # remote hosts - ssh needn't wait for the process to finish.
    COMMAND_TEMPLATE = "%s </dev/null 1>%s 2>%s &"

    def construct_jobfile_submission_command( self ):
        command_template = self.job_submit_command_template
        if not command_template:
            command_template = self.COMMAND_TEMPLATE
        self.command = command_template % ( self.jobfile_path,
                                             self.stdout_file,
                                             self.stderr_file )
```

Here is the `at` method:

```
#!/usr/bin/env python

from job_submit import job_submit

class at( job_submit ):
    """
Submit the task job script to the simple 'at' scheduler. The 'atd' daemon
service must be running.
    """

    COMMAND_TEMPLATE = "echo \"%s 1>%s 2>%s\" | at now"

    def construct_jobfile_submission_command( self ):
        command_template = self.job_submit_command_template
        if not command_template:
            command_template = self.COMMAND_TEMPLATE
        self.command = command_template % ( self.jobfile_path,
                                             self.stdout_file,
                                             self.stderr_file )
```

And here is the `pbs` method:

```
#!/usr/bin/env python
from job_submit import job_submit

class pbs( job_submit ):
    """
PBS qsub job submission.
    """

    COMMAND_TEMPLATE = "qsub %s"

    def set_directives( self ):
        self.directive_prefix = "#PBS"
        self.final_directive = None
        self.directive_connector = " "

        defaults = {}
        defaults[ '-N' ] = self.task_id
        defaults[ '-o' ] = self.stdout_file
        defaults[ '-e' ] = self.stderr_file

        # In case the user wants to override the above defaults:
        for d in self.directives:
            defaults[ d ] = self.directives[ d ]
        self.directives = defaults

    def construct_jobfile_submission_command( self ):
        command_template = self.job_submit_command_template
        if not command_template:
```

```
    command_template = self.COMMAND_TEMPLATE
    self.command = command_template % ( self.jobfile_path )
```

10.5.1 An Example

The following user-defined job submission class, called *qsub*, overrides the built-in *pbs* class to change the directive prefix from `#PBS` to `#QSUB`:

```
#!/usr/bin/env python

# to import from outside of the cylc source tree:
from cylc.job_submission.pbs import pbs
# OR, from $CYLC_DIR/lib/cylc/job_submission
# from pbs import pbs

class qsub( pbs ):
    """
    This is a user-defined job submission method that overrides the '#PBS'
    directive prefix of the built-in pbs method.
    """

    def set_directives( self ):
        pbs.set_directives( self )
        # override the '#PBS' directive prefix
        self.directive_prefix = "#QSUB"
```

To check that this works correctly save the new source file to `qsub.py` in one of the allowed locations (see just below), use it in a suite definition:

```
# SUITE.rc
# $HOME/test/suite.rc
[scheduling]
[[dependencies]]
graph = "a"
[runtime]
[[root]]
[[[job submission]]]
method = qsub
[[[directives]]]
-I = bar=baz
-l = 'nodes=1,walltime=00:01:00'
-cwd = ''
```

and generate a job script to see the resulting directives:

```
$ cylc db reg test $HOME/test
$ cylc jobscontrol test a | grep QSUB
#QSUB -e /home/oliverh/cylc-run/pbs/log/job/a%1-1321046049.703576.err
#QSUB -l nodes=1,walltime=00:01:00
#QSUB -o /home/oliverh/cylc-run/pbs/log/job/a%1-1321046049.703576.out
#QSUB -N a%1
#QSUB -I bar=baz
#QSUB -cwd
```

10.5.2 Where To Put New Job Submission Modules

You new job submission class code should be saved to a file with the same name as the class (plus “.py” extension). It can reside in any of the following locations, depending on how generally useful the new method is and whether or not you have write-access to the cylc source tree:

- a `python` sub-directory of your suite definition directory.
- any directory in, or added to, your `PYTHONPATH` environment variable.
- in the `lib/cylc/job_submission` directory of the cylc source tree.

Note that the form of the import statement at the top of the new user-defined Python module differs depending on whether or not the file is installed in the cylc source tree (see the comment at the top of the example file above).

11 Running Suites

This section is still incomplete - please see also the Quick Start Guide (Section 6), command documentation (Section B), and play with the example suites.

11.1 Task States

As a suite runs its tasks (or rather the task proxies that represent the real tasks) move through a series of defined states:

- **waiting** - a *waiting* task's prerequisites have not yet been satisfied (clock-triggered tasks may also be waiting on their trigger time).
- **queued** - a *queued* task is ready to run but has been temporarily prevented from doing so by the internal cylc queue to which it is assigned (see *Internal Queues*, Section 11.2.2). Note this does mean the task is sitting in an external batch scheduler queue (that, depending on the chosen job submission method, would be the *submitted* state).
- **submitted** - a *submitted* task has been submitted by its defined job submission method but has not yet started executing (it may be waiting in an external batch scheduler queue, depending on job submission method).
- **running** - a *running* task has started executing (reported via a messaging call near the top of the task job script) but has not yet finished.
- two finished states:
 - **succeeded** - a *succeeded* task has finished successfully. It will remain in the suite so long as its outputs could be needed to satisfy the prerequisites of waiting tasks, and then it will be removed from the suite.
 - **failed** - a *failed* task has finished but reported failure. Failed tasks are not automatically removed from the suite, they must be handled by the suite operator (e.g. by retriggering after fixing the problem).

In addition, there are two held (paused) task states:

- **held** - a *held* task will not be submitted, even if it is ready to run, until it is released. Tasks that are already running cannot be held. The suite operator can deliberately hold tasks, and cylc itself puts tasks on hold under certain circumstances. For example, if the suite has been told to "shut down after currently running tasks have finished" waiting tasks will be put on hold while currently running tasks are finishing.
- **runahead** - a task in the *runahead* state has exceeded the suite runahead limit, i.e. it has got "too far ahead" of the slowest tasks and will be held back until they have caught up sufficiently. This state is automatically managed by cylc. See *The Suite Runahead Limit*, Section 11.2.1.

Finally, you may see references to the following pseudo states:

- **ready** - to set a task to *ready* means to set all of its prerequisites satisfied - i.e. ready to run (except clock-triggered tasks, which also wait on a trigger time).
- **base** (Graph Control GUI only) - graph nodes displayed in the off-white base color are displayed in order to show the proper graph structure but the corresponding task proxies are not actually present in the suite at the moment, because they have not been spawned yet, or they have been removed by cylc as spent tasks, or removed by the suite operator.

11.2 Limiting Suite Activity

Cylc suites are self-limiting to the extent that tasks are only submitted to run if they are actually ready to run. But even so some suites may generate too much activity at once, particularly in delayed/catch-up/historical operation when suites are not constrained by the clock-triggered tasks that normally have to wait on real time data. There are two issues to be aware of here: overburdening an external batch queueing system, or the task host hardware, by submitting too many tasks at once - this can be prevented by cylc's *internal queues*; and overburdening cylc itself by letting unconstrained or loosely constrained tasks run off far into the future (which would require cylc to keep vast numbers of "succeeded" tasks around until other tasks, which may depend on them, have caught up) - this can be prevented by the suite runahead limit.

11.2.1 The Suite Runahead Limit

In a cylc suite each task cycles independently constrained only by dependence on other tasks or by clock triggers. A cycling task with no prerequisites and no clock-trigger will tend to run off into the future (it is unlikely that you will need such a task but it is possible to define one). Similarly, in delayed or historical operation a clock-triggered task with no dependence on other tasks will run off ahead because the clock trigger provides no constraint until the task catches up to the wall clock. There is no advantage in letting a few fast data retrieval tasks, say, get way ahead of the slower tasks that actually use the data, and in large suites this could eventually swamp cylc because a large number of "succeeded" tasks would have to be maintained until the slower tasks, which might depend on them, have caught up.

To control this sort of behaviour cylc has a configurable *suite runahead limit* that prevents the fastest tasks getting too far ahead of the slowest tasks:

```
#SUITE.RC
[scheduling]
runahead_limit = 48 # (hours; default 24)
```

A cycling task spawns its successor when it enters the submitted state or, for sequential tasks, when it finishes. If the successor's cycle time is ahead of the oldest non-failed task by more than the runahead limit it is put into the special "runahead" held state until other tasks catch up sufficiently. In real time operation the runahead limit is of little consequence, because the suite will be constrained by its clock-triggered tasks, but the limit must be long enough to cover the minimum cycle time range present in the suite. A task that only runs once per day, for example, must be able to spawn 24 hours ahead.

Failed tasks, which are not automatically removed from a suite, are ignored when computing the runahead limit (but tasks that can't run because they depend on a failed task are not ignored, of course).

Note that it may occasionally be useful to define tasks that are entirely unconstrained even in real time operation. For instance, consider a tide prediction model that contributes, along with weather and sea state models, to a "seagram" product, but which does not wait on any real time data (tide is astronomically determined) and does not take inputs from other tasks in the suite. Such a task could be prevented from running off into the future by giving it artificial dependence on some other task, but the runahead limit will have the same effect.

11.2.2 Internal Queues

Large suites could potentially swamp the task host hardware or external batch queueing system, depending on the chosen job submission method, by submitting too many tasks at once. Cylc's internal queues prevent this by limiting the number of tasks, within defined groups, that are active (submitted or running) at once.

A queue is defined by a name; a *limit*, which is the maximum number of active tasks allowed for the queue; and a list of member tasks, which are assigned by name to the queue.

Queue configuration is done under the [scheduling] section of the suite.rc file, not as part of the runtime namespace hierarchy, because like dependencies queues constrain *when* a task runs rather than *what* runs after it is submitted. When runtime family relationships and queues do coincide you can assign task family members en masse to queues by using the family name, as shown in the example suite listing below.

By default every task is assigned to a *default* queue, which by default has a zero limit (interpreted by cylc as no limit). To use a single queue for the whole suite just set the default queue limit:

```
#SUITE.RC
[scheduling]
[[queues]]
    # limit the entire suite to 5 active tasks at once
    [[[default]]]
        limit = 5
```

To use other queues just name each one, set the limit, and assign member tasks:

```
#SUITE.RC
[scheduling]
[[queues]]
    [[[q_foo]]]
        limit = 5
        members = foo, bar, baz
```

Any tasks not assigned to a particular queue will remain in the default queue. The *queues* example suite illustrates how queues work by running two task trees side by side (as seen in the graph GUI) each limited to 2 and 3 tasks respectively:

```
title = demonstrates internal queueing
description = """
Two trees of tasks: the first uses the default queue set to a limit of
two active tasks at once; the second uses another queue limited to three
active tasks at once. Run via the graph control GUI for a clear view.

"""

[scheduling]
[[queues]]
    [[[default]]]
        limit = 2
    [[[foo]]]
        limit = 3
        members = n, o, p, fam2, u, v, w, x, y, z
[[dependencies]]
graph =
    a => b & c => fam1 => h & i & j & k & l & m
    n => o & p => fam2 => u & v & w & x & y & z
    """

[runtime]
[[fam1,fam2]]
[[d,e,f,g]]
    inherit = fam1
[[q,r,s,t]]
    inherit = fam2
```

Note assignment of runtime task family members to queues using the family name.

11.3 Other Topics In Brief

The following topics have yet to be documented in detail.

- The difference between cold-, warm-, raw-, and re-starting, a suite: see [cylc run help](#).
- Intervening in suites, e.g. stopping, removing, inserting tasks: see [cylc control help](#).

12 SUITE DESIGN PRINCIPLES

- Interrogating suites and tasks: see `cylc info help`, `cylc show help`, and `cylc discovery help`.
- Understanding suite evolution, particularly in delayed/catchup operation: the *Quick Start Guide* helps here, along with playing with example suites.
- Suite security, use of secure passphrases: this is trivial to configure, see documentation of the `use secure passphrase` configuration item in the *Suite.rc Reference*, Appendix A.
- Automatic state dump backups, named pre-intervention state dumps: these are used to reset a suite from a previous state of operation. They are mentioned in the *Quick Start Guide*. Watch the suite log after intervening in a suite, to get the filenames.
- Centralized alerting and timeouts: see documentation of `task event hooks` in the *Suite.rc Reference*, Appendix A.
- Recursive purge: this is a powerful intervention but you need to understand how it works before using it. See `cylc purge help` for details.
- Handling spin-up processes via temporary tasks and adding prerequisites on-the-fly: see `cylc depend --help`, and note that when you insert a task into a running suite (a) the initial cycle time can be in the past; and (b) you can give a final cycle time after which the task will be eliminated from the suite.
- Sub-suites: to run another suite inside a task, just invoke the subsuite, with appropriate start and end cycles (probably a single cycle), in the host task's command scripting:

```
[runtime]
  [[foo]]
    command scripting = \
      "cylc run SUITE ${CYLC_TASK_CYCLE_TIME} --until=${CYLC_TASK_CYCLE_TIME}"
```

12 Suite Design Principles

12.1 Make Fine-Grained Suites

A suite can contain a small number of large, internally complex tasks; a large number of small, simple tasks; or anything in between. CycL can easily handle a large number of tasks, however, so there are definite advantages to fine-graining:

- a more modular and transparent suite.
- better functional parallelism (multiple tasks running at the same time).
- faster debugging and failure recovery: rerun just the task(s) that failed.
- code reuse: similar tasks can often call the same script or command with differing task-specific input parameters (consider tasks that move files around, for example).

12.2 Make Tasks Rerunnable

It should be possible to rerun a task by simply resubmitting it for the same cycle time. In other words, failure at any point during execution of a task should not render a rerun impossible by corrupting the state of some internal-use file, or whatever. It's difficult to overstate the usefulness of being able to rerun the same task multiple times, either outside of the suite with `cylc submit`, or by retriggering it within the running suite, when debugging a problem.

12.3 Make Models Rerunnable

If a warm-cycled model simply overwrites its restart files in each run, the only cycle that can subsequently run is the next one. This is dangerous because if, accidentally or otherwise, the task runs for the wrong cycle time, its restart files will be corrupted such that the correct cycle

can no longer run (probably necessitating a cold-start). Instead, consider organising restart files by cycle time, through a file or directory naming convention, and keep them in a simple rolling archive (cylc's filename templating and housekeeping utilities can easily do this for you). Then, given availability of any external inputs, you can easily rerun the task for any cycle still in the restart archive.

12.4 Limit Previous-Instance Dependence

Cylc does not require that successive instances of the same task run sequentially. In order to task advantage of this and achieve maximum functional parallelism whenever the opportunity arises (usually when catching up from a delay) you should ensure that tasks that in principle do not depend on their own previous instances (the vast majority of tasks in most suites, in fact) do not do so in practice. In other words, they should be able to run as soon as their prerequisites are satisfied regardless of whether or not their predecessors have finished yet. This generally just means ensuring that all file I/O contains the generating task's cycle time in the file or directory name so that there is no interference between successive instances. If this is difficult to achieve in particular cases, however, you can declare the offending tasks to be *sequential*.

12.5 Put Task Cycle Time In All Output File Paths

Having all filenames, or perhaps the names of their containing directories, stamped with the cycle time of the generating task greatly aids in managing suite disk usage, both for archiving and cleanup. It also enables the aforementioned task rerunnability recommendation by avoiding overwrite of important files from one cycle to the next. Cylc has powerful utilities for cycle time offset based filename templating and housekeeping.

12.5.1 Use Cylc's Cycle Time Filename Template Utility

The command line utility program `cylc template` determines filenames based on a template string containing YYYYMMDDHH (or variations thereof) by substituting the current cycle time, or some offset from it, into the template. This can be used in the suite.rc environment sections, or in task implementation scripts if necessary, to instantly generate cycle time appropriate filenames for any purpose in the suite.

See `cylc util template help` for more information.

12.6 How To Manage Input/Output File Dependencies

Dependencies between tasks usually, though not always, take the form of files generated by one task that are used by other tasks. It is possible to manage these files across a suite without hard wiring I/O locations and therefore comprising suite flexibility and portability.

- **Use A Common I/O Workspace**

For small suites you may be able to have all tasks read and write from a common workspace, thereby avoiding the need to move common files around. You should be able to define the workspace location once in the suite.rc file rather than hard wiring it into the task implementations.

- **Add Connector Tasks To The Suite**

Tasks can be added to a suite to move files from A's output directory to B's input directory, and so on. These connector tasks may all be able to call the same file transfer script or command, with differing input parameters defined in the suite.rc file.

- **Dynamic Configuration Of I/O Paths**

Whether or not your suite uses a single common workspace, passing common I/O paths to tasks via variables defined once in the suite.rc file should allow you to avoid using connector tasks at all, except where it is necessary to transfer files between machines, or similar.

12.7 Use Generic Task Scripts

If your suite contains multiple logically distinct tasks that actually have similar functionality (e.g. for moving files around, or for generating similar products from the output of several similar models) have the corresponding cylc tasks all call the same command, script, or executable - just provide different input parameters via the task command scripting and/or execution environment, in the suite.rc file.

12.8 Make Suites Portable

If every task in a suite is configured to put its output under `$HOME` (i.e. the environment variable, literally, not the explicit path to your home directory; and similarly for temporary directories, etc.) then other users will be able to copy the suite and run it immediately, after merely ensuring that any external input files are in the right place.

For the ultimate in portability, construct suites in which all task I/O paths are dynamically configured to be user and suite (registration) specific, e.g.

```
$HOME/output/$CYLC_SUITE_REG_PATH
```

(these variables are automatically exported to the task execution environment by cylc - see *Task Execution Environment*, Section 8.4.4). Then you can run multiple instances of the suite at once (even under the same user account) without changing anything, and they will not interfere with each other.

You can test changes to a portable suite safely by making a quick copy of it in a temporary directory, then modifying and running the test copy without fear of corrupting the output directories, suite logs, and suite state, of the original.

12.9 Make Tasks As Self-Contained As Possible

Where possible, no task should rely on the action of another task, except for the prerequisites embodied in the suite dependency graph that it has no choice but to depend on. If this rule is followed, your suite will be as flexible as possible in terms of being able to run single tasks, or subsets of the suite, whilst debugging or developing new features.¹⁷ For example, every task should create its own output directories if they do not already exist, instead of assuming their existence due to the action of some another task; then you will be able to run single tasks without having to manually create output directories first.

```
# manual task scripting:
# 1/ create $OUTDIR if it doesn't already exist:
mkdir -p $OUTDIR
# 2/ create the parent directory of $OUTFILE if it doesn't exist:
mkdir -p $( dirname $OUTFILE )

# OR using the cylc checkvars utility:
# 1/ check vars are defined, and create directories if necessary:
cylc util checkvars -c OUTDIR1 OUTDIR2 ...
# 2/ check vars are defined, and create parent dirs if necessary:
cylc util checkvars -p OUTFILE1 OUTFILE2 ...
```

¹⁷The `cylc submit` command runs a single task exactly as its suite would, in terms of both job submission method and execution environment.

12.10 Make Suites As Self-Contained As Possible

The only compulsory content of a cylc suite definition directory is the suite.rc file (and you'll almost certainly have a suite `bin` sub-directory too). However, you can store whatever you like in a suite definition directory;¹⁸ other files there will be ignored by cylc but suite tasks can access them via the `$CYLC_SUITE_DEF_PATH` variable that cylc automatically exports into the task execution environment. Disk space is cheap - if all programs, ancillary files, control files (etc.) required by the suite are stored in the suite definition directory instead of having the suite reference external build directories (etc.), you can turn the directory into a revision control repository and be virtually assured of the ability to exactly reproduce earlier versions, regardless of suite complexity.

12.11 Orderly Product Generation?

Correct scheduling is not equivalent to “orderly generation of products by cycle time”. Under cylc, a product generation task will trigger as soon as its prerequisites are satisfied (i.e. when its input files are ready, generally) regardless of whether other tasks with the same cycle time have finished or have yet to run. If your product delivery or presentation system demands that all products for one cycle time are uploaded (or whatever) before any from the next cycle, then be aware that this may be quite inefficient if your suite is ever faced with catching up from a significant delay or running over historical data.

If you must, however, you can introduce artificial dependencies into your suite to ensure that the final products never arrive out of sequence. One way of doing this would be to have a final “product upload” task that depends on completion of all the real product generation tasks at the same cycle time, and then declare it to be sequential.

12.12 Clock-triggered Tasks Wait On External Data

All tasks in a cylc suite know their own private cycle time, but most don't care about the wall clock time - they just run when their prerequisites are satisfied. The exception to this is *clock-triggered* tasks, which wait on a wall clock time expressed as an offset from their own cycle time, in addition to any other prerequisites. The usual purpose of these tasks is to retrieve real time data from the external world, triggering at roughly the expected time of availability of the data. Triggering the task at the right time is up to cylc, but the task itself should go into a check-and-wait loop in case the data is delayed; only on successful detection or retrieval should the task report success and then exit (or perhaps report failure and then exit if the data has not arrived by some cutoff time).

12.13 Do Not Treat Real Time Operation As Special

Cylc suites, without modification, can handle real time and delayed operation equally well.

In real time operation clock-triggered tasks constrain the behaviour of the whole suite, or at least of all tasks downstream of them in the dependency graph.

In delayed operation (whether due to an actual delay in an operational suite or because you're running an historical trial) clock-triggered tasks will not constrain the suite at all, and cylc's cycle interleaving abilities come to the fore, because their trigger times have already passed. But if a clock-triggered task happens to catch up to the wall clock, it will automatically wait again. In this way a cylc suite naturally and seamlessly transitions between delayed and real time operation as required.

¹⁸If you copy a suite using cylc commands or gcycle, the entire suite definition directory will be copied.

A Suite.rc Reference

This appendix documents the legal content of raw cylc suite.rc files. Most suites will only need to explicitly configure a few of these items - many have sensible default values, some may only be needed for critical operational suites (e.g. secure passphrases and lock server settings), and some are primarily used for cylc development.

In addition to the raw syntax, Jinja2 expressions can also be embedded to programmatically generate the final suite definition seen by cylc. Use of Jinja2 is documented in Section [8.6](#).

See also *Suite Definition - Suite.rc Overview* (Section [8.2](#)) for a descriptive overview of suite.rc files.

A.1 Top Level Items

The only top level configuration items at present are the suite title and description.

A.1.1 title

The suite title is displayed in the g cylc suite database window. It can also be retrieved from a suite at run time with `cylc show` (or use `cylc get-config`).

- *type*: string
- *default*: “No title provided”

A.1.2 description

The suite description can be retrieved by g cylc right-click menu. It can also be retrieved from a suite at run time with `cylc show` (or use `cylc get-config`).

- *type*: string
- *default*: “No description provided”

A.2 [cylc]

This section is for suite configuration that is not specifically task-related.

A.2.1 [cylc] → UTC mode

Cylc runs off the suite host’s system clock by default. This item allows you to run the suite in UTC even if the system clock is set to local time. Clock-triggered tasks will trigger when the current UTC time is equal to their cycle time plus offset; other time values used, reported, or logged by cylc will also be in UTC.

- *type*: boolean
- *default*: False

A.2.2 [cylc] → simulation mode only

This prevents a suite from running in real mode - use for demo suites created by copying real suites out of their normal operating environment.

- *type*: boolean
- *default*: False

A.2.3 [cylc] → use secure passphrase

Critical operational suites can be made to ignore commands unless the originating user account has a special passphrase (the same one used by the suite owner at startup) with secure permissions (as for ssh keys) in the file `$HOME/.cylc/security/$CYLC_SUITE_REG_NAME` (remotely hosted tasks also need the passphrase in their host accounts). The passphrase itself is never transferred across the network (a secure MD5 checksum is). Note that cylc's normal owner-only suite access is implemented by means of a simple username comparison, which could in principle be subverted by a malicious user copying and modifying the cylc code base. Secure passphrases, on the other hand, should guarantee suite security so long as your user accounts aren't breached.

- *type*: boolean
- *default*: False

A.2.4 [cylc] → [[logging]]

This section configures cylc's logging functionality, which records time-stamped events to a special log file.

A.2.4.1 [cylc] → [[logging]] → directory

The cylc log and its backups are stored in this directory. If you change the directory make sure it remains suite-specific by using suite identity environment variables in the path.

- *type*: string (directory path, may contain environment variables)
- *default*: `$HOME/cylc-run/$CYLC_SUITE_REG_NAME/log/suite`

A.2.4.2 [cylc] → [[logging]] → roll over at start-up

Suite logs roll over (start anew) automatically when they reach a certain size - currently hard-wired to 1MB in `$CYLC_DIR/lib/cylc/pimp_my_logger.py`. They can also be rolled automatically whenever a suite is started or restarted.

- *type*: boolean
- *default*: True

A.2.5 [cylc] → [[state dumps]]

State dump files allow cylc to restart suites from previous states of operation.

A.2.5.1 [cylc] → [[state dumps]] → directory

The rolling archive of suite state dump files, backups of the default state dump, and any special pre-intervention state dumps, are stored under this directory. If you change this directory make sure it remains suite-specific by using suite identity environment variables in the path.

- *type*: string (directory path, may contain environment variables)
- *default*: `$HOME/cylc-run/$CYLC_SUITE_REG_NAME/state`

A.2.5.2 [cylc] → [[state dumps]] → number of backups

This is the length, in number of changes, of the automatic rolling archive of state dump files that allows you to restart a suite from a previous state. Every time a task changes state cylc updates the state dump and rolls previous states back one on the archive. You'll probably only ever need the latest (most recent) state dump, which is automatically used in a restart, but any previous state still in the archive can be used. Additionally, special labeled state dumps are written out prior to actioning any suite intervention - their filenames are logged by cylc.

- *type*: integer (≥ 1)
- *default*: 10

A.2.6 [cylc] → [[event hooks]]

Cylc can call a nominated “event handler” script when certain suite events occur. This is intended to facilitate centralized alerting and automated handling of critical events. Event handlers can do anything you like, such as send emails or SMS, call pagers, or intervene in the operation of their parent suite with cylc commands. The command `cylc [hook] email-suite` is a ready made suite event handler.

A single script can be nominated to handle a list of chosen events, and the event name is passed to the handler to allow it to distinguish between them.

Custom suite event handlers can be located in the suite bin directory. They are passed the following arguments by cylc:

```
<hook-script> EVENT SUITE MESSAGE
```

Currently there is only one suite EVENT defined:

- ‘shutdown’ - the suite stopped, normally or not.

MESSAGE, if provided, describes what has happened.

Note that event handlers are called by cylc itself so if you wish to pass them additional information via the environment you must use [cylc] → [[environment]], not task runtime environments.

See also Section A.4.1.10 for task event hooks.

A.2.6.1 [cylc] → [[event hooks]] → events

The list of events to handle, as documented for EVENT just above.

- *type*: list of strings
- *default*: (none)

A.2.6.2 [cylc] → [[event hooks]] → script

The handler script to call when one of the nominated events occurs.

- *type*: string
- *default*: (none)

A.2.7 [cylc] → [[lockserver]]

The cylc lockserver brokers suite and task locks on the network (these are somewhat analogous to traditional local *lock files*). It prevents multiple instances of a suite or task from being invoked at the same time (via scheduler instances or `cylc submit`).

See `cylc lockserver --help` for how to run the lockserver, and `cylc lockclient --help` for occasional manual lock management requirements.

A.2.7.1 [cylc] → [[lockserver]] → enable

The lockserver is currently disabled by default. It is intended mainly for operational use.

- *type*: boolean
- *default*: False

A.2.7.2 [cylc] → [[lockserver]] → simultaneous instances

By default the lockserver prevents multiple simultaneous instances of a suite from running even under different registered names. But allowing this may be desirable if the I/O paths of every task in the suite are dynamically configured to be suite specific (and similarly for the suite state dump and logging directories, by using suite identity variables in their directory paths). Note that *the lockserver cannot protect you from running multiple distinct copies of a suite simultaneously*.

- *type*: boolean
- *default*: False

A.2.8 [cylc] → [[environment]]

Variables defined here are exported into the environment in which cylc itself runs. They are then available to *local processes spawned directly by cylc*. Any variables read by task event handlers must be defined here, for instance, because event handlers are executed directly by cylc, not by running tasks. And similarly the command lines issued by cylc to invoke event handlers or to submit task job scripts could, in principle, make use of environment variables defined here.

A.2.8.1 Warnings

- Cylc local variables are not available to executing tasks unless you happen to choose a local direct job submission method. *Always use task runtime environments to pass variables into the task execution environment* ([A.4.1.11.1](#)).
- Unlike task execution environment variables, which are written to job scripts and interpreted by the shell at run time, cylc local environment variables are written directly to the environment by Python as literal strings *so shell variable expansion expressions cannot be used here*.

A.2.8.2 [cylc] → [[environment]] → __VARIABLE__

Replace `__VARIABLE__` with any number of environment variable assignment expressions. Values may refer to other local environment variables (order of definition is preserved) and are not evaluated or manipulated by cylc, so any variable assignment expression that is legal in the shell in which cylc is running can be used (but see the warning above on variable expansions, which will not be evaluated). White space around the '=' is allowed (as far as cylc's suite.rc parser is concerned these are normal configuration items).

- *type*: string
- *default*: (none)
- *examples*:
 - `FOO = $HOME/foo`

A.2.9 [cylc] → [[simulation mode]]

Items specific to running suites in simulation mode.

A.2.9.1 [cylc] → [[simulation mode]] → clock rate

This determines the speed at which the simulation mode clock runs, in real seconds per simulated hour. A value of 10, for example, means it will take 10 real seconds to simulate one hour of operation.

- *type*: integer (≥ 0 , real seconds per simulated hour)
- *default*: 10

A.2.9.2 [cylc] → [[simulation mode]] → clock offset

The clock offset determines the initial time on the simulation clock, at suite startup, relative to the initial cycle time. An offset of 0 simulates real time operation; greater offsets simulate catch up from a delay and subsequent transition to real time operation.

- *type*: integer (≥ 0 , hours behind initial cycle time)
- *default*: 24

A.2.9.3 [cylc] → [[simulation mode]] → command scripting

The command scripting to execute for all tasks when running in simulation mode.

- *type*: string (scripting valid in job submission shell; triple quote for multiple lines)
- *default*: `echo SIMULATION MODE $CYLC_TASK_ID; sleep 10; echo BYE`

A.2.9.4 [cylc] → [[simulation mode]] → [[[job submission]]]

Configure job submission for simulation mode.

A.2.9.4.1 [cylc] → [[simulation mode]] → [[[job submission]]] → method

The job submission method to use for all tasks in simulation mode. Any available method can be used but the default is probably sufficient for simulation mode.

- *type*: string (a job submission method name - see Section A.4.1.8.1)
- *default*: `background`

A.2.9.5 [cylc] → [[simulation mode]] → [[[event hooks]]]

Configure event hooks for simulation mode.

A.2.9.5.1 [cylc] → [[simulation mode]] → [[[event hooks]]] → enable

Event hooks are disabled by default in simulation mode. They can be enabled in order to test automated alerts, for example, without running the real suite tasks, but be aware that timeouts will be relative to the accelerated simulation mode clock which by default runs very quickly.

- *type*: boolean
- *default*: False

A.3 [scheduling]

This section allows cylc to determine when tasks are ready to run.

A.3.1 [scheduling] → initial cycle time

At startup each cycling task (unless specifically excluded under [special tasks]) will be inserted into the suite with this cycle time, or with the closest subsequent valid cycle time for the task. Note that whether or not *cold-start tasks*, specified under [special tasks], are inserted, and in what state they are inserted, depends on the start up method - cold, warm, or raw. If this item is provided you can override it on the command line or in the gcylc suite start panel.

- *type*: integer (YYYYMMDDHH)
- *default*: (none)

A.3.2 [scheduling] → final cycle time

Cycling tasks are held (i.e. not allowed to spawn a successor) once they pass the final cycle time, if one is specified. Once all tasks have achieved this state the suite will shut down. If this item is provided you can override it on the command line or in the gcylc suite start panel.

- *type*: integer (YYYYMMDDHH)
- *default*: (none)

A.3.3 [scheduling] → runahead limit

The suite runahead limit prevents the fastest tasks in a suite from getting too far ahead of the slowest ones, as documented in Section 11.2.1. Tasks exceeding the limit will be put into a special runahead held state until slower tasks have caught up sufficiently. The limit must be long enough to cover the minimum cycle time range of tasks present in the suite: a task that only runs once per day, for instance, needs to spawn 24 hours ahead. Failed tasks, which are not automatically removed from a suite, are ignored when computing the runahead limit.

- *type*: integer (≥ 0 , hours)
- *default*: 24

A.3.4 [scheduling] → [[queues]]

Configuration of internal queues, by which the number of simultaneously active tasks (submitted or running) can be limited, per queue. By default a single queue called *default* is defined, with all tasks assigned to it and no limit. To use a single queue for the whole suite just set the limit on the *default* queue as required. See also Section 11.2.2.

A.3.4.1 [scheduling] → [[queues]] → [[[QUEUE]]]

Section heading for configuration of a single queue. Replace `QUEUE` with a queue name, and repeat the section as required.

- *type*: string
- *default*: “default”

A.3.4.2 [scheduling] → [[queues]] → [[[QUEUE]]] → limit

The maximum number of active tasks allowed at any one time, for this queue.

- *type*: integer
- *default*: 0 (i.e. no limit)

A.3.4.3 [scheduling] → [[queues]] → [[[QUEUE]]] → members

A list of member tasks, or task family names, to assign to this queue (assigned tasks will automatically be removed from the default queue).

- *type*: list of strings
- *default*: none for user-defined queues; all tasks for the “default” queue

A.3.5 [scheduling] → [[special tasks]]

This section identifies any tasks with special behaviour. By default (i.e. non “special” behaviour) tasks submit (or queue) as soon as their prerequisites are satisfied, and they spawn a successor at the next valid cycle time for the task as soon as they enter the submitted state¹⁹

A.3.5.1 [scheduling] → [[special tasks]] → clock-triggered

Clock-triggered tasks wait on a wall clock time specified as an offset *in hours* relative to their own cycle time, in addition to any dependence they have on other tasks. Generally speaking, only tasks that wait on external real time data need to be clock-triggered. Note that in computing the trigger time the full wall clock time and cycle time are compared, not just hours and minutes of the day, so when running a suite in catchup/delayed operation, or over historical periods, clock-triggered tasks will not constrain the suite at all until they catch up to the wall clock.

- *type*: list of tasknames with offsets (hours, positive or negative)
- *default*: (none)
- *example*: `clock-triggered = foo(1.5), bar(2.25)`

A.3.5.2 [scheduling] → [[special tasks]] → start-up

Start-up tasks are one-off tasks that are only used when a suite is cold started (i.e. starting up without assuming any previous cycle). They can be used to clean out or prepare a suite workspace, for example, before other tasks run.

- *type*: list of task names
- *default*: (none)

¹⁹Spawning any earlier than this brings no advantage in terms of functional parallelism and would cause uncontrolled proliferation of waiting tasks.

A.3.5.3 [scheduling] → [[special tasks]] → cold-start

A cold-start task (or possibly a sequence of them) is used to satisfy the dependence of an associated task with the same cycle time, on outputs from a previous cycle - when those outputs are not available. The primary use for this is to cold-start a warm-cycled forecast model that normally depends on restart files (e.g. model background fields) generated by its previous forecast, when there is no previous forecast. This is required when cold-starting the suite, but cold-start tasks can also be inserted into a running suite to restart a model that has had to skip some cycles after running into problems (e.g. critical inputs not available). Cold-start tasks can invoke real cold-start processes, or they can just be dummy tasks that represent some external process that has to be completed before the suite is started.

- *type*: list of task names
- *default*: (none)

A.3.5.4 [scheduling] → [[special tasks]] → sequential

By default, a task spawns a successor as soon as it is submitted to run so that successive instances of the same task can run in parallel if the opportunity arises (i.e. if their prerequisites happen to be satisfied before their predecessor has finished). *Sequential tasks*, however, will not spawn a successor until they have finished successfully. This should be used for (a) *tasks that cannot run in parallel with their own previous instances* because they would somehow interfere with each other (use cycle time in all I/O paths to avoid this); and (b) *warm cycled forecast models that write out restart files for multiple cycles ahead* (exception: see “explicit restart outputs” below).²⁰

- *type*: list of task names
- *default*: (none)

A.3.5.5 [scheduling] → [[special tasks]] → one-off

One-off tasks do not spawn a successor; they run once and are then removed from the suite when they are no longer needed. *Start-up* and *cold-start* tasks are automatically one-off tasks and do not need to be listed here.

- *type*: list of task names
- *default*: (none)

A.3.5.6 [scheduling] → [[special tasks]] → explicit restart outputs

This is only required in the event that you need a warm cycled forecast model to start at the instant its restart files are ready (if other prerequisites are satisfied) *even if its previous instance has not finished yet*. If so, the model task has to depend on special output messages emitted by the previous instance as soon as its restart files are ready, instead of just on the previous instance finishing. *Tasks in this category must define special restart output messages containing the word “restart”, in [runtime] → [[TASK]] → [[[outputs]]]* - see Section 9.4.2.

- *type*: list of task names
- *default*: (none)

²⁰This is because you don't want Model[T] waiting around to trigger off Model[T-12] if Model[T-6] has not finished yet. If Model is forced to be sequential this can't happen because Model[T] won't exist in the suite until Model[T-6] has finished. But if Model[T-6] fails, it can be spawned-and-removed from the suite so that Model[T] can then trigger off Model[T-12], which is the correct behaviour.

A.3.5.7 [scheduling] → [[special tasks]] → exclude at start-up

Any task listed here will be excluded from the initial task pool (this goes for suite restarts too). If an *inclusion* list is also specified, the initial pool will contain only included tasks that have not been excluded. Excluded tasks can still be inserted at run time. Other tasks may still depend on excluded tasks if they have not been removed from the suite dependency graph, in which case some manual triggering, or insertion of excluded tasks, may be required.

- *type*: list of task names
- *default*: (none)

A.3.5.8 [scheduling] → [[special tasks]] → include at start-up

If this list is not empty, any task *not* listed in it will be excluded from the initial task pool (this goes for suite restarts too). If an *exclusion* list is also specified, the initial pool will contain only included tasks that have not been excluded. Excluded tasks can still be inserted at run time. Other tasks may still depend on excluded tasks if they have not been removed from the suite dependency graph, in which case some manual triggering, or insertion of excluded tasks, may be required.

- *type*: list of task names
- *default*: (none)

A.3.6 [scheduling] → [[dependencies]]

The suite dependency graph is defined under this section. You can plot the dependency graph as you work on it, with `cylc graph` or by right clicking on the suite in cylc. See also Section 8.3.

A.3.6.1 [scheduling] → [[dependencies]] → graph

The dependency graph for any one-off asynchronous (non-cycling) tasks in the suite goes here. This can be used to construct a suite of one-off tasks (e.g. build jobs and related processing) that just completes and then exits, or an initial suite section that completes prior to the cycling tasks starting (if you make the first cycling tasks depend on the last one-off ones). But note that synchronous *start-up* tasks can also be used for the latter purpose. See Section A.3.6.2.1 below for graph string syntax, and Section 8.3.

- *type*: string
- *example*: (see Section A.3.6.2.1 below)

A.3.6.2 [scheduling] → [[dependencies]] → [[[VALIDITY]]]

Replace `VALIDITY` with a comma-separated list of integer hours, $0 \leq H \leq 23$, defining the valid cycle times for the subsequent graph of cycling tasks, or multiple such subsections as required for different dependencies at different hours; or with ‘`ASYNCCID:pattern`’, where *pattern* is a regular expression that matches an asynchronous task ID, for each graph of repeating asynchronous tasks.

- *examples*:
 - (cycling tasks) `[[0,6,12,18]]`
 - (repeating asynchronous tasks) `[[ASYNCCID:SAT-\d+]]`
- *default*: (none)

For a cycling graph with multiple validity sections for different hours of the day, the different sections *add* to generate the complete graph. Different graph sections can overlap (i.e. the same hours may appear in multiple section headings) and the same tasks may appear in multiple sections, but the individual dependencies must be unique across the entire graph. Duplicate dependencies will cause a run time error at start-up, as documented in *How Multiple Graph Strings Combine* (Section 8.3.3.2.1).

A.3.6.2.1 [scheduling] → [[dependencies]] → [[[VALIDITY]]] → graph

The dependency graph for the specified validity section (described just above) goes here. Syntax examples follow; see also Section 8.3.

- *type*: string
- *example*:

```
graph = """
    foo => bar => baz & waz    # baz and waz both trigger off bar
    baz:out1 => faz            # faz triggers off an internal output of baz
    ColdFoo | foo(T-6) => foo # cold-start or restart for foo
    X:fail => Y              # Y triggers if X fails
    X => !Y                  # Y suicides if X succeeds
    X | X:fail => Z          # Z triggers if X finishes or fails
    (A | B & C ) | D => foo  # general conditional triggers
    # comment
    """
```

- *default*: (none)

A.3.6.2.2 [scheduling] → [[dependencies]] → [[[VALIDITY]]] → daemon

For [[[ASYNCID:pattern]]] validity sections only, list *asynchronous daemon* tasks by name. This item is located here rather than under [scheduling] → [[special tasks]] because a damon task is associated with a particular asynchronous ID.

- *type*: list of task names
- *default*: (none)

A.4 [runtime]

This section defines how, where, and what to execute when tasks are ready to run. Runtime subsections define an inheritance hierarchy of *namespaces*, each of which represents a family of tasks or an individual task, as described in Section 8.4).

A.4.1 [runtime] → [[NAME]]

Replace *NAME* with a namespace name, or a comma separated list of names, and repeat as needed to define all tasks in the suite.

Namespace names may contain letters, digits, underscores, and hyphens. They may not contain colons (which would preclude use of directory paths involving the registration name in *SPATH* variables). They may not contain the ‘.’ character (it will be interpreted as the namespace hierarchy delimiter, separating groups and names). *Task names should not be hardwired into task implementations*. Rather, task and suite identity should be extracted portably from the task execution environment supplied by cylc (Section 8.4.4) - then to rename a task you can just change its name in the suite.rc file.

A namespace represents a family if other namespaces inherit from it; it represents a task if it is a leaf on the inheritance tree (i.e. no other namespaces inherit from it).

- *legal values:*
 - [[`foo`]]
 - [[`foo, bar, baz`]]

If multiple names are listed, the subsequent namespace configuration items apply to each member, but any instance of `<TASK>` in any value will be replaced by cylc with the actual namespace name. This can be used to define many tasks that are almost identical.

All namespaces inherit initially from `root`, which you can explicitly configure to override or provide default runtime settings for all tasks in the suite.

A.4.1.1 [runtime] → [[_NAME_]] → **inherit**

Specify here the namespace from which this namespace should inherit all of its runtime configuration; specific items can then be overridden as required. Note that many of the available items are left undefined even in the root namespace.

- *type:* string (another namespace name)
- *default:* root

A.4.1.2 [runtime] → [[_NAME_]] → **description**

A description of this namespace, retrievable from running tasks via `cylc show`.

- *type:* string
- *root default:* “No description provided”

A.4.1.3 [runtime] → [[_NAME_]] → **initial scripting**

Scripting defined here is put at the top of the task job script, before the task execution environment is configured, and before configuring task access to cylc. This can be used when there is a need for remotely hosted tasks to source login scripts - when Pyro has been installed locally rather than at system level, for example - because ssh does not do this for you. Use of initial scripting is not restricted to remote tasks, but pre-command or command scripting, which has access to the task execution environment, should normally be used in preference to this item.

- *type:* string
- *default:* (none)
- *example:* `initial scripting = ". $HOME/.profile"`

A.4.1.4 [runtime] → [[_NAME_]] → **command scripting**

The scripting to execute when the associated task is ready to run. The value can be a single command, or a multiline string of scripting, or a list of commands or multiline strings. If a list is provided the task will automatically resubmit with the second list member if the task fails, and then the third if the task fails again, and so on until no command scripting list members remain - this is one of the ways you can do automated failure recovery in cylc.

- *type:* (a list of) strings (each)
- *root default:* `echo DUMMY $CYLC_TASK_ID; sleep 10; echo BYE`

A.4.1.5 [runtime] → [[_NAME_]] → pre-command scripting

Scripting to be executed immediately *before* the command scripting. This would typically be used to add scripting to every task in a family (for individual tasks you could just incorporate the extra commands into the main command scripting). See also *post-command scripting*, below.

- *type*: string, or a list of strings
- *default*: (none)
- *example*:

```
pre-command scripting = """
    . $HOME/.profile
    echo Hello from suite ${CYLC_SUITE_REG_NAME} !"""
```

A.4.1.6 [runtime] → [[_NAME_]] → post-command scripting

Scripting to be executed immediately *after* the command scripting. This would typically be used to add scripting to every task in a family (for individual tasks you could just incorporate the extra commands into the main command scripting). See also *pre-command scripting*, above.

- *type*: string, or a list of strings
- *default*: (none)

A.4.1.7 [runtime] → [[_NAME_]] → manual completion

If a task's initiating process detaches and exits before task processing is finished then cylc cannot arrange for the task to automatically signal when it has succeeded or failed. In such cases you must use this configuration item to tell cylc not to arrange for automatic completion messaging, and insert some minimal completion messaging yourself in appropriate places in the task implementation (see Section 9.4.4).

- *type*: boolean
- *default*: False

A.4.1.8 [runtime] → [[_NAME_]] → [[[job submission]]]

This is where to configure the means by which cylc submits task job scripts to run.

A.4.1.8.1 [runtime] → [[_NAME_]] → [[[job submission]]] → method

See *Task Job Submission* (Section 10) for how job submission works, and how to define new methods. Cylc has a number of built in job submission methods:

- *type*: string
- *legal values*:
 - *background* - direct background execution
 - *at* - the rudimentary Unix *at* scheduler
 - *loadleveler* - *llsubmit*, with directives defined in the suite.rc file
 - *pbs* - PBS *qsub*, with directives defined in the suite.rc file
 - *sge* - Sun Grid Engine *qsub*, with directives defined in the suite.rc file
- *default*: *background*

A.4.1.8.2 [runtime] → [[_NAME_]] → [[[job submission]]] → command template

This allows you to override the actual command used by the chosen job submission method. The template's first %s will be substituted by the job file path. Where applicable the second and third %s will be substituted by the paths to the job stdout and stderr files.

- *type*: string
- *legal values*: a string template
- *example*: `llsubmit %s`

A.4.1.8.3 [runtime] → [[_NAME_]] → [[[job submission]]] → shell

This is the shell used to interpret the job script submitted by cylc when a task is ready to run. *It has no bearing on the shell used in task implementations.* Command scripting and suite environment variable assignment expressions must be valid for this shell. The latter is currently hardwired into cylc as `export item=value` - valid for both bash and ksh because `value` is entirely user-defined - but cylc would have to be modified slightly to allow use of the C shell.

- *type*: string
- *root default*: `/bin/bash`

A.4.1.8.4 [runtime] → [[_NAME_]] → [[[job submission]]] → log directory

This is where task job scripts, and the stdout and stderr logs for local tasks, are written. The directory path may contain environment variables, including suite identity variables to make the path suite-specific (as the default value does). The job script filename is constructed, just before job submission, from the task ID and *seconds since epoch*, and then *.out* and *.err* are appended to construct the stdout and stderr log names, respectively. These filenames are thus unique even if a task gets retriggered, and yet will be correctly time ordered if the log directory is listed. The filenames are also recorded by each task proxy for access via cylc commands and the suite control GUIs.

- *type*: string (directory path, may contain environment variables)
- *default*: `$HOME/cylc-run/$CYLC_SUITE_REG_NAME/log/job`

A.4.1.8.5 [runtime] → [[_NAME_]] → [[[job submission]]] → work directory

Each task runs in its own working directory. The work directory is created on the fly by the task job script prior to executing the task command scripting. For non-detaching tasks the work directory is removed, if it is empty, before the job script exits. It can be accessed by tasks at run time via the environment variable `$CYLC_TASK_WORK_PATH`.

- *type*: string (directory path, may contain environment variables)
- *default*: `$CYLC_SUITE_DEF_PATH/work/$CYLC_TASK_ID`

A.4.1.8.6 [runtime] → [[_NAME_]] → [[[job submission]]] → share directory

This is intended as a shared data directory for multiple tasks. As such it should be defined at root level, if for all tasks; or in a family namespace, if for a subset of tasks; however, it can be made unique for every task if you like. The share directory is created on the fly by the task job script prior to executing the task command scripting. It can be accessed by tasks at run time via the environment variable `$CYLC_SUITE_SHARE_PATH`.

- *type*: string (directory path, may contain environment variables)
- *default*: `$CYLC_SUITE_DEF_PATH/share`

A.4.1.9 [runtime] → [[_NAME_]] → [[[remote]]]

Remote hosting configuration for tasks that need to run on platforms other than the suite host. If a remote host is specified cylc will attempt to execute the task on that host by passwordless ssh. Relevant task scripts and executables, cylc itself, and possibly all or part of the suite definition directory, must be installed on the remote host. Passwordless ssh must be configured between the local suite owner and remote task owner accounts.

A.4.1.9.1 [runtime] → [[_NAME_]] → [[[remote]]] → host

The remote host for this task or family.

- *type*: string (a valid hostname on the network)
- *default*: (none)

A.4.1.9.2 [runtime] → [[_NAME_]] → [[[remote]]] → owner

The task owner username. This is (only) used in the passwordless ssh command line invoked by cylc to submit the remote task (consequently it may be defined using local environment variables (i.e. the shell in which cylc runs, and [cylc] → [[environment]]).

- *type*: string (a valid username on the remote host)
- *default*: (none)

A.4.1.9.3 [runtime] → [[_NAME_]] → [[[remote]]] → cylc directory

The path to the remote cylc installation, required if cylc is not in the default search path on the remote host.

- *type*: string (a valid directory path on the remote host)
- *default*: (none)

A.4.1.9.4 [runtime] → [[_NAME_]] → [[[remote]]] → suite definition directory

The path to the suite definition directory on the remote host, needed if remote tasks require access to files stored there (via `$CYLC_SUITE_DEF_PATH`) or in the suite bin directory (via `$PATH`). If this item is not defined, the local suite definition directory path will be assumed, with the suite owner's home directory, if present, replaced by '`$HOME`' for interpretation on the remote host.

- *type*: string (a valid directory path on the remote host)
- *default*: (local suite definition path with `$HOME` replaced)

A.4.1.9.5 [runtime] → [[_NAME_]] → [[[remote]]] → remote shell template

A template for the remote shell command for a submitting a remote task. The template's first %s will be substituted by the remote user@host.

- *type*: string (a string template)
- *root default*: `ssh -oBatchMode=yes %s`

A.4.1.9.6 [runtime] → [[_NAME_]] → [[[remote]]] → log directory

This log directory is used for the stdout and stderr logs of remote tasks. The directory will be created on the fly if necessary. If not specified, the local log path will be used (see [A.4.1.8.4](#)) with the suite owner's home directory path, if present, replaced by '`$HOME`' for interpretation on the remote host. The stdout and stderr log file names are the same as for local tasks, and are recorded for access via `gcylc`.

- *type*: string (a valid directory path on the remote host)
- *default*: (local log path with `$HOME` replaced)

A.4.1.9.7 [runtime] → [[_NAME_]] → [[[remote]]] → work directory

Use this item if you need to override the local task work directory (see [A.4.1.8.5](#)). If omitted, the local directory will be used with the suite owner's home directory path, if present, replaced by '`$HOME`' for interpretation on the remote host.

- *type*: string (directory path, may contain environment variables)
- *default*: (local task work path with `$HOME` replaced)

A.4.1.9.8 [runtime] → [[_NAME_]] → [[[remote]]] → share directory

Use this item if you need to override the local share directory (see [A.4.1.8.6](#)). If omitted, the local directory will be used with the suite owner's home directory path, if present, replaced by '`$HOME`' for interpretation on the remote host.

- *type*: string (directory path, may contain environment variables)
- *default*: (local task share path with `$HOME` replaced)

A.4.1.10 [runtime] → [[_NAME_]] → [[[event hooks]]]

Cylc can call a nominated “event handler” script when certain task events, such as task failure or timeout, occur. This is intended to facilitate centralized alerting and automated handling of critical events. Event handlers can do anything you like, such as send emails or SMS, call pagers, or intervene in the operation of their parent suite with cylc commands. The command `cylc [hook] email-task` is a ready made task event handler.

Currently a single script can be nominated to handle a list of chosen events, and the event name is passed to the handler to allow it to distinguish between them (cylc once allowed distinct handlers for each event but this required verbose configuration for little gain).

Custom task event handlers can be located in the suite bin directory. They are passed the following arguments by cylc:

```
<hook-script> EVENT SUITE TASKID MESSAGE
```

where EVENT is one of the following strings:

- ‘submitted’ - the task was submitted
- ‘started’ - the task started running
- ‘succeeded’ - the task succeeded
- ‘submission_failed’ - the task failed in job submission
- ‘warning’ - the task reported a warning message
- ‘failed’ - the task failed
- ‘submission_timeout’ - task job submission timed out

- ‘execution_timeout’ - task execution timed out

MESSAGE, if provided, describes what has happened, and TASKID identifies the task (`NAME%CYCLE` for cycling tasks).

To configure timeout handling, list ‘submission_timeout’ and/or ‘execution_timeout’ in the events to handle, and then specify corresponding timeout values in minutes using the timeout config items below.

Note that event handlers are called by cylc itself, not by the running tasks so if you wish to pass them additional information via the environment you must use `[cylc] → [[environment]]`, not task runtime environments.

See also Section [A.2.6](#) for suite event hooks.

A.4.1.10.1 [runtime] → [[_NAME_]] → [[[event hooks]]] → events

The list of events to handle, as documented for EVENT just above.

- *type*: list of strings
- *default*: (none)

A.4.1.10.2 [runtime] → [[_NAME_]] → [[[event hooks]]] → script

The handler script to call when one of the nominated events occurs.

- *type*: string
- *default*: (none)

A.4.1.10.3 [runtime] → [[_NAME_]] → [[[event hooks]]] → submission timeout

If a task has not started the specified number of minutes after it was submitted, the event handler will be called by cylc with *submission_timeout* as the EVENT argument:

- *type*: float (minutes)
- *default*: (none)

A.4.1.10.4 [runtime] → [[_NAME_]] → [[[event hooks]]] → execution timeout

If a task has not finished the specified number of minutes after it started running, the event handler will be called by cylc with *execution_timeout* as the EVENT argument:

- *type*: float (minutes)
- *default*: (none)

A.4.1.10.5 [runtime] → [[_NAME_]] → [[[event hooks]]] → reset timer

If you set an execution timeout the timer can be reset to zero every time a message is received from the running task (which indicates the task is still alive). Otherwise, the task will timeout if it does not finish in the allotted time regardless of incoming messages.

- *type*: boolean
- *default*: False

A.4.1.11 [runtime] → [[_NAME_]] → [[[environment]]]

The user defined task execution environment. Variables defined here can refer to cylc suite and task identity variables, which are exported earlier in the task job script, and variable assignment expressions can use cylc utility commands because access to cylc is also configured earlier in the script. See also *Task Execution Environment*, Section 8.4.4.

A.4.1.11.1 [runtime] → [[_NAME_]] → [[[environment]]] → __VARIABLE__

Replace __VARIABLE__ with any number of environment variable assignment expressions. Order of definition is preserved so values can refer to previously defined variables. Values are passed through to the task job script without evaluation or manipulation by cylc, so any variable assignment expression that is legal in the job submission shell can be used. White space around the '=' is allowed (as far as cylc's suite.rc parser is concerned these are just normal configuration items).

- *type*: string
- *default*: (none)
- *legal values*: depends to some extent on the task job submission shell (Section A.4.1.8.3).
- *examples*, for the bash shell:

```
— FOO = $HOME/bar/baz
— BAR = ${FOO}${GLOBALVAR}
— BAZ = $( echo "hello world")
— WAZ = ${FOO%.jpg}.png
— NEXT_CYCLE = $( cylc cycletime -a 6 )
— PREV_CYCLE = `cylc cycletime -s 6'
— ZAZ = "${FOO#bar}"#<-- QUOTED to escape the suite.rc comment character
```

A.4.1.12 [runtime] → [[_NAME_]] → [[[directives]]]

Batch queue scheduler directives. Whether or not these are used depends on the job submission method. For the built-in loadleveler, pbs, and sge methods directives are written to the top of the task job script in the correct format for the method. Specifying directives individually like this allows use of default directives that can be individually overridden at lower levels of the runtime namespace hierarchy.

A.4.1.12.1 [runtime] → [[_NAME_]] → [[[directives]]] → __DIRECTIVE__

Replace __DIRECTIVE__ with each directive assignment, e.g. `class = parallel`

- *type*: string
- *default*: (none)

Example directives for the built-in job submission methods are shown in Section 10.2.

A.4.1.13 [runtime] → [[_NAME_]] → [[[outputs]]]

This section is only required if other tasks need to trigger off specific internal outputs of this task (as opposed to triggering off it finishing). The task implementation must report the specified output messages by calling `cylc task message` when the corresponding real outputs have been completed.

A.4.1.13.1 [runtime] → [[_NAME_]] → [[[outputs]]] → _OUTPUT_

Replace _OUTPUT_ with any number of labelled output messages.

- *type*: string (a message containing <CYLC_TASK_CYCLE_TIME> with an optional offset as shown below. Note that you cannot use the corresponding shell variable \$CYLC_TASK_CYCLE_TIME here. The string substitution (replacing this special variable with the actual task cycle time) is done inside cylc, not in the task execution environment.
- *default*: (none)
- *examples*:

```
foo = "sea state products ready for <CYLC_TASK_CYCLE_TIME>"  
bar = "nwp restart files ready for <CYLC_TASK_CYCLE_TIME+6>"
```

where the item name must match the output label associated with this task in the suite dependency graph, e.g.:

```
[scheduling]  
[dependencies]  
graph = TaskA:foo => TaskB
```

A.5 [visualization]

Configuration of suite graphing and, where explicitly stated, the graph-based suite control GUI.

A.5.1 [visualization] → initial cycle time

The first cycle time to use when plotting the suite dependency graph.

- *type*: integer
- *default*: 2999010100

A.5.2 [visualization] → final cycle time

The last cycle time to use when plotting the suite dependency graph. Typically this should be just far enough ahead of the initial cycle to show the full suite.

- *type*: integer
- *default*: 2999010123

A.5.3 [visualization] → collapsed families

A list of family (namespace) names to be shown in the collapsed state (i.e. the family members will be replaced by a single family node) when the suite is plotted in the graph viewer or the graph-based suite control GUI. This item determines how family groups are shown *initially* in the suite control GUI; subsequently you can use the interactive controls to group and ungroup nodes at will. For the same reason (presence of interactive grouping controls) this item is ignored if the suite is reparsed during graph viewing (other changes to graph styling will be picked up and applied if the graph viewer detects that the suite.rc file has changed).

- *type*: list of family names
- *default*: (none)

A.5.4 [visualization] → use node color for edges

Graph edges (dependency arrows) can be plotted in the same color as the upstream node (task or family) to make paths through a complex graph easier to follow.

- *type*: boolean
- *default*: True

A.5.5 [visualization] → default node attributes

Set the default attributes (color and style etc.) of graph nodes (tasks and families). Attribute pairs must be quoted to hide the internal = character.

- *type*: list of quoted '`attribute=value`' pairs
- *legal values*: see graphviz or pygraphviz documentation
- *default*: '`style=unfilled`', '`color=black`', '`shape=box`'

A.5.6 [visualization] → default edge attributes

Set the default attributes (color and style etc.) of graph edges (dependency arrows). Attribute pairs must be quoted to hide the internal = character.

- *type*: list of quoted '`attribute=value`' pairs
- *legal values*: see graphviz or pygraphviz documentation
- *default*: '`color=black`'

A.5.7 [visualization] → [[node groups]]

Define named groups of graph nodes (tasks and families) that can have attributes assigned to them en masse in the [visualization] → [[node attributes]] section.

A.5.7.1 [visualization] → [[node groups]] → __GROUP__

Replace __GROUP__ with each named group of task or family names.

- *type*: comma separated list of task or family names
- *default*: (none)
- *example*:

```
PreProc = foo, bar
PostProc = baz, waz
```

A.5.8 [visualization] → [[node attributes]]

Here you can assign graph node attributes to specific tasks, or named groups defined in [visualization] → [[node groups]].

A.5.8.1 [visualization] → [[node attributes]] → __NAME__

Replace __NAME__ with any task, family, or group that you want to assign attributes to.

- *type*: list of quoted '`attribute=value`' pairs
- *legal values*: see graphviz or pygraphviz documentation

- *default*: (none)
- *example*: (with reference to the node groups defined above)

```
PreProc = 'style=filled', 'color=blue'
PostProc = 'color=red'
foo = 'style=unfilled'
```

A.5.9 [visualization] → [[run time graph]]

Cylc can generate graphs of dependencies resolved at run time, i.e. what actually triggers off what as the suite runs. This feature is retained mainly for development and debugging purposes. You can use simulation mode to generate run time graphs very quickly.

A.5.9.1 [visualization] → [[run time graph]] → enable

Run time graphing is now disabled by default.

- *type*: boolean
- *default*: False

A.5.9.2 [visualization] → [[run time graph]] → cutoff

New nodes will be added to the run time graph as the corresponding tasks trigger, until their cycle time exceeds the initial cycle time by more than this cutoff, in hours.

- *type*: integer (≥ 0 , hours)
- *default*: 24

A.5.9.3 [visualization] → [[run time graph]] → directory

Where to put the run time graph file, `runtime-graph.dot`.

- *type*: string (a valid directory path, may contain environment variables)
- *default*: `$CYLC_SUITE_DEF_PATH/graphing`

A.6 Special Placeholder Variables

A small number of special variables are used as placeholders in cylc suite definitions:

- `<CYLC_TASK_CYCLE_TIME>` and `<CYLC_TASK_CYCLE_TIME+/-offset>`

These are replaced with the actual cycle time of the task (with optional computed offset) in the message strings registered for explicit internal (pre-completion) task outputs. They correspond to `$CYLC_TASK_CYCLE_TIME` in the task execution environment.

- `<TASK>` and `<NAMESPACE>`

These placeholder variables are replaced with actual task or namespace names in runtime settings. The difference between the two is explained in Section 8.4.3, *Defining Multiple Namespaces At Once*.

- `<ASYNCID>`

This placeholder variable is replaced with the actual “asynchronous ID” (e.g. satellite pass ID) for repeating asynchronous tasks.

To use real (programming language) variables in suite definitions, use the Jinja2 template processor (Section 8.6).

A.7 Default Suite Configuration

Cylc provides, via the suite.rc spec file, sensible default values for many configuration items so that users may not need to explicitly configure log directories and the like. The defaults are sufficient, in fact, to define simple test suites by dependency graph alone (command scripting, for example, defaults to printing a simple message, sleeping for ten seconds, and then exiting). The following listing shows all current legal items and any default values:

```
# SUITE.RC DEFAULTS
# This file shows ALL cylc configuration items in context, with defaults.
# Below, 'None' is the Python undefined value; '(none)' is more general.
#
# DEVELOPER NOTE: this file must be kept up to date with the suite.rc
# specification file ${CYLC_DIR}/conf/suiterc.spec, AND with higher level
# handling of configuration items (e.g. remote logging directories
# default to the local logging directories; this is effected within cylc
# not via the specification file).

title = "No title provided" # DEFAULT
description = "No description provided" # DEFAULT

[cylc]
    UTC mode = False # DEFAULT
    simulation mode only = False # DEFAULT
    use secure passphrase = False # DEFAULT
    [[logging]]
        directory = '$HOME/cylc-run/${CYLC_SUITE_REG_NAME}/log/suite' # DEFAULT
        roll over at start-up = True # DEFAULT
    [[state dumps]]
        directory = '$HOME/cylc-run/${CYLC_SUITE_REG_NAME}/state' # DEFAULT
        number of backups = 10 # DEFAULT
    [[event hooks]]
        script = None
        events = (none)
    [[lockserver]]
        enable = False # DEFAULT
        simultaneous instances = False # DEFAULT
    [[environment]]
        # (none)
    [[simulation mode]]
        clock offset = 24 # DEFAULT
        clock rate = 10 # DEFAULT
        command scripting = "echo SIMULATION MODE ${CYLC_TASK_ID}; sleep 10; echo BYE" # DEFAULT
        [[[job submission]]]
            method = background # DEFAULT
        [[[event hooks]]]
            enable = False # DEFAULT
[scheduling]
    initial cycle time = None
    final cycle time = None
    runahead limit = 25 # DEFAULT
    [[queues]]
        [[[default]]]
            limit = 0 (no limit) # DEFAULT
            members = (all tasks) # DEFAULT
    [[special tasks]]
        clock-triggered = (none)
        start-up = (none)
        cold-start = (none)
        sequential = (none)
        one-off = (none)
        explicit restart outputs = (none)
        exclude at start-up = (none)
        include at start-up = (none)
    [[dependencies]]
        graph = None
        [[[many]]]
            graph = None
            daemon = None
```

```
[runtime]
  [[root]]
    inherit = None # (other namespaces inherit first from root)
    description = "No description provided" # DEFAULT
    initial scripting = None
    pre-command scripting = None
    command scripting = "echo DUMMY ${CYLC_SUITE_REG_NAME}; sleep 10; echo BYE" # DEFAULT
    post-command scripting = None
    manual completion = False # DEFAULT
  [[[job submission]]]
    method = background # DEFAULT
    command template = None
    shell = /bin/bash # DEFAULT
    log directory = '$HOME/cylc-run/${CYLC_SUITE_REG_NAME}/log/job' # DEFAULT
    share directory = '${CYLC_SUITE_DEF_PATH}/share' # DEFAULT
    work directory = '${CYLC_SUITE_DEF_PATH}/work/${CYLC_TASK_ID}' # DEFAULT
  [[[remote]]]
    host = None
    owner = None
    cylc directory = None
    suite definition directory = (local directory with '$HOME' replaced) # DEFAULT
    log directory = (local directory with '$HOME' replaced) # DEFAULT
    share directory = (local directory with '$HOME' replaced) # DEFAULT
    work directory = (local directory with '$HOME' replaced) # DEFAULT
    remote shell template = 'ssh -oBatchMode=yes %s' # DEFAULT
  [[[event hooks]]]
    script = None
    events = (none)
    submission timeout = None
    execution timeout = None
    reset time = False # DEFAULT
  [[[environment]]]
    # (SUITE AND TASK IDENTITY VARIABLES ARE PROVIDED)
  [[[directives]]]
    # (none)
  [[[outputs]]]
    # (none)

[visualization]
  initial cycle time = 2999010100 # DEFAULT (via config.py, not suiterc.spec)
  final cycle time = 2999010123 # DEFAULT (via config.py, not suiterc.spec)
  collapsed families = (none)
  use node color for edges = True # DEFAULT
  default node attributes = 'style=unfilled', 'color=black', 'shape=box' # DEFAULT
  default edge attributes = 'color=black' # DEFAULT
  [[node groups]]
    # (none)
  [[node attributes]]
    # (none)
  [[run time graph]]
    enable = False # DEFAULT
    cutoff = 24 # DEFAULT
    directory = '${CYLC_SUITE_DEF_PATH}/graphing' # DEFAULT
```

B COMMAND REFERENCE

B Command Reference

```
Cylc is a (meta)scheduler for complex cycling meteorological forecasting systems and related data processing. This is the cylc command line interface. The graphical interface is g cylc.
```

```
Release version: cylc-4.2.0
```

USAGE:

```
% cylc -v,--version          # cylc version
% cylc help,--help,-h,?        # this help page
% cylc help CATEGORY          # help by category
% cylc CATEGORY help          # (ditto)
% cylc help [CATEGORY] COMMAND # command help
% cylc [CATEGORY] COMMAND help,--help # (ditto)
```

Command syntax:

```
% cylc [CATEGORY] COMMAND [options] SUITE [arguments]
% cylc [CATEGORY] COMMAND [options] SUITE TASK [arguments]
```

Commands and categories can be abbreviated; categories can disambiguate highly abbreviated commands but their use is optional:

```
% cylc control trigger SUITE TASK      # trigger TASK in SUITE
% cylc trigger SUITE TASK              # ditto
% cylc con trig SUITE TASK            # ditto
% cylc c t SUITE TASK                # ditto
```

(these commands target a particular TASK in a particular SUITE).

Suites are referred to by registered names that allow hierarchical grouping, e.g. nwp.test.LAM1, nwp.test.LAM2, nwp.oper.global.

Tasks are identified by NAME%TAG where TAG is a cycle time for cycling tasks (GetData%2011080806) or 'a:Int' for asynchronous tasks (foo%a:1).

Here's how to drill down to usage help for a particular command:

```
% cylc help           # list all available categories (this page)
% cylc help prep      # list commands in category 'preparation'
% cylc help prep edit # command usage help for 'cylc [prep] edit'
```

Command CATEGORIES:

```
all ..... the complete cylc command set
db|database ... private and central suite registration.
preparation ... suite editing, graphing, validation etc.
information ... retrieve and print information about a suite.
discovery ..... what suites are running at the moment?
control ..... running, monitoring, and controlling suites.
utility ..... cycle time arithmetic, filename templating, etc.
task ..... running single tasks, and task messaging commands.
hook ..... suite and task event hook scripts.
admin ..... commands for use by the cylc administrator.
license|GPL ... Software licensing information (GPL v3.0).
```

B.1 Command Categories

B.1.1 admin

```
CATEGORY: admin - commands for use by the cylc administrator.
```

```
HELP: cylc [admin] COMMAND help,--help
You can abbreviate admin and COMMAND.
The category admin may be omitted.
```

COMMANDS:

```
check-examples ... (ADMIN) Check all example suites validate
create-cdb ..... (ADMIN) Create the central suite database
test-db ..... (ADMIN) Automated suite database test
test-suite ..... (ADMIN) Automated cylc scheduler test
```

B.1.2 all

```
CATEGORY: all - the complete cylc command set

HELP: cylc [all] COMMAND help,--help
You can abbreviate all and COMMAND.
The category all may be omitted.

COMMANDS:
alias ..... Register an alias for another registered suite
block ..... Do not comply with subsequent intervention commands
browse ..... Browse cylc documentation.
check-examples ..... (ADMIN) Check all example suites validate
checkvars ..... Check required environment variables en masse
conditions ..... Print the GNU General Public License v3.0
copy|cp ..... Copy a suite or group of suites.
create-cdb ..... (ADMIN) Create the central suite database
cycletime ..... Cycle time arithmetic
depend ..... Add prerequisites to tasks on the fly
diff|compare ..... Compare two suite definitions and print differences
dump ..... Print the state of each task in a running suite
edit ..... Edit a suite.rc file in $EDITOR, optionally inlined
export ..... Export private registrations to the central database
failed|task-failed ..... Release task lock and report failure
gcylc ..... The cylc Graphical User Interface
get-config ..... Retrieve any configuration data from a suite.
get-directory ..... Print a suite definition directory path
graph ..... Plot suite dependency graphs or runtime namespace hierarchies
hold ..... Put a hold on a suite or a single task
housekeeping ..... Parallel archiving and cleanup on cycle time offsets
import ..... Import central registrations to your private database
insert ..... Insert a task or group into a running suite
jobscript ..... Generate a task job script and print it to stdout
list|ls ..... Print a suite's task list or runtime namespace hierarchy
lockclient|lc ..... Manual suite and task lock management
lockserver ..... The cylc lockserver daemon
log ..... Print or view suite logs, with filtering
maxrunahead ..... Change the runahead limit in a running suite.
message|task-message ..... Report progress and completion of outputs
monitor ..... An in-terminal suite monitor (see also gcylc)
nudge ..... Cause the cylc task processing loop to be invoked
ping ..... Check that a suite is running
print ..... Print private or central suite registrations
purge ..... Remove a full dependency tree from a running suite
refresh ..... Report invalid registrations and update suite titles
register ..... Add a suite to your private registration database
release|unhold ..... Release a hold on a suite or a single task
remove|kill ..... Remove a task from a running suite
reregister|rename ..... Change a suite's registered name.
reset ..... Force a task to waiting, ready, or succeeded state
restart ..... Restart a suite from a previous state
run|start ..... Start a suite running at a specified cycle time
scan ..... Scan a host for running suites and lockservers
scp-transfer ..... Scp-based file transfer for cylc suites
search|grep ..... An intelligent search tool for cylc suites
show ..... Print task-specific information (prerequisites...)
started|task-started ..... Acquire a task lock and report started
stop|shutdown ..... Stop a suite running by various means
submit|single ..... Run a single task just as its parent suite would
succeeded|task-succeeded ..... Release task lock and report succeeded
template ..... Powerful cycle time offset filename templating
test-db ..... (ADMIN) Automated suite database test
test-suite ..... (ADMIN) Automated cylc scheduler test
trigger ..... Cause a task to trigger immediately
unblock ..... Comply with subsequent intervention commands
unregister ..... Remove suites from the private or central database
validate ..... Parse and validate a suite config (suite.rc) file
verbosity ..... Change a suite's logging verbosity level
view ..... View a temporary suite.rc file, inlined or processed
warranty ..... Print the GPLv3 disclaimer of warranty
```

B.1.3 control

```
CATEGORY: control - running, monitoring, and controlling suites.

HELP: cylc [control] COMMAND help,--help
You can abbreviate control and COMMAND.
The category control may be omitted.

COMMANDS:
block ..... Do not comply with subsequent intervention commands
depend ..... Add prerequisites to tasks on the fly
gcylc ..... The cylc Graphical User Interface
hold ..... Put a hold on a suite or a single task
insert ..... Insert a task or group into a running suite
lockclient|lc .... Manual suite and task lock management
lockserver ..... The cylc lockserver daemon
maxrunahead ..... Change the runahead limit in a running suite.
nudge ..... Cause the cylc task processing loop to be invoked
purge ..... Remove a full dependency tree from a running suite
release|unhold ... Release a hold on a suite or a single task
remove|kill ..... Remove a task from a running suite
reset ..... Force a task to waiting, ready, or succeeded state
restart ..... Restart a suite from a previous state
run|start ..... Start a suite running at a specified cycle time
stop|shutdown .... Stop a suite running by various means
trigger ..... Cause a task to trigger immediately
unblock ..... Comply with subsequent intervention commands
verbosity ..... Change a suite's logging verbosity level
```

B.1.4 database

```
CATEGORY: db|database - private and central suite registration.

HELP: cylc [db|database] COMMAND help,--help
You can abbreviate db|database and COMMAND.
The category db|database may be omitted.

COMMANDS:
alias ..... Register an alias for another registered suite
browse ..... Browse cylc documentation.
copy|cp ..... Copy a suite or group of suites.
export ..... Export private registrations to the central database
gcylc ..... The cylc Graphical User Interface
get-directory ..... Print a suite definition directory path
import ..... Import central registrations to your private database
print ..... Print private or central suite registrations
refresh ..... Report invalid registrations and update suite titles
register ..... Add a suite to your private registration database
reregister|rename ... Change a suite's registered name.
 unregister ..... Remove suites from the private or central database
```

B.1.5 discovery

```
CATEGORY: discovery - what suites are running at the moment?

HELP: cylc [discovery] COMMAND help,--help
You can abbreviate discovery and COMMAND.
The category discovery may be omitted.

COMMANDS:
ping ... Check that a suite is running
scan ... Scan a host for running suites and lockservers
```

B.1.6 hook

```
CATEGORY: hook - suite and task event hook scripts.

HELP: cylc [hook] COMMAND help,--help
You can abbreviate hook and COMMAND.
```

The category hook may be omitted.

COMMANDS:

```
email-suite ... A suite event hook script that sends email alerts
email-task .... A task event hook script that sends email alerts
```

B.1.7 information

CATEGORY: information – retrieve and print information about a suite.

HELP: cylc [information] COMMAND help,--help
 You can abbreviate information and COMMAND.
 The category information may be omitted.

COMMANDS:

```
browse ..... Browse cylc documentation.
dump ..... Print the state of each task in a running suite
gcylc ..... The cylc Graphical User Interface
get-config .... Retrieve any configuration data from a suite.
list|ls ..... Print a suite's task list or runtime namespace hierarchy
log ..... Print or view suite logs, with filtering
monitor ..... An in-terminal suite monitor (see also gcylc)
nudge ..... Cause the cylc task processing loop to be invoked
show ..... Print task-specific information (prerequisites...)
```

B.1.8 license

CATEGORY: license|GPL – Software licensing information (GPL v3.0).

HELP: cylc [license|GPL] COMMAND help,--help
 You can abbreviate license|GPL and COMMAND.
 The category license|GPL may be omitted.

COMMANDS:

```
conditions ... Print the GNU General Public License v3.0
warranty ..... Print the GPLv3 disclaimer of warranty
```

B.1.9 preparation

CATEGORY: preparation – suite editing, graphing, validation etc.

HELP: cylc [preparation] COMMAND help,--help
 You can abbreviate preparation and COMMAND.
 The category preparation may be omitted.

COMMANDS:

```
diff|compare ... Compare two suite definitions and print differences
edit ..... Edit a suite.rc file in $EDITOR, optionally inlined
gcylc ..... The cylc Graphical User Interface
graph ..... Plot suite dependency graphs or runtime namespace hierarchies
jobscript ..... Generate a task job script and print it to stdout
list|ls ..... Print a suite's task list or runtime namespace hierarchy
search|grep .... An intelligent search tool for cylc suites
validate ..... Parse and validate a suite config (suite.rc) file
view ..... View a temporary suite.rc file, inlined or processed
```

B.1.10 task

CATEGORY: task – running single tasks, and task messaging commands.

HELP: cylc [task] COMMAND help,--help
 You can abbreviate task and COMMAND.
 The category task may be omitted.

COMMANDS:

```
failed|task-failed ..... Release task lock and report failure
message|task-message ..... Report progress and completion of outputs
started|task-started ..... Acquire a task lock and report started
submit|single ..... Run a single task just as its parent suite would
succeeded|task-succeeded ... Release task lock and report succeeded
```

B.1.11 utility

```
CATEGORY: utility - cycle time arithmetic, filename templating, etc.

HELP: cylc [utility] COMMAND help,--help
  You can abbreviate utility and COMMAND.
  The category utility may be omitted.

COMMANDS:
  checkvars ..... Check required environment variables en masse
  cycletime ..... Cycle time arithmetic
  housekeeping ... Parallel archiving and cleanup on cycle time offsets
  scp-transfer ... Scp-based file transfer for cylc suites
  template ....... Powerful cycle time offset filename templating
```

B.2 Commands**B.2.1 alias**

```
Usage: cylc [db] alias [options] SUITE ALIAS
```

Register an alias for suite. Using the alias on the command line is entirely equivalent to using target suite name (and if you run a suite via an alias, run time identity will correspond to the target suite).

```
$ cylc alias global.ensemble.parallel.test3 bob
$ cylc edit bob
$ cylc run bob      # work with global.ensemble.parallel.test3
$ cylc stop bob
```

This differs from registering the target suite definition under another name with 'cylc register' - in that case run time suite identity will correspond to the new registered name.

Arguments:

SUITE	- Target suite.
ALIAS	- An alternative name for the same suite.

Options:

-h, --help	show this help message and exit
-v, --verbose	Print extra information.

B.2.2 block

```
Usage: cylc [control] block [options] SUITE
```

A blocked suite refuses to comply with intervention commands until deliberately unblocked. This is a crude security measure to guard against accidental intervention in your own suites. It may be a useful aid when running multiple suites at once.

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
-------	---------------

Options:

-h, --help	show this help message and exit
--debug	Turn on exception tracebacks.
-f, --force	Do not ask for confirmation before acting.

B.2.3 browse

```
Usage: cylc [info] browse [options]
```

Browse cylc documentation:

1/ The User Guide in PDF format (default)
2/ The User Guide in HTML format

```

3/ The change log (plain text format)
4/ The cylc internet homepage.

Depending on the chosen browse target you must define one or more of:
$PDF_READER    (e.g. evince)
$html_READER   (e.g. firefox)
$EDITOR        (e.g. vim)
$GEDITOR       (e.g. 'gvim -f')

Options:
-h, --help      show this help message and exit
--html         Browse the HTML User Guide (default is PDF).
--www          Browse the cylc internet homepage.
--log          Browse the cylc repository change log.
-g             Use $EDITOR instead of $EDITOR to view the change log.

```

B.2.4 check-examples

```

USAGE: cylc [admin] check-examples

Check that all cylc example suites validate successfully.
Do this immedietly prior to rolling a new cylc release.

```

B.2.5 checkvars

```

Usage: cylc checkvars [options] VAR NAMES

Check that each member of a list of environment variables is defined,
and then optionally check their values according to the chosen
commandline option. Note that THE VARIABLES MUST BE EXPORTED AS THIS
SCRIPT NECESSARILY EXECUTES IN A SUBSHELL.

All of the input variables are checked in turn and the results printed.
If any problems are found then, depending on use of '-w,--warn-only',
this script either aborts with exit status 1 (error) or emits a stern
warning and exits with status 0 (success).

Arguments:
  VAR NAMES      Space-separated list of environment variable names.

Options:
-h, --help      show this help message and exit
-d, --dirs-exist  Check that the variables refer to directories that
                  exist.
-c, --create-dirs  Attempt to create the directories referred to by the
                  variables, if they do not already exist.
-p, --create-parent-dirs  Attempt to create the parent directories of files
                  referred to by the variables, if they do not already
                  exist.
-f, --files-exist  Check that the variables refer to files that exist.
-i, --int          Check that the variables refer to integer values.
-s, --silent       Do not print the result of each check.
-w, --warn-only    Print a warning instead of aborting with error status.

```

B.2.6 conditions

```

USAGE: cylc [license] warranty [--help]
Cylc is release under the GNU General Public License v3.0
This command prints the GPL v3.0 license in full.

```

```

Options:
--help  Print this usage message.

```

B.2.7 copy

```
Usage: cylc [db] copy|cp [OPTIONS] SOURCE TARGET TOPDIR

Make copies of suites, or groups of suites, in your private database
(to get copies of suites in the central database use 'cylc import').
```

Cylc suites have hierarchical "registration names" that are analogous to file paths in a directory tree (but which are not necessarily related to the actual directory paths in which the suite definitions are stored!). Consider the following three suites:

```
% cylc db print '^foo'      # printed in flat form
foo.bag      | "Test Suite Zero" | /tmp/oliverh/zero
foo.bar.waz  | "Test Suite Two"  | /tmp/oliverh/two
foo.bar.baz  | "Test Suite One"   | /tmp/oliverh/one

% cylc db print -t '^foo'  # printed in tree from
foo
  |-bag    "Test Suite Zero" | /tmp/oliverh/zero
  `|-bar
    |-baz  "Test Suite One"  | /tmp/oliverh/one
    `|-waz  "Test Suite Two"  | /tmp/oliverh/two
```

These suites are stored in a flat directory structure under \$TMPDIR, but they are organised in the suite database as a group 'foo' that contains the suite 'foo.bag' and a group 'foo.bar', which in turn contains the suites 'foo.bar.baz' and 'foo.bar.waz'.

When you copy suites with this command, the target registration names are determined by TARGET and the name structure underneath SOURCE, and the suite definition directories are copied into a directory tree under TOPDIR whose structure reflects the target registration names. If this is not what you want, you can copy suite definition directories manually and then register the copied directories manually with 'cylc register'.

EXAMPLES (using the three suites above):

```
% cylc db copy foo.bar.baz red $TMPDIR          # suite to suite
Copying suite definition for red
% cylc db print '^red'
red | "Test Suite One" | /tmp/oliverh/red

% cylc copy foo.bar.baz blue.green $TMPDIR      # suite to group
Copying suite definition for blue.green
% cylc db pr '^blue'
blue.green | "Test Suite One" | /tmp/oliverh/blue/green

% cylc copy foo.bar orange $TMPDIR              # group to group
Copying suite definition for orange.waz
Copying suite definition for orange.baz
% cylc db pr '^orange'
orange.waz | "Test Suite Two" | /tmp/oliverh/orange/waz
orange.baz | "Test Suite One" | /tmp/oliverh/orange/baz
```

Arguments:

```
SOURCE - Source suite or group
TARGET - Target suite or group
TOPDIR - Destination for copied suite definition directories
```

Options:

```
-h, --help      show this help message and exit
-v, --verbose   Print extra information.
```

B.2.8 create-cdb

```
USAGE: cylc [admin] create-cdb
```

This command is used by the cylc administrator to create the central suite database immediately after installing cylc, or to upgrade the centrally held cylc example suites after upgrading to a new version.

The example suites are registered with the cylc version number in the registration path so that multiple versions can be retained.

Seconds since epoch are also appended to the cylc version number to allow multiple uploads at the same cylc version (in case the examples need to be modified). You can then rename the uploaded examples group and unregister any older uploads if necessary.

B.2.9 cylctime

Usage: cylc [util] cylctime [options] [CYCLE]

Perform arithmetic operations on cycle times and print the results to stdout. Examples:

```
# offset from explicit cycle time:  
$ cylc cylctime -s 6 2010082318  
2010082312  
  
# offset from $CYLC_TASK_CYCLE_TIME:  
$ export CYLC_TASK_CYCLE_TIME=2010082318  
$ cylc cylctime -s 6  
2010082312
```

Arguments:

CYCLE	(YYYYMMDDHH) defaults to \$CYLC_TASK_CYCLE_TIME.
-------	--

Options:

-h, --help	show this help message and exit
-s HOURS, --subtract=HOURS	Subtract HOURS from CYCLE
-a HOURS, --add=HOURS	Add HOURS to CYCLE
-o HOURS, --offset=HOURS	Apply an offset of +/-HOURS to CYCLE
--year	Print only YYYY of result
--month	Print only MM of result
--day	Print only DD of result
--hour	Print only HH of result

B.2.10 depend

Usage: cylc [control] depend [options] SUITE TASK DEP

Add new dependencies on the fly to tasks in a running suite. If DEP is a task ID the target TASK will depend on that task finishing, otherwise DEP can be an explicit quoted message such as

"Data files uploaded for 2011080806"

(presumably there will be another task in the suite, or you will insert one, that reports that message as an output).

Prerequisites added on the fly are not propagated to the successors of TASK, and they will not persist in TASK across a suite restart.

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
TASK	Target task.
DEP	New dependency for the target task.

Options:

-h, --help	show this help message and exit
--debug	Turn on exception tracebacks.
-f, --force	Do not ask for confirmation before acting.

B.2.11 diff

```
Usage: cylc [prep] diff|compare [options] SUITE1 SUITE2

Compare two suite definitions and display any differences.

Differencing is done after parsing the suite.rc files so it takes
account of default values that are not explicitly defined, it disregards
the order of configuration items, and it sees any include-file content
after inlining has occurred.

Seemingly identical suites can differ slightly if they use default
configuration items, such as the default job submission log directory,
that are suite-specific (i.e. they include $CYLC_SUITE_REG_NAME etc.). 

Files in the suite bin directory and other sub-directories of the
suite definition directory are not currently differenced.

Arguments:
  SUITE1, SUITE2 - suites to compare.

Options:
  -h, --help      show this help message and exit
  -n, --nested    print suite.rc section headings in nested form.
  -c, --central   target the central database.
```

B.2.12 dump

```
Usage: cylc [info] dump [options] SUITE

Print current suite state information (e.g. the state of each task). For
small suites 'watch cylc [info] dump SUITE' is an effective non-GUI real
time monitor (but see also 'cylc monitor').

For more information about a specific task, such as the current state of
its prerequisites and outputs, see 'cylc [info] show'.

Examples:
  Display the state of all running tasks, sorted by cycle time:
  % cylc [info] dump -ts SUITE | grep running

  Display the state of all tasks in a particular cycle:
  % cylc [info] dump -t SUITE | grep 2010082406

Arguments:
  SUITE           Target suite.

Options:
  -h, --help      show this help message and exit
  -o USER, --owner=USER
                  Owner of the target suite (defaults to $USER).
  --host=HOST     Cylc suite host (defaults to local host).
  -f, --force     (No effect; for consistency with interactive commands)
  --debug         Turn on exception tracebacks.
  -g, --global    Global information only.
  -t, --tasks     Task states only.
  -s, --sort      Task states only; sort by cycle time instead of name.
```

B.2.13 edit

```
Usage: 1/ cylc [prep] edit SUITE
       2/ cylc [prep] edit -i,--inline SUITE
       3/ cylc [prep] edit --cleanup SUITE

This command facilitates editing of suite definitions (suite.rc files),
with optional reversible inlining of include-files. Note that suites
requiring processing with a template engine such as Jinja2 can only
be edited in unprocessed form, but the processed suite definition can be
viewed using 'cylc [prep] view'.
```

1/ Change to the suite definition directory and load the suite.rc file

in your \$EDITOR. This is a convenience so that you can edit the suite without having to remember the suite definition location.

2/ Edit suite.rc files WITH INCLUDE-FILES INLINED between special markers. The real suite.rc file is temporarily replaced so that THE INLINED SUITE.RC FILE IS THE "LIVE" FILE DURING EDITING. The inlined suite.rc file gets split up again when you exit the editor. Include-files can be nested or multiply-included; in the latter case only the first inclusion is inlined (this serves to prevent conflicting changes being made to the same file).

3/ Remove backup files left by previous INLINED edit sessions.

INLINED EDITING SAFETY: The suite.rc file and its include-files are automatically backed up prior to an inlined editing session. If the editor dies mid-session just invoke 'cylc edit -i' again to recover from the last saved inlined file. On exiting the editor, if any of the original include-files are found to have changed due to external intervention during editing you will be warned and the affected files will be written to new backups instead of overwriting the originals. Finally, the inlined suite.rc file is also backed up on exiting the editor, to allow recovery in case of accidental corruption of the inlined include-file boundary marker lines.

The edit process is spawned in the foreground as follows:

```
$ (G)EDITOR suite.rc  
$EDITOR or $EDITOR, and $TMDPIR, must be in your environment.
```

Examples:

```
export EDITOR=vim  
export GEDITOR='gvim -f'      # -f: do not detach from parent shell!!  
export EDITOR='xterm -e vim'  # for gcylc, if gvim is not available  
export GEDITOR=emacs  
export EDITOR='emacs -nw'
```

You can set both \$GEDITOR and \$EDITOR to a GUI editor if you like, but \$EDITOR at least ***must*** be a GUI editor, or an in-terminal invocation of a non-GUI editor, if you want to spawn editing sessions via gcylc.

Arguments:

SUITE Target suite.

Options:

-h, --help	show this help message and exit
-i, --inline	Edit with include-files inlined as described above.
--central	Target the central suite database.
--cleanup	Remove backup files left by previous inlined edit sessions.
-g, --gui	Use GUI editor \$GEDITOR instead of \$EDITOR. This option is automatically used when an editing session is spawned by gcylc.

B.2.14 email-suite

USAGE: cylc [hook] email-suite EVENT SUITE MESSAGE

This is a simple suite event hook script that sends an email. The command line arguments are supplied automatically by cylc.

For example, to get an email alert when a suite shuts down:

```
# SUITE.RC  
[cylc]  
  [[environment]]  
    MAIL_ADDRESS = foo@bar.baz.waz  
  [[event hooks]]  
    events = shutdown  
    script = cylc email-suite
```

See the Suite.rc Reference (Cylc User Guide) for more information on suite and task event hooks and event handler scripts.

B.2.15 email-task

```
USAGE: cylc [hook] email-task EVENT SUITE TASKID MESSAGE
```

This is a simple task event hook handler script that sends an email. The command line arguments are supplied automatically by cylc.

For example, to get an email alert whenever any task fails:

```
# SUITE.RC
[cylc]
  [[environment]]
    MAIL_ADDRESS = foo@bar.baz.waz
[runtime]
  [[root]]
    [[[event hooks]]]
      events = failed
      script = cylc email-task
```

See the Suite.rc Reference (Cylc User Guide) for more information on suite and task event hooks and event handler scripts.

B.2.16 export

```
Usage: cylc [db] export [OPTIONS] SOURCE [TARGET]
```

Export suites or groups of suites to the central suite database.

Central registrations always start with the suite owner's username. If a target name is not given ...

Arguments:

```
SOURCE - Suite to export to the central database
TARGET - Target central suite name, if different from SOURCE.
```

Options:

```
-h, --help      show this help message and exit
-v, --verbose   Print extra information.
```

B.2.17 failed

```
Usage: cylc [task] failed [options] [REASON]
```

This is part of the cylc external task interface.

Release my lock to the lockserver, and report that I have failed.

Arguments:

```
REASON - message explaining why the task failed.
```

Options:

```
-h, --help      show this help message and exit
```

B.2.18 gcylc

Note that you can invoke gcylc directly.

```
Usage: cylc gcylc [SUITE]
      gcylc [SUITE]
```

This is the cylc graphical user interface. It is functionally equivalent to the command line interface ('cylc help') in most respects.

```
1/ gcylc &
This invokes the gcylc main window, which displays suites in your private registration database, and in the central database available to all users. By right-clicking on suites or registration groups you can gain access to all cylc functionality, from editing and graphing through to suite control and monitoring.
```

```

2/ gcylc [-g,--graph] SUITE &
This directly invokes a suite control and monitoring application for a
particular suite. Alternatively you can get this by right-clicking on
the suite in the main gcylc suite database viewer (1/ above). Without
the -g option you'll get a filtered text treeview GUI; with it, a
dependency graph based GUI.

NOTE: daemonize important suites with the POSIX nohup command:
$ nohup gcylc [options] SUITE &

Arguments:
  SUITE      Target suite.

Options:
  -h, --help   show this help message and exit
  -g, --graph  With SUITE - invoked the new dependency graph based suite
               control and monitoring interface.

```

B.2.19 get-config

```

Usage: cylc [info] get-config SUITE [HIERARCHY]

Print configuration items parsed from a suite definition (including any
defaults not explicitly set in the suite.rc file). This enables scripts
to control suites or process their output without assuming directory
locations and so on.

Config items containing spaces must be quoted. If the requested config
item is not a single value the corresponding internal data structure (a
nested Python dict) will be printed verbatim.

#_____EXAMPLE
# excerpt from a suite definition registered as foo.bar.baz:
title = local area implementation of modelX
[cylc]
  [[lockserver]]
    enable = True
[runtime]
  [[modelX]]
    command scripting = run-model.sh
    [[[environment]]]
      OUTPUT_DIR=/oper/live/modelX/output
#_____
$ cylc get-config foo.bar.baz title
$ local area implementation of modelX

$ cylc get-config foo.bar.baz cylc lockserver enable
$ True

$ cylc get-config foo.bar.baz runtime modelX environment OUTPUT_DIR
$ /oper/live/modelX/output

Arguments:
  SUITE - target suite.
  HIERARCHY - list of suite.rc item hierarchy elements.

Options:
  -h, --help       show this help message and exit
  -c, --central   Target the central suite database.
  -d, --directories Print all configured suite directory paths.

```

B.2.20 get-directory

```

Usage: cylc [db] get-directory SUITE

Retrieve and print a suite definition directory path.
How to move to a suite definition directory, for the lazy:
$ cd $(cylc get SUITE)

```

```
Arguments:
  SUITE - target suite.

Options:
  -h, --help      show this help message and exit
  -c, --central   Target the central suite database.
```

B.2.21 graph

```
Usage: 1/ cylc [prep] graph [options] SUITE [START [STOP]]
  Plot the suite.rc dependency graph for SUITE.
  2/ cylc [prep] graph [options] -f,--file FILE
  Plot the specified dot-language graph file.

Plot cylc dependency graphs in a pannable, zoomable viewer.

The viewer updates automatically when the suite.rc file is saved during
editing. By default the full cold start graph is plotted; you can omit
cold start tasks with the '-w,--warmstart' option. Specify the optional
initial and final cycle time arguments to override the suite.rc defaults.
If you just override the initial cycle, only that cycle will be plotted.

GRAPH VIEWER CONTROLS:
* Left-click to center the graph on a node.
* Left-drag to pan the view.
* Zoom buttons, mouse-wheel, or ctrl-left-drag to zoom in and out.
* Shift-left-drag to zoom in on a box.
* Also: "Best Fit" and "Normal Size".
Family (namespace) grouping controls:
Toolbar:
* "group" - group all families up to root.
* "ungroup" - recursively ungroup all families.
Right-click menu:
* "group" - close this node's parent family.
* "ungroup" - open this family node.
* "recursive ungroup" - ungroup all families below this node.

Arguments:
  SUITE    - Target suite.
  START    - Initial cycle time to plot (default=2999010100)
  STOP     - Final cycle time to plot (default=2999010123)

Options:
  -h, --help          show this help message and exit
  -w, --warmstart    Plot the mid-stream warm start (raw start) dependency
                     graph (the default is cold start).
  -n, --namespaces   Plot the suite namespace inheritance hierarchy (task
                     run time properties).
  -f FILE, --file=FILE View a specific dot-language graphfile.
  -c, --central       Target the central database.
  -o FILE, --output=FILE
                     Write an image file with format determined by file
                     extension. The image will be rewritten automatically,
                     for the configured (suite.rc) graph as you edit the
                     suite. Available formats may include png, svg, jpg,
                     gif, ps, ..., depending on your graphviz build; to see
                     what's available specify a non-existent format and
                     read the resulting error message.
```

B.2.22 hold

```
Usage: cylc [control] hold [options] SUITE [TASK]

Put a hold on a suite or a single waiting task. Putting a suite on hold
stops it from submitting any tasks that are ready to run, until it is
released. Putting a waiting task on hold prevents it from running and
spawning successors, until it is released.
```

See also 'cylc [control] release'.

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
TASK	Task to hold (NAME%YYYYMMDDHH)

Options:

-h, --help	show this help message and exit
--debug	Turn on exception tracebacks.
-f, --force	Do not ask for confirmation before acting.

B.2.23 housekeeping

```
Usage: 1/ cylc [util] housekeeping [options] SOURCE MATCH OPER OFFSET [TARGET]
Usage: 2/ cylc [util] housekeeping [options] FILE
```

Parallel archiving and cleanup of files or directories with names that contain a cycle time. Matched items are grouped into batches in which members are processed in parallel, by spawned sub-processes. Once all batch members have completed, the next batch is processed.

OPERATE ('delete', 'move', or 'copy') on items (files or directories) matching a Python-style regular expression MATCH in directory SOURCE whose names contain a cycle time (as YYYYMMDDHH, or YYYYMMDD and HH separately) more than OFFSET (integer hours) earlier than a base cycle time (which can be \$CYLC_TASK_CYCLE_TIME if called by a cylc task, or otherwise specified on the command line).

FILE is a housekeeping config file containing one or more of lines of:

```
VARNAME=VALUE
# comment
SOURCE      MATCH      OPERATION      OFFSET      [TARGET]
```

(example: \$CYLC_DIR/conf/housekeeping.eg)

MATCH must be a Python-style regular expression (NOT A SHELL GLOB EXPRESSION!) to match the names of items to be operated on AND to extract the cycle time from the names via one or two parenthesized sub-expressions - '(\d{10})' for YYYYMMDDHH, '(\d{8})' and '(\d{2})' for YYYYMMDD and HH in either order. Partial matching can be used (partial: 'foo-(\d{10})'; full: '^foo-(\d{10})\$'). Any additional parenthesized sub-expressions, e.g. for either-or matching, MUST be of the (?....) type to avoid creating a new match group.

SOURCE and TARGET must be on the local filesystem and may contain environment variables such as \$HOME or \${FOO} (e.g. as defined in the suite.rc file for suite housekeeping tasks). Variables defined in the housekeeping file itself can also be used, as above.

TARGET may contain the strings YYYYMMDDHH, YYYY, MM, DD, HH; these will be replaced with the extracted cycle time for each matched item, e.g. \$ARCHIVE/oper/YYYYMM/DD.

If TARGET is specified for the 'delete' operation, matched items in SOURCE will not be deleted unless an identical item is found in TARGET. This can be used to check that important files have been successfully archived before deleting the originals.

The 'move' and 'copy' operations are aborted if the TARGET/item already exists, but a warning is emitted if the source and target items are not identical.

To implement a simple ROLLING ARCHIVE of cycle-time labelled files or directories: just use 'delete' with OFFSET set to the archive length.

SAFE ARCHIVING: The 'move' operation is safe - it uses Python's

```
shutils.move() which renames files on the local disk partition and
otherwise copies before deleting the original. But for extra safety
consider two-step archiving and cleanup:
1/ copy files to archive, then
2/ delete the originals only if identicals are found in the archive.

Options:
-h, --help      show this help message and exit
--cycletime=YYYYMMDDHH    Cycle time, defaults to $CYLC_TASK_CYCLE_TIME
--mode=MODE    Octal umask for creating new destination
               directoriesE.g. 0775 for drwxrwxr-x
-o LIST, --only=LIST Only action config file lines matching any member of a
               comma-separated list of regular expressions.
-e LIST, --except=LIST   Only action config file lines NOT matching any member
               of a comma-separated list of regular expressions.
-v, --verbose   print the result of every action
-d, --debug     print item matching output.
-c, --cheapdiff Assume source and target identical if the same size
-b INT, --batchsize=INT   Batch size for parallel processing of matched files.
               Members of each batch (matched items) are processed in
               parallel; when a batch completes, the next batch
               starts. Defaults to a batch size of 1, i.e. sequential
               processing.
```

B.2.24 import

```
Usage: cylc [db] import [OPTIONS] SOURCE [TARGET] TOPDIR

Import suites or groups of suites from the central registration database
to your private database. Import semantics are the same as for the copy
command. The directory path of imported suite definitions under TOPDIR
will reflect the corresponding suite names, with '.' replaced by '/'.
See 'cylc db copy --help' for examples.

Arguments:
SOURCE      - Central suite or group (e.g. alice.lam.test).
[TARGET]     - Private suite or group (omit to use SOURCE: lam.test).
TOPDIR      - Top directory for imported suite definitions.

Options:
-h, --help      show this help message and exit
-v, --verbose   Print extra information.
```

B.2.25 insert

```
Usage: cylc [control] insert [options] SUITE TASK[%STOP]

Insert a task into a running suite. Inserted tasks will spawn successors
as normal unless they are 'one-off' tasks.
See also 'cylc [task] submit', for running single tasks without the scheduler.

You must be the owner of the target suite to use this command.

Arguments:
SUITE        Target suite.
TASK         Task to insert (NAME%TAG).
STOP         Optional stop TAG (e.g. a final cycle time).

Options:
-h, --help      show this help message and exit
--debug       Turn on exception tracebacks.
-f, --force    Do not ask for confirmation before acting.
```

B.2.26 jobscript

```
USAGE: cylc [prep] jobscrip [options] SUITE TASK

Generate a task job script and print it to stdout.

Here's how to capture the script in the vim editor:
% cylc jobscrip SUITE TASK | vim -
Emacs unfortunately cannot read from stdin:
% cylc jobscrip SUITE TASK > tmp.sh; emacs tmp.sh

This command wraps 'cylc [control] submit --dry-run'.

Options:
-h,--help - print this usage message.

Arguments:
SUITE      - Target suite.
TASK       - Task NAME or NAME%YYYYMMDDHH (if you omit
              the cycle time 2999010100 will be used).
```

B.2.27 list

```
Usage: cylc [info|prep] list|ls [OPTIONS] SUITE [FILTER]

Print a suite's task list; or print or graph its runtime namespace
inheritance hierarchy.

Arguments:
SUITE      - Target suite
FILTER     - regular expression filter on task names.

Options:
-h, --help   show this help message and exit
-c, --central Target the central database.
-t, --tree    Print the full runtime inheritance hierarchy.
-g, --graph   Graph the full runtime inheritance hierarchy.
-p, --pretty   Use unicode box drawing characters when printing the suite
               runtime inheritance hierarchy.
```

B.2.28 lockclient

```
Usage: cylc [control] lockclient|lc [options]

This is the command line client interface to the cylc lockserver daemon,
for server interrogation and manual lock management.

Use of the lockserver is optional (see suite.rc documentation)

Manual lock acquisition is mainly for testing purposes, but manual
release may be required to remove stale locks if a suite or task dies
without cleaning up after itself.

See also:
  cylc lockserver

Options:
-h, --help           show this help message and exit
--acquire-task=SUITE:TASK%CYCLE
                    Acquire a task lock.
--release-task=SUITE:TASK%CYCLE
                    Release a task lock.
--acquire-suite=SUITE
                    Acquire an exclusive suite lock.
--acquire-suite-nonex=SUITE
                    Acquire a non-exclusive suite lock.
--release-suite=SUITE
                    Release a suite and associated task locks
-p, --print          Print all locks.
-l, --list           List all locks (same as -p).
-c, --clear          Release all locks.
-f, --filenames     Print lockserver PID, log, and state filenames.
```

B.2.29 lockserver

```
Usage: cylc [control] lockserver [-f CONFIG] ACTION
```

The cylc lockserver daemon brokers suite and task locks for a single user. These locks are analogous to traditional lock files, but they work even for tasks that start and finish executing on different hosts. Suite locks prevent multiple instances of the same suite from running at the same time (even if registered under different names) unless the suite allows that. Task locks do the same for individual tasks (even if submitted outside of their suite using 'cylc submit').

The command line user interface for interrogating the daemon, and for manual lock management, is 'cylc lockclient'.

Use of the lockserver is optional (see suite.rc documentation).

The lockserver reads a config file that specifies the location of the daemon's process ID, state, and log files. The default config file is '\$CYLC_DIR/conf/lockserver.conf'. You can specify an alternative config file on the command line, but then all subsequent interaction with the daemon via the lockclient command must also specify the same file (this is really only for testing purposes). The default process ID, state, and log files paths are relative to \$HOME so this should be sufficient for all users.

The state file records currently held locks and, if it exists at startup, is used to initialize the lockserver (i.e. suite and task locks are not lost if the lockserver is killed and restarted). All locking activity is recorded in the log file.

Arguments:

ACTION	-	'start', 'stop', 'status', 'restart', or 'debug'
		In debug mode the server does not daemonize so its the stdout and stderr streams are not lost.

Options:

-h, --help	show this help message and exit
-c CONFIGFILE, --config-file=CONFIGFILE	Config file (default /tmp/oliverh/20348/cylc-4.2.0/conf/lockserver.conf)

B.2.30 log

```
Usage: cylc [info] log [options] SUITE
```

Print, or view in \$EDITOR, local suite log files, with optional filtering. This command is a convenience: you don't need to know the log file location on disk.

Arguments:

SUITE	- Target suite.
-------	-----------------

Options:

-h, --help	show this help message and exit
-p, --print	Print the suite log file location and exit. This is equivalent to 'cylc get-config SUITE cylc logging directory'.
-t TASK, --task=TASK	Filter the log for messages from a specific task
-f RE, --filter=RE	Filter the log with a Python-style regular expression e.g. '\[(foo bar)\].*(started succeeded)'
-r INT, --rotation=INT	Rotation number (to view older, rotated logs)
-e, --edit	View a temporary copy of the (filtered) log file in \$EDITOR instead of printing it to stdout.

B.2.31 maxrunahead

```
Usage: cylc [control] maxrunahead [options] SUITE [HOURS]
```

Change the suite runahead limit in a running suite. This is the number of hours that the fastest task is allowed to get ahead of the slowest. If a task spawns beyond that limit it will be held back from running until the slowest tasks catch up enough. **WARNING:** if you omit HOURS no limit will be used!

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
HOURS	New runahead limit (no limit if omitted).

Options:

-h, --help	show this help message and exit
--debug	Turn on exception tracebacks.
-f, --force	Do not ask for confirmation before acting.

B.2.32 message

Usage: cylc [task] message [options]

This is part of the cylc external task interface.

Report completed outputs, progress, or any other messages.

Options:

-h, --help	show this help message and exit
-p PRIORITY	message priority: NORMAL, WARNING, or CRITICAL; default NORMAL.
--next-restart-completed	Report next restart file(s) completed
--all-restart-outputs-completed	Report all restart outputs completed at once.
--all-outputs-completed	Report all internal outputs completed at once.

B.2.33 monitor

Usage: cylc [info] monitor [options] SUITE

A terminal-based suite monitor that updates the current state of all tasks in real time. It is effective even for quite large suites if '--align' is not used. Being a passive monitor that cannot intervene in a suite's operation, it is allowed to monitor suites owned by others and/or running on remote hosts.

Arguments:

SUITE	Target suite.
-------	---------------

Options:

-h, --help	show this help message and exit
-o USER, --owner=USER	Owner of the target suite (defaults to \$USER).
--host=HOST	Cylc suite host (defaults to local host).
-f, --force	(No effect; for consistency with interactive commands)
--debug	Turn on exception tracebacks.
-a, --align	Align columns by task name. This option is only useful for small suites.

B.2.34 nudge

Usage: cylc [control] nudge [options] SUITE

Cause the cylc task processing loop to be invoked in a running suite.

This happens automatically when the state of any task changes such that task processing (dependency negotiation etc.) is required, or if a clock-triggered task is ready to run.

The main reason to use this command is to update the "estimated time till completion" intervals shown in the tree-view suite control GUI, during periods when nothing else is happening.

You must be the owner of the target suite to use this command.

Arguments:

SUITE Target suite.

Options:

-h, --help	show this help message and exit
-f, --force	(No effect; for consistency with interactive commands)
--debug	Turn on exception tracebacks.

B.2.35 ping

Usage: cylc [discover] ping [options] SUITE

Check that SUITE is running.

Arguments:

SUITE Target suite.

Arguments:

SUITE Target suite.

Options:

-h, --help	show this help message and exit
-o USER, --owner=USER	Owner of the target suite (defaults to \$USER).
--host=HOST	Cylc suite host (defaults to local host).
-f, --force	(No effect; for consistency with interactive commands)
--debug	Turn on exception tracebacks.
--print-ports	Print cylc's configured port range.

B.2.36 print

Usage: cylc [db] print [options] [FILTER]

Print private or central suite registrations.

FILTERING:

(a) The filter patterns are Regular Expressions, not shell globs, so the general wildcard is '.*' (match zero or more of anything), NOT '*'.
 (b) For printing purposes there is an implicit wildcard at the end of each pattern ('foo' is the same as 'foo.*'); use the string end marker to stop this ('foo\$' matches only literal 'foo').

Options:

-h, --help	show this help message and exit
-c, --central	Target the the central database.
-t, --tree	Print registrations in nested tree form.
-p, --pretty	Use unicode box drawing characters in tree views.
-a, --align	Align columns.
-x	don't print suite definition directory paths.
-y	Don't print suite titles.
--fail	Fail (exit 1) if no matching suites are found.

B.2.37 purge

Usage: cylc [control] purge [options] SUITE TASK STOP

Remove an entire tree of dependent tasks from a running suite. The root task will be forced to spawn and will then be removed, then so will every task that depends on it, and every task that depends on those, and so on until the given stop cycle time.

WARNING: THIS COMMAND IS DANGEROUS but in case of disaster you can restart the suite from the automatic pre-purge state dump (the filename

```
will be logged by cylc before the purge is actioned.)
```

UNDERSTANDING HOW PURGE WORKS: cylc identifies tasks that depend on the root task, and then on its downstream dependents, and then on theirs, etc., by simulating what would happen if the root task were to trigger: it artificially sets the root task to the "succeeded" state then negotiates dependencies and artificially sets any tasks whose prerequisites get satisfied to "succeeded"; then it negotiates dependencies again, and so on until the stop cycle is reached or nothing new triggers. Finally it marks "virtually triggered" tasks for removal. Consequently:

- * Dependent tasks will only be identified as such if they have already spawned into the root cycle, otherwise they will be missed by the purge. To avoid this, wait until all tasks that depend on the root have caught up to it before purging.
- * If you purge a task that has already finished, only it and its own successors will be purged (other downstream tasks will already have triggered if they were able to).

[development note: post cylc-3.0 we should be able to identify tasks that depend on the purge root by using the suite dependency graph, even if they have not spawned into the cycle time of the root yet.]

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
TASK	Task (NAME%CYCLE) at which to start the purge.
STOP	Cycle time (inclusive!) at which to stop purging.

Options:

-h, --help	show this help message and exit
--debug	Turn on exception tracebacks.
-f, --force	Do not ask for confirmation before acting.

B.2.38 refresh

```
Usage: cylc [db] refresh [options] SUITE
```

Check a suite database for invalid registrations (no suite definition directory or suite.rc file) and refresh suite titles in case they have changed since the suite was registered (for the central database this also updates the titles of suites owned by others).

Arguments:

SUITE	- suite name, or match pattern
-------	--------------------------------

Options:

-h, --help	show this help message and exit
-c, --central	Print suite registrations from the central database.
-u, --unregister	Automatically unregister invalid registrations.
-v, --verbose	Print extra information.

B.2.39 register

```
Usage: cylc [db] register [options] SUITE PATH
```

Register a suite in your private suite database.

Cylc commands target suites via names registered in a suite database so that you don't need to continually re-type the actual location of the suite definition directory on disk. Suite names are hierarchical like directory paths but delimited by '.' (foo.bar.baz), allowing suites to be grouped and organised into tree-like structures. Groups are virtual and do not need to be explicitly created before use.

Legal name characters are letters, digits, underscore, and hyphen. The period '.' is the registration hierarchy delimiter. Colon cannot be used because of the potential for trouble with PATH variables if the suite name is used in certain directory paths.

```
EXAMPLES, for suite definition directories $TMPDIR/(one,two,three,four):

% cylc db reg bob      $TMPDIR/one
% cylc db reg foo.bag   $TMPDIR/two
% cylc db reg foo.bar.baz $TMPDIR/three
% cylc db reg foo.bar.waz $TMPDIR/four

% cylc db pr '^foo'      # print in flat form
  bob          | "Test Suite One"    | /tmp/oliverh/one
  foo.bag      | "Test Suite Two"   | /tmp/oliverh/two
  foo.bar.baz | "Test Suite Four"  | /tmp/oliverh/three
  foo.bar.waz | "Test Suite Three" | /tmp/oliverh/four

% cylc db pr -t '^foo'    # print in tree form
  bob        "Test Suite One"    | /tmp/oliverh/one
  foo
    |-bag     "Test Suite Two"   | /tmp/oliverh/two
    '-bar
      |-baz    "Test Suite Three" | /tmp/oliverh/three
      '-waz    "Test Suite Four"  | /tmp/oliverh/four

Arguments:
  SUITE  -  The new hierarchical suite registration name.
  PATH   -  A cylc suite definition directory.

Options:
  -h, --help      show this help message and exit
  -v, --verbose   Print extra information.
```

B.2.40 release

```
Usage: cylc [control] release|unhold [options] SUITE [TASK]

Release a suite or a single task from a hold, allowing it run as normal.
Putting a suite on hold stops it from submitting any tasks that are
ready to run, until it is released. Putting a waiting task on hold
prevents it from running and spawning successors, until it is released.

See also 'cylc [control] hold'.

You must be the owner of the target suite to use this command.

Arguments:
  SUITE          Target suite.
  TASK           Task to release (NAME%YYYYMMDDHH)

Options:
  -h, --help      show this help message and exit
  --debug        Turn on exception tracebacks.
  -f, --force    Do not ask for confirmation before acting.
```

B.2.41 remove

```
Usage: cylc [control] remove|kill [options] SUITE TARGET

Remove a single task, or all tasks with a common TAG (cycle time or
asynchronous 'a:INT'), from a running suite.

Target tasks will be forced to spawn successors before being removed if
they have not done so already, unless you use '--no-spawn'.

You must be the owner of the target suite to use this command.

Arguments:
  SUITE          Target suite.
  TARGET         NAME%CYCLE or NAME%a:TAG to remove a single task;
                CYCLE or a:TAG to remove all tasks with the same tag.
```

```
Options:
-h, --help    show this help message and exit
--debug      Turn on exception tracebacks.
-f, --force   Do not ask for confirmation before acting.
--no-spawn   Do not spawn successors before removal.
```

B.2.42 reregister

```
Usage: cylc [db] reregister|rename [options] SOURCE TARGET

Change the hierarchical registration name of a suite or group.
Example:
  cylc db rereg foo.bar.baz test.baz

Arguments:
  SOURCE, TARGET - suite or group names.

Options:
-h, --help    show this help message and exit
-c, --central Target the central suite database.
-v, --verbose  Print extra information.
```

B.2.43 reset

```
Usage: cylc [control] reset [options] SUITE TASK

Force a task's state to:
1/ 'ready' .... (default) ..... all prerequisites satisfied (default)
2/ 'waiting' ... (--waiting) ..... prerequisites not satisfied yet
3/ 'succeeded' (--) ..... all outputs completed
4/ 'failed' ... (--failed) ..... use to test failure recovery

Resetting a task to 'ready' will cause it to trigger immediately unless
the suite is held, in which case the task will trigger when normal
operation is resumed.

You must be the owner of the target suite to use this command.

Arguments:
  SUITE          Target suite.
  TASK           The target task.

Options:
-h, --help    show this help message and exit
--debug      Turn on exception tracebacks.
-f, --force   Do not ask for confirmation before acting.
--ready     Force task to the 'ready' state.
--waiting   Force task to the 'waiting' state.
--succeeded Force task to 'succeeded' state.
--failed    Force task to 'failed' state.
```

B.2.44 restart

```
Usage: cylc [control] restart [options] SUITE [FILE]

Restart a cylc suite from a previous state. Tasks in the 'submitted',
'running', or 'failed' states will immediately be resubmitted at
start up unless you specify '--no-reset'. Any 'held' tasks will be
released unless you specify the '--no-release' option. A final cycle
time that was set prior to shutdown will be ignored on restart unless
you specify '--keep-finalcycle'

By default the suite will restart from the suite state dump file, which is
updated whenever a task changes state and thus records the most recent
previous state of the suite. However, cylc also records a special named
state dump, and logs its filename, before actioning any intervention
command, and you can choose to restart from one of these (just
cut-and-paste the filename from the log to the command line).
```

```
NOTE: daemonize important suites with the POSIX nohup command:
      nohup cylc [con] restart SUITE YYYYMMDDHH > suite.out 2> suite.err &
```

Arguments:

SUITE	Target suite.
FILE	Optional non-default state dump file (assumed to reside in the suite state dump directory unless you give the full path).

Options:

-h, --help	show this help message and exit
--no-reset	Do not reset failed tasks to ready on restarting
--no-release	Do not release held tasks on restarting.
--keep-finalcycle	Do not ignore a previously set final cycle time on restarting
--until=YYYYMMDDHH	Shut down after all tasks have PASSED this cycle time.
--hold	Hold (don't run tasks) immediately on starting.
--hold-after=YYYYMMDDHH	Hold (don't run tasks) AFTER this cycle time.
-s, --simulation-mode	Use dummy tasks that masquerade as the real thing, and accelerate the wall clock: get the scheduling right without having to run the real suite tasks.
--fail=NAME%YYYYMMDDHH	(SIMULATION MODE) get the specified task to report failure and then abort.
--debug	Turn on 'debug' logging and full exception tracebacks.
--timing	Turn on main task processing loop timing, which may be useful for testing very large suites of 1000+ tasks.
--gcylc	(DO NOT USE THIS OPTION).

B.2.45 run

Usage: cylc [control] run|start [options] SUITE [START]

Start a suite running at a specified initial cycle time. To restart a suite from a previous state (which may contain tasks at multiple cycle times), see 'cylc restart SUITE'.

There are three start up modes that differ in how initial intercycle dependencies are handled (at start up there is no previous cycle for tasks, such as warm cycled forecast models, that normally depend on tasks in previous cycles):

- 1/ Cold start (the default) -- start from scratch with special tasks
- 2/ Warm start (-w,--warm) -- assume a previous cycle
- 3/ Raw start (-r,--raw) -- assume nothing

These are all the same for a suite that has no intercycle dependencies. Otherwise, cold start tasks should be defined in the suite.rc file:

```
# SUITE.RC:  
[scheduling]->[[special tasks]]->cold-start = task1, task2, ...
```

1/ COLD START -- start from scratch, with special one off "cold start tasks" to satisfy the initial dependencies of any tasks that normally depend on tasks from a previous cycle (most notably the restart dependencies of warm cycled forecast models). Cold start tasks can run real processes (e.g. a special forecast to generate the initial restart files for a warm cycled model) or, with no task command defined, just act as a dummy proxy for some external process that has to occur before the suite is started (e.g. a spinup run of some kind). For a suite with no intercycle dependencies there is no distinction between the cold, warm, and raw start methods. For a cold start, each task, including designated cold start tasks, starts in the 'waiting' state (i.e. prerequisites not satisfied) at the initial cycle time or at the next valid cycle time thereafter for the task.

2/ WARM START -- this assumes that there was a previous cycle (from a previous suite run - in which case a 'cylc restart' may be an option too) or that files required from previous cycles have been put in place by some external means. Start up is as for a cold start, except that designated cold start tasks are inserted in the 'succeeded' state (i.e.

```
outputs completed).

3/ RAW START -- this starts the suite as if in mid run without any
special handling - any tasks that depend on previous cycles will have to
be triggered manually. Start up is as for a cold start, except that
designated cold start tasks are excluded from the suite.

NOTE: daemonize important suites with the POSIX nohup command:
$ nohup cylc [con] run SUITE [START] > suite.out 2> suite.err &

Arguments:
  SUITE           Target suite.
  START           Initial cycle time. This is optional if defined in the
                  suite.rc file, in which case the command line
                  takes priority (and a suite.rc final cycle time
                  will be ignored), and is not required if the
                  suite contains no cycling tasks.

Options:
  -h, --help      show this help message and exit
  -w, --warm      Warm start the suite
  -r, --raw       Raw start the suite
  --until=YYYYMMDDHH Shut down after all tasks have PASSED this cycle time.
  --hold          Hold (don't run tasks) immediately on starting.
  --hold-after=YYYYMMDDHH
                    Hold (don't run tasks) AFTER this cycle time.
  -s, --simulation-mode
                    Use dummy tasks that masquerade as the real thing, and
                    accelerate the wall clock: get the scheduling right
                    without having to run the real suite tasks.
  --fail=NAME%YYYYMMDDHH
                    (SIMULATION MODE) get the specified task to report
                    failure and then abort.
  --debug         Turn on 'debug' logging and full exception tracebacks.
  --timing        Turn on main task processing loop timing, which may be
                    useful for testing very large suites of 1000+ tasks.
  --gcylc        (DO NOT USE THIS OPTION).
```

B.2.46 scan

```
Usage: cylc [discover] scan [options]

Scan cylc ports for running suites and lockservers. Connection Denied
indicates a secure suite owned by somebody else.

Options:
  -h, --help      show this help message and exit
  --host=HOST    Cylc suite host (defaults to localhost).
  --print-ports  Print cylc's configured port range.
```

B.2.47 scp-transfer

```
Usage: cylc [util] scp-transfer [options]

An scp wrapper for transferring a list of files and/or directories
at once. The source and target scp URLs can be local or remote (scp
can transfer files between two remote hosts). Passwordless ssh must
be configured between the source and target hosts.

Environment inputs:
$SRCE  - list of sources (files or directories) as scp URLs.
$DEST  - parallel list of targets as scp URLs.

We let scp determine the validity of source and target URLs.
Target directories are created pre-copy if they don't exist.

Options:
  -v   - verbose: print scp stdout.
  --help - print this usage message.
```

B.2.48 search

```
Usage: cylc [prep] search|grep [options] PATTERN SUITE
```

Find matches to PATTERN in the suite.rc file, and any contained include-files, and in any files in the suite bin directory. For the suite.rc and include-files, matches are reported by file line number and suite section (even if include files are nested).

An unquoted list of PATTERNS will be converted to an OR'd pattern.

Note that the order of command line arguments conforms to grep command usage ('grep PATTERN FILE') not normal cylc command usage ('command SUITE ARGS'). However, if the second argument is not found to be a registered suite, an attempt will be made to action the command with the arguments reversed.

Note that this command performs a text search on the suite definition, rather than search all item names and values in the data structure that results from parsing the suite definition. Both types of search may be implemented in the future.

For case insensitive matching use '(?i)PATTERN'.

Arguments:

SUITE	- Suite database registration
PATTERN	(list of) Python-style regular expression(s).

Options:

-h, --help	show this help message and exit
-x	Do not search in the suite bin directory
-c, --central	Target the central database.

B.2.49 show

```
Usage: cylc [info] show [options] SUITE [NAME[%TAG]]
```

Print global or task-specific information from a running suite: title and task list, task descriptions, current state of task prerequisites and outputs and, for clock-triggered tasks, whether the delayed start time is up yet.

Arguments:

SUITE	Target suite.
NAME	Name of a task type.
TAG	Cycle time, or asynchronous tag 'a:INT'.

Options:

-h, --help	show this help message and exit
-o USER, --owner=USER	Owner of the target suite (defaults to \$USER).
--host=HOST	Cylc suite host (defaults to local host).
-f, --force	(No effect; for consistency with interactive commands)
--debug	Turn on exception tracebacks.

B.2.50 started

```
Usage: cylc [task] started [options]
```

This is part of the cylc external task interface.

Acquire a lock from the lockserver, and report that I have started.

Options:

-h, --help	show this help message and exit
------------	---------------------------------

B.2.51 stop

```
Usage: cylc [control] stop|shutdown [options] SUITE [STOP]

1/ Shut down a suite when all currently running tasks have finished.
   No other tasks will be submitted to run in the meantime.

2/ With [STOP], shut down a suite AFTER one of the following events:
   a/ all tasks have passed the TAG STOP (cycle time or async tag)
   b/ the clock time has reached STOP (YYYY/MM/DD-HH:mm)
   c/ the task STOP (TASK%TAG) has finished

3/ With [--now], shut down immediately, regardless of tasks still running.
   WARNING: beware of orphaning tasks that are still running at shutdown;
   these may need to be killed manually, and they will (by default) be
   resubmitted if the suite is restarted.

You must be the owner of the target suite to use this command.

Arguments:
  SUITE          Target suite.
  STOP           a/ task TAG (cycle time or 'a:INTEGER'), or
                 b/ YYYY/MM/DD-HH:mm (clock time), or
                 c/ TASK (task ID).

Options:
  -h, --help    show this help message and exit
  --debug      Turn on exception tracebacks.
  -f, --force   Do not ask for confirmation before acting.
  --now        Shut down immediately; see WARNING above.
```

B.2.52 submit

```
Usage: cylc [task] submit|single [options] SUITE TASK

Submit a single task to run exactly as it would be submitted by its
parent suite, in terms of both execution environment and job submission
method. This can be used as an easy way to run single tasks for any
reason, but it is particularly useful during suite development.

If the parent suite is running at the same time and it has acquired an
exclusive suite lock (which means you cannot run multiple instances
of the suite at once, even under different registrations) then the
lockserver will let you 'submit' a task from the suite only under the
same registration, and only if the task is not locked (i.e. only if
the same task, NAME%CYCLE, is not currently running in the suite).

Arguments:
  SUITE          Registered name of the target suite.
  TASK           Identity of the task to run (NAME%YYYYMMDDHH).

Options:
  -h, --help    show this help message and exit
  -d, --dry-run Generate the cylc task execution file for the task and report
                 how it would be submitted to run.
  --scheduler   (EXPERIMENTAL) tell the task to run as a scheduler task, i.e.
                 to attempt to communicate with a task proxy in a running cylc
                 suite (you probably do not want to do this).
```

B.2.53 succeeded

```
Usage: cylc [task] succeeded [options]

This is part of the cylc external task interface.

Release my lock to the lockserver, and report that I have succeeded.

Options:
  -h, --help    show this help message and exit
```

B.2.54 template

```
Usage: cylc [util] template [options] STRING [CYCLE]

Compute cycle time-dependent file names based on template strings, and
print the result to stdout. The base cycle time can be taken from the
environment or the command line, and offsets can be computed a la
'cylc cycletime'. Suite file naming conventions can thus be encoded
in template variables defined in the suite config file.

STRING can be the name of an environment variable containing a template
string OR the template string itself. In the template, 'YYYY', 'MM',
'DD', or 'HH' will be replaced with computed cycle time components.

$ export CYLC_TASK_CYCLE_TIME=2010082318

$ cylc template -s 6 fooHH-YYYYMMDD.nc #explicit name convention
foo12-20100823.nc

$ export NCONV=fooHH-YYYYMMDD.nc      # (as if from system config file)

$ cylc template NCONV #.....implicit filename convention
foo18-20100823.nc
$ cylc template -s 6 NCONV_FOO #..same, with current cycle offset
foo12-20100823.nc

Note that 'cylc template NCONV' and 'cylc template $NCONV' will
generate the same result, but the former is preferred so that we can
detect accidental use of undefined environment variables.

Arguments:
  STRING      A template string, or the name of an env variable that
              contains a template.
  CYCLE       (YYYYMMDDHH) defaults to $CYLC_TASK_CYCLE_TIME.

Options:
  -h, --help          show this help message and exit
  -s HOURS, --subtract=HOURS
                      Subtract HOURS from CYCLE
  -a HOURS, --add=HOURS
                      Add HOURS to CYCLE
  -o HOURS, --offset=HOURS
                      Apply an offset of +/-HOURS to CYCLE
```

B.2.55 test-db

```
USAGE: cylc [admin] test-db [--help]
This is a thorough test of cylc suite registration
database functionality (private and central).
Options:
  --help  Print this usage message.
```

B.2.56 test-suite

```
USAGE: cylc [admin] test-suite [--help]

Run an automated test of core cylc functionality using a new copy of
'/tmp/oliverh/20348/cylc-4.2.0/examples/admin/test'. This should be used to check that new
developments in the cylc codebase have not introduced serious bugs.
The test runs a suite registered as 'test'; to watch its progress
use 'cylc view'. Aside from timing differences results should be the
same in live or simulation mode.

Currently the test does the following:
- Copies the 'intro' example suite definition directory;
- Registers the new suite as ;
- Starts the suite at T0=06Z, with task X set to fail at 12Z;
- Unlocks the running suite;
- Sets a stop time at 12Z (i.e. T0+30 hours);
- Waits for the suite to stall as result of X failing;
- Inserts a new coldstart task at 06Z (T0+24 hours);
```

- Purges the failed X and dependants through to 00Z (T0+18 hours) inclusive, which allows the suite to get going again at 06Z;
- Waits for the suite to shut itself down at the 12Z stop time.
- Run a single task (called prep) from the suite with submit.

Options:

-h, --help Print this help message and exit.

B.2.57 trigger

Usage: cylc [control] trigger [options] SUITE TASK

Get a task to trigger immediately (unless the suite is paused, in which case it will trigger when normal operation is resumed). This is effected by setting the task to the 'ready' state (all prerequisites satisfied) and, for clock-triggered tasks, ignoring the designated trigger time.

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
TASK	The target task.

Options:

-h, --help show this help message and exit
--debug Turn on exception tracebacks.
-f, --force Do not ask for confirmation before acting.

B.2.58 unblock

Usage: cylc [control] unblock [options] SUITE

A blocked suite refuses to comply with intervention commands until deliberately unblocked. This is a crude security measure to guard against accidental intervention in your own suites. It may be a useful aid when running multiple suites at once.

You must be the owner of the target suite to use this command.

Arguments:

SUITE	Target suite.
-------	---------------

Options:

-h, --help show this help message and exit
--debug Turn on exception tracebacks.
-f, --force Do not ask for confirmation before acting.

B.2.59 unregister

Usage: cylc [db] unregister [OPTIONS] SUITE

Unregister suites from your private database, or remove your suites from the central suite database. This does not delete the corresponding suite definition directories unless you use the '--delete' option.

If you accidentally unregister a running suite, just register it again under the original name to regain access to it.

Arguments:

SUITE	- suite or group name, or filter expression.
-------	--

Options:

-h, --help show this help message and exit
-d, --delete Delete the suite definition directory too (!DANGEROUS!).
-f, --force Don't ask for confirmation before deleting suites.
-c, --central Target the central suite database.
--dry-run Just show what I would do.
--filter A regular expression to select multiple suites. This must be

```
    explicit for unregister, for safety.
-v, --verbose Print extra information.
```

B.2.60 validate

```
Usage: cylc [prep] validate [options] SUITE

Parse and validate a suite config (suite.rc) file to check that all
entries conform to the $CYLC_DIR/conf/suiterc.spec specification, and
then attempt to instantiate a proxy object for each task in the suite.

IF THE SUITE.RC FILE USES INCLUDE-FILES: line numbers reported with
validation errors will be wrong because the parser sees an inlined
version of the file. You can use 'cylc prep inline SUITE' to trace
errors to the correct source line, although the extra information
reported by the validator should be sufficient to make this unnecessary.

Arguments:
  SUITE      - Suite database registration.

Options:
  -h, --help    show this help message and exit
  -c, --central target the central database.
  -v, --verbose Print extra information from the validation process.
  -d, --debug   Print full exception traceback and abort on error.
  -p, --pretty   (with -v,--verbose) Use unicode box drawing characters when
                 printing the suite runtime inheritance hierarchy.
```

B.2.61 verbosity

```
Usage: cylc [control] verbosity [options] SUITE LEVEL

Change the logging priority level of a running suite. Only messages at
or above the chosen priority level will be logged; for example, if you
choose 'warning', only warning, error, and critical messages will be
logged. The 'info' level is appropriate under most circumstances.

You must be the owner of the target suite to use this command.

Arguments:
  SUITE          Target suite.
  LEVEL          debug, info, warning, error, or critical

Options:
  -h, --help    show this help message and exit
  --debug     Turn on exception tracebacks.
  -f, --force   Do not ask for confirmation before acting.
```

B.2.62 view

```
Usage: cylc [prep] view [options] SUITE

View a read-only temporary copy of a suite definition (suite.rc file) in
your editor. If the suite uses include-files you can view it inlined. If
it uses a template engine such as Jinja2 you can view the processed file.

The edit process is spawned in the foreground as follows:
$(G)EDITOR suite.rc
$GEDITOR or $EDITOR, and $TMDPIR, must be in your environment.

Examples:
  export EDITOR=vim
  export GEDITOR='gvim -f'      # -f: do not detach from parent shell!!
  export EDITOR='xterm -e vim'  # for gcylc, if gvim is not available
  export GEDITOR=emacs
  export EDITOR='emacs -nw'

You can set both $GEDITOR and $EDITOR to a GUI editor if you like, but
$EDITOR at least *must* be a GUI editor, or an in-terminal invocation
of a non-GUI editor, if you want to spawn editing sessions via gcylc.
```

C THE CYLC LOCKSERVER

```
See also 'cylc [prep] edit'.
```

Arguments:

```
SUITE      Target suite.
```

Options:

-h, --help	show this help message and exit
-i, --inline	Inline any include-files.
-p, --processed	View the suite after template engine processing. This necessarily implies '-i' as well.
-m, --mark	(With '-i') Mark inclusions in the left margin.
-l, --label	(With '-i') Label file inclusions with the file name. Line numbers will not correspond to those reported by the parser.
-s, --single	(With '-i') Inline only the first instances of any multiply-included files. Line numbers will not correspond to those reported by the parser.
-n, --nojoin	Do not join continuation lines (line numbers will not correspond to those reported by the parser).
--central	Target the central suite database.
-g, --gui	Use GUI editor \$GEDITOR instead of \$EDITOR. This option is automatically used when an editing session is spawned by cylc.
-o, --stdout	Print the suite definition to stdout instead of editing it.

B.2.63 warranty

```
USAGE: cylc [license] warranty [--help]
Cylc is released under the GNU General Public License v3.0
This command prints the GPL v3.0 disclaimer of warranty.
Options:
  --help  Print this usage message.
```

C The Cylc Lockserver

Each cylc user can optionally run his/her own lockserver to prevent accidental invocation of multiple instances of the same suite or task at the same time. The suite and task locks brokered by the lockserver are analogous to traditional lock files, but they work across a network, even for distributed suites containing tasks that start executing on one host and finish on another.

Accidental invocation of multiple instances of the same suite or task at the same time can have serious consequences, so use of the lockserver should be considered for important operational suites, but it may be considered an unnecessary complication for general less critical usage, so it is currently disabled by default.

To enable the lockserver:

```
# SUITE.RC
use lockserver = True
```

The suite will now abort at start-up if it cannot connect to the lockserver. To start your lockserver daemon,

```
% cylc lockserver start
```

To check that it is running,

```
% cylc lockserver status
```

For detailed usage information,

```
% cylc lockserver --help
```

There is a command line client interface,

```
% cylc lockclient --help
```

E SIMULATION MODE

for interrogating the lockserver and managing locks manually (e.g. releasing locks if a suite was killed before it could clean up after itself).

To watch suite locks being acquired and released as a suite runs,

```
% watch cylc lockclient --print
```

D The Graph-Based Suite Control GUI

The graph-based suite control GUI has the advantage that it shows the structure of a suite very clearly as it evolves. It works remarkably well even for very large suites, on the order of one hundred tasks or more *but* on the downside, the graphviz engine does a new global optimization every time the graph changes, so the layout is often not very stable. This may or may not be a solvable problem - it seems likely that making continual incremental changes to an existing graph without redoing the global layout would inevitably result in some kind of horrible mess.

The following features of the graph-based control GUI go a long way toward mitigating the changing layout problem:

- The disconnect button can be used to temporarily prevent the graph from changing as the content of the suite changes (and in real time operation suites evolve quite slowly anyway)
- Right-click on a task and choose the “Focus” option to restrict the graph display to that task’s cycle time. Anything interesting happening in other cycles will show up as disconnected rectangular nodes to the right of the graph (and you can click on those to instantly refocus to their cycles).
- Task filtering is the ultimate quick route to temporarily focusing on just the tasks you’re interested in (but this will destroy the graph structure, to state the obvious).

In future cylc releases we plan to keep the graph centered, after layout changes, on the most recently clicked-on task.

E Simulation Mode

If you run a suite in simulation mode its tasks will be replaced by dummy tasks that just print a message and then exit after a few seconds, and the suite will run quickly on an accelerated clock. As far as cylc is concerned this is essentially identical to real operation. Simulation mode can be used to get the scheduling of a complex suite right without having to run real tasks, and to test recovery strategies for simulated task failure scenarios.

By configuring how the accelerated clock is initialized you can watch suites catch up and transition from delayed to real time operation.

Remote task hosting and job submission configuration is ignored in simulation mode so you can simulate any suite, even outside of its normal operating environment.

A more realistic but less portable simulation mode, with full remote job submission etc., can be achieved by simply commenting out task command scripting so that it defaults to dummy command scripting, and running the suite in live mode (note however that this won’t work for tasks with internal outputs unless you provide your own dummy command scripting that reports the internal outputs finished).

Simulation mode was, and remains, an important aid to cylc development because it allows testing of all aspects of scheduling without having to run real tasks in real time. Prior to cylc-3 it was also a critical aid to suite construction, quickly identifying any mismatch between the user-defined task prerequisites and outputs across a suite. Post cylc-3.0 this is less important

because task prerequisites and outputs are now implicitly defined by the suite dependency graph and, short of a bug in cylc, the suite will run according to the graph.

E.1 Clock Rate and Offset

You can set the simulation mode clock rate and offset with respect to the initial cycle time with options to the `cylc run` command. An offset of 10 hours, say, means that the clock starts at 10 hours prior to the suite's initial cycle time. You can thus simulate the behaviour of the suite as it catches up from a delay and transitions to real time operation. By default, the clock runs at a rate of 10 seconds real time to 1 hour suite time, and with an initial offset of 10 hours.

E.2 Switching A Suite Between Simulation And Live Modes?

The scheduler mode (simulation or live) is recorded in the suite state dump file. Cylc will not let you *restart* a simulation mode suite in live mode, or vice versa - any attempt to do the former would certainly be a mistake (because the simulation mode dummy tasks do not generate any of the real outputs depended on by downstream live tasks), and the latter, while feasible, would corrupt the live state dump and turn it over to simulation mode. The easiest way to test a live suite in simulation mode, if you don't want to obliterate the current state dump by doing a cold or warm start (as opposed to a restart from the previous state) is to take a quick copy of the suite and run the copy in simulation mode. However, if you really want to run a live suite forward in simulation mode without copying it, do this:

1. Backup the live state dump file.
2. Edit the mode line in the state dump file and restart in simulation mode.
3. Later, restart the live suite from the restored live state dump backup.

Finally, it should be noted that cylc previously had the ability to create an instant safe simulation mode clone of the current state of any suite, running or stopped, but this “practice mode” has been disabled since the upgrade to cylc-3, pending testing; it could easily be restored if needed.

F Cylc Development History

F.1 Pre-3.0

Early versions of cylc were focused on developing and testing the new scheduling algorithm, and the suite design interface at the time was essentially the quickest route to that end. A suite was a collection of “task definition files” that encoded the prerequisites and outputs of each task in a direct reflection of cylc’s internal task proxies. This way of defining suites exposed cylc’s self-organising nature to the user, and it did have some nice properties. For instance a group of tasks could be transferred directly from one suite to another by simply copying the taskdef files over (and checking that prerequisite and output messages were consistent with the new suite). However, ensuring consistency of prerequisites and outputs across a large suite could be tedious; a few edge cases associated with suite start-up and forecast model restart dependencies were, arguably, difficult to understand; and the global structure of a suite was not readily apparent until run time (although to counter this cylc 2.x could generate run-time resolved dependency graphs very quickly in simulation mode).

F.2 Version 3.0

Version 3.0 implemented an entirely new suite design interface in which one defines the suite dependency graph, execution environment, and command scripting for each task, in a single structured, validated, configuration file - the suite.rc file. This *really* makes suite structure apparent at a glance, and task prerequisites and outputs (and some other important parameters besides) no longer need to be specified by the user because they are implied by the graph.

F.3 Version 4.0

Version 4.0 has the following major improvements over cylc-3.x, along with many refinements:

- Suite registration has been generalized from GROUP:NAME to a hierarchy of arbitrary depth, e.g foo.bar.baz, allowing suites to be organized in a tree-like structure.
- The suite.rc file has been extensively reorganized. In particular it now defines a *namespace hierarchy* of families and tasks to allow common runtime properties to be grouped naturally among related tasks.
- The Jinja2 template engine is now supported, adding general programmability to cylc suite definitions.

G Pyro

Pyro (Python Remote Objects) is a widely used open source object oriented Remote Procedure Call technology developed by Irmel de Jong.

Earlier versions of cylc used the Pyro Nameserver to marshal communication between client programs (tasks, commands, viewers, etc.) and their target suites. This worked well, but in principle it provided a route for one suite or user on the subnet to bring down all running suites by killing the nameserver. Consequently cylc now uses Pyro simply as a lightweight object oriented wrapper for direct network socket communication between client programs and their target suites - all suites are thus entirely isolated from one another.

H Troubleshooting

H.1 pre-3.9.1 Pyro Incompatibility

Attempting to run `cylc admin test-suite` on a system using Pyro 3.9 or earlier will result in the following Python traceback:

```
Traceback (most recent call last):
  File "/home/oliverh/cylc/bin/_run", line 232, in <module>
    server = start()
  File "/home/oliverh/cylc/bin/_run", line 92, in __init__
    scheduler.__init__( self )
  File "/home/oliverh/cylc/lib/cylc/scheduler.py", line 141, in __init__
    self.load_tasks()
  File "/home/oliverh/cylc/bin/_run", line 141, in load_tasks_cold
    itask = self.config.get_task_proxy( name, tag, 'waiting',
    stopctime=None, startup=True )
  File "/home/oliverh/cylc/lib/cylc/config.py", line 1252, in get_task_proxy
    return self.taskdefs[name].get_task_class()( ctime, state,
    stopctime, startup )
  File "/home/oliverh/cylc/lib/cylc/taskdef.py", line 453, in tclass_init
    print '-' , sself.__class__.__name__, sself.__class__.__bases__
```

J GNU GENERAL PUBLIC LICENSE V3.0

```
AttributeError: type object 'A' has no attribute '_taskdef__bases_'
/run --debug testsuite.1322742021 2010010106 failed: 1
```

Solution: upgrade Pyro - see *Requirements*, Section 3.

I Acknowledgements

Bernard Miville and Phil Andrews (NIWA) for discussion, bug finding, and many suggestions that have improved cylc's usability and functionality; David Matthews and Matthew Shin (Met Office UK) for the same, and much code development as well.

J GNU GENERAL PUBLIC LICENSE v3.0

Copyright © 2007 Free Software Foundation, Inc. <http://fsf.org/>

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

PREAMBLE

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of

such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS

0. Definitions.

“This License” refers to version 3 of the GNU General Public License.

“Copyright” also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

“The Program” refers to any copyrightable work licensed under this License. Each licensee is addressed as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users’ Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights

J GNU GENERAL PUBLIC LICENSE V3.0

under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- (a) The work must carry prominent notices stating that you modified it, and giving a relevant date.
- (b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".
- (c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- (d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an "aggregate" if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation's users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- (a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.

- (b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- (c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- (d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- (e) Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a

fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- (a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- (b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- (c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or
- (d) Limiting the use for publicity purposes of names of licensors or authors of the material; or
- (e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- (f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of

liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor’s “contributor version”.

A contributor’s “essential patent claims” are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, “control” includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor’s essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a “patent license” is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To “grant” such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. “Knowingly relying” means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient’s use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does.>

Copyright (C) <textyear> <name of author>

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

J GNU GENERAL PUBLIC LICENSE V3.0

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

```
<program> Copyright (C) <year> <name of author>
```

```
This program comes with ABSOLUTELY NO WARRANTY; for details type 'show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.
```

The hypothetical commands `show w` and `show c` should show the appropriate parts of the General Public License. Of course, your program's commands might be different; for a GUI interface, you would use an “about box”.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see <http://www.gnu.org/licenses/>.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read <http://www.gnu.org/philosophy/why-not-lgpl.html>.